
Amazon Kendra

Developer Guide



Amazon Kendra: Developer Guide

Copyright © 2020 Amazon Web Services, Inc. and/or its affiliates. All rights reserved.

Amazon's trademarks and trade dress may not be used in connection with any product or service that is not Amazon's, in any manner that is likely to cause confusion among customers, or in any manner that disparages or discredits Amazon. All other trademarks not owned by Amazon are the property of their respective owners, who may or may not be affiliated with, connected to, or sponsored by Amazon.

Table of Contents

.....	viii
What is Amazon Kendra?	1
Benefits of Amazon Kendra	1
Amazon Kendra Developer Edition	1
Amazon Kendra Enterprise Edition	2
Pricing for Amazon Kendra	2
Are you a first-time Amazon Kendra user?	2
How Amazon Kendra works	3
Index	3
Index Fields	4
Searching Indexes	4
Documents	4
Types of documents	5
Document attributes	6
Data sources	6
Queries	7
Tags	7
Tagging resources	7
Tag restrictions	8
Setting up Amazon Kendra	9
Sign up for AWS	9
Regions and endpoints	9
Setting up the AWS CLI	9
Setting up the AWS SDKs	10
Getting started	11
Prerequisites	11
Prerequisites for the AWS CLI and SDK	11
Getting started with an S3 bucket (Console)	14
Getting started (AWS CLI)	15
Getting started (SDK for Python (Boto 3))	16
Getting started (SDK for Java)	18
Getting started with OneDrive (Console)	21
Getting started with an S3 data source (Console)	22
Getting started with Salesforce (Console)	23
Getting started with ServiceNow (Console)	24
Getting started with SharePoint Online (Console)	25
IAM access roles for Amazon Kendra	27
IAM roles for indexes	27
IAM Roles for the BatchPutDocument operation	28
IAM roles for data sources	29
IAM roles for Amazon S3 data sources	29
IAM roles for database data sources	30
IAM roles for OneDrive data sources	32
IAM role for Salesforce data sources	34
IAM role for ServiceNow data sources	35
IAM roles for SharePoint Online data sources	36
IAM roles for frequently asked questions	37
Creating an index	39
Adding documents directly to an index	42
Adding documents with the API	42
Adding documents from an Amazon S3 bucket	43
Adding questions and answers	45
Adding documents from a data source	46
Setting an update schedule	47

Using an Amazon S3 data source	47
Using a database data source	50
Using a OneDrive data source	52
Using a Salesforce data source	54
Using a ServiceNow data source	57
Using a SharePoint data source	60
Deleting Data Sources	61
Creating custom document attributes	62
Adding custom attributes with the BatchPutDocument operation	63
Adding custom attributes to an S3 data source	63
Mapping data source fields	63
Configuring Amazon Kendra to use a VPC	65
Connecting to a database in a VPC	66
Searching indexes	69
Querying an index	69
Prerequisites	70
Searching an index (Console)	70
Searching an index (SDK)	70
Filtering queries	72
Facets	72
Using document attributes to filter queries	73
Filtering a document's attributes	74
Filtering on user context	74
User context filtering for documents added directly to an index	75
User context filtering for frequently asked questions	75
User context filtering for database data sources	75
User context filtering for Microsoft OneDrive data sources	76
User context filtering for Amazon S3 data sources	76
User context filtering for Salesforce data sources	77
User context filtering for ServiceNow data sources	77
User context filtering for Microsoft SharePoint data sources	77
Query responses	77
Types of response	79
Answer	79
Document	80
Question and answer	81
Submitting feedback	82
Deploying Amazon Kendra	85
Overview	85
Prerequisites	85
Setting up the example	86
Main search page	86
Search component	86
Results component	86
Facets component	87
Pagination component	87
Troubleshooting	88
My synchronization job failed	88
My synchronization job is incomplete	88
My synchronization job succeeded but there are no indexed documents	89
General troubleshooting	89
Manually tuning an index	91
Adjusting capacity	93
Viewing capacity	93
Adding and removing capacity	93
Monitoring and Logging	95
Monitoring indexes	95

Monitoring Amazon Kendra API calls with CloudTrail	98
Amazon Kendra Information in CloudTrail	99
Example: Amazon Kendra log file Entries	99
Monitoring Amazon Kendra with CloudWatch	100
Viewing Amazon Kendra metrics	100
Creating an alarm	100
CloudWatch Metrics for index synchronization Jobs	101
Metrics for Amazon Kendra data sources	102
Metrics for indexed documents	104
Monitoring Amazon Kendra with CloudWatch Logs	104
data source log streams	105
Document log streams	106
Security	107
Data protection	107
Encryption at rest	108
Encryption in transit	108
Key management	108
Identity and access management	109
Audience	109
Authenticating with identities	109
Managing access using policies	111
How Amazon Kendra works with IAM	113
Identity-based policy examples	116
Logging and monitoring in Amazon Kendra	121
Compliance validation	121
Resilience	122
Infrastructure security	122
Quotas	123
Supported regions	123
Quotas	123
Document history	124
API Reference	125
Actions	125
BatchDeleteDocument	126
BatchPutDocument	129
CreateDataSource	132
CreateFaq	139
CreateIndex	142
DeleteDataSource	145
DeleteFaq	147
DeleteIndex	149
DescribeDataSource	151
DescribeFaq	158
DescribeIndex	162
ListDataSources	166
ListDataSourceSyncJobs	169
ListFaqs	172
ListIndices	175
ListTagsForResource	177
Query	179
StartDataSourceSyncJob	185
StopDataSourceSyncJob	187
SubmitFeedback	189
TagResource	191
UntagResource	193
UpdateDataSource	195
UpdateIndex	201

Data Types	203
AccessControlListConfiguration	206
AclConfiguration	207
AdditionalResultAttribute	208
AdditionalResultAttributeValue	209
AttributeFilter	210
BatchDeleteDocumentResponseFailedDocument	212
BatchPutDocumentResponseFailedDocument	213
CapacityUnitsConfiguration	214
ClickFeedback	215
ColumnConfiguration	216
ConnectionConfiguration	218
DatabaseConfiguration	220
DataSourceConfiguration	221
DataSourceSummary	223
DataSourceSyncJob	225
DataSourceSyncJobMetrics	227
DataSourceSyncJobMetricTarget	229
DataSourceToIndexFieldMapping	230
DataSourceVpcConfiguration	231
Document	232
DocumentAttribute	234
DocumentAttributeValue	235
DocumentAttributeValueCountPair	236
DocumentMetadataConfiguration	237
DocumentsMetadataConfiguration	238
Facet	239
FacetResult	240
FaqStatistics	241
FaqSummary	242
Highlight	244
IndexConfigurationSummary	245
IndexStatistics	247
OneDriveConfiguration	248
OneDriveUsers	250
Principal	251
QueryResultItem	252
Relevance	254
RelevanceFeedback	256
S3DataSourceConfiguration	257
S3Path	259
SalesforceChatterFeedConfiguration	260
SalesforceConfiguration	262
SalesforceCustomKnowledgeArticleTypeConfiguration	265
SalesforceKnowledgeArticleConfiguration	267
SalesforceStandardKnowledgeArticleTypeConfiguration	268
SalesforceStandardObjectAttachmentConfiguration	269
SalesforceStandardObjectConfiguration	270
Search	272
ServerSideEncryptionConfiguration	273
ServiceNowConfiguration	274
ServiceNowKnowledgeArticleConfiguration	276
ServiceNowServiceCatalogConfiguration	278
SharePointConfiguration	280
Tag	283
TextDocumentStatistics	284
TextWithHighlights	285

TimeRange	286
Common Errors	286
Common Parameters	288
AWS glossary	290

What is Amazon Kendra?

Amazon Kendra is an enterprise search service that enables your users to intuitively search unstructured data using natural language. It returns specific answers to questions, giving users an experience that's close to interacting with a human expert. It is highly available and scalable, tightly integrated with other AWS services, and offers enterprise-grade security.

Amazon Kendra users can ask the following types of questions, or queries:

- **Factoid questions** — Simple who, what, when, or where questions, such as *Who is Amazon's CEO?* or *What is the height of the Space Needle?*. Factoid questions have fact-based answers that can be returned in the form of a single word or phrase. The precise answer, however, must be explicitly stated in the ingested text content.
- **Descriptive questions** — Questions whose answer could be a sentence, passage, or an entire document. For example, *How do I connect my Echo Plus to my network?* or *How do I get tax benefits for lower income families?*.
- **Keyword searches** — For questions where the intent and scope isn't clear, for example, *vacation policy* or *health benefits*, Amazon Kendra uses its deep learning models to return relevant documents.

Benefits of Amazon Kendra

Amazon Kendra has the following benefits:

- **Accuracy** — Unlike traditional search services that use keyword searches where results are based on basic keyword matching and ranking, Amazon Kendra attempts to understand the content, the user context, and the question. Amazon Kendra searches across your data and goes beyond traditional search to return the most relevant word, snippet, or document for your query. Amazon Kendra uses machine learning to improve search results over time.
- **Simplicity** — Amazon Kendra provides a console and API for managing the documents that you want to search. You can use a simple search API to integrate Amazon Kendra into your client applications, such as websites or mobile applications.
- **Connectivity** — Amazon Kendra can connect to third-party data sources to provide search across documents managed in different environments.

Amazon Kendra Developer Edition

The Amazon Kendra Developer Edition provides all of the features of Amazon Kendra at a lower cost. It includes a free tier that provides 750 hours of use. The Developer Edition is ideal to explore how Amazon Kendra indexes your documents, to try out features, and to develop applications that use Amazon Kendra.

The developer edition provides the following:

- Up to 5 indexes with up to 5 data sources each
- 10,000 documents or 3 Gb of extracted text
- 4,000 queries per day
- Runs in 1 availability zone (AZ)

You should not use the Developer Edition for a production application. The Developer Edition doesn't provide any guarantees of latency or availability.

Amazon Kendra Enterprise Edition

Use Amazon Kendra Enterprise Edition when you want to index your entire enterprise document library or for when your application is ready for use in a production environment. The enterprise edition provides the following. You can increase this quota using the [Service Quotas console](#).

- Up to 5 indexes with up to 50 data sources each.
- 500,000 documents or 150 Gb of extracted text.
- 40,000 queries per day.
- Runs in 3 availability zones (AZ)

Pricing for Amazon Kendra

You can get started for free with the Amazon Kendra Developer Edition that provides usage of up to 750 hours for the first 30 days. After your trial expires, you are charged for all provisioned Amazon Kendra indexes, even if they are empty and no queries are executed. After the trial expires, there are additional charges for scanning and syncing documents using the Amazon Kendra data sources.

For a complete list of charges and prices, see [Amazon Kendra pricing](#).

Are you a first-time Amazon Kendra user?

If you are a first-time user of Amazon Kendra, we recommend that you read the following sections in order:

1. [How Amazon Kendra works \(p. 3\)](#) – Introduces the Amazon Kendra components and describes how you use them to create a search solution.
2. [Getting started \(p. 11\)](#) – Explains how to set up your account and test the Amazon Kendra search API.
3. [Creating an index \(p. 39\)](#) – Explains how to use Amazon Kendra to create a search index and to add data sources to sync your documents.
4. [Adding documents directly to an index \(p. 42\)](#) – Explains how to add documents directly to an Amazon Kendra index.
5. [Searching indexes \(p. 69\)](#) – Explains how to use the Amazon Kendra search API to search an index.
6. [Deploying Amazon Kendra \(p. 85\)](#) – Provides a sample application you can use to deploy Amazon Kendra to your website.

How Amazon Kendra works

Amazon Kendra provides an interface for indexing and searching documents. You can use Amazon Kendra to create an updatable index of documents of a variety of types, including plain text, HTML files, Microsoft Word documents, Microsoft PowerPoint presentations, and PDF files. It has a search API that you can use from a variety of client applications, such as websites or mobile applications.

Amazon Kendra has the following components:

- The *index*, which provides a search API for client queries. You create the index from source documents.
- A *source repository*, which contains the documents to index.
- A *data source* that syncs the documents in your source repositories to an Amazon Kendra index. You can automatically synchronize a data source with an Amazon Kendra index so that new, updated, and deleted files in the source repository are updated in the index.
- A *document addition API*, that adds documents directly to the index.

To manage indexes and data sources, you can use the Amazon Kendra console or the API. You can create, update, and delete indexes. Deleting an index deletes all data sources and permanently deletes all of your document information from Amazon Kendra.

Topics

- [Index \(p. 3\)](#)
- [Documents \(p. 4\)](#)
- [Data sources \(p. 6\)](#)
- [Queries \(p. 7\)](#)
- [Tags \(p. 7\)](#)

Index

An *index* provides search results for documents and frequently asked questions (FAQ) that it has indexed. The way you add documents to the index depends on the type of document and where they are stored.

- If your documents are in a store, such as an S3 bucket or a Microsoft SharePoint site, you use a data source .
- You use the Amazon Kendra API to add documents
- For FAQ questions and answers, which must be stored in an Amazon Simple Storage Service ((Amazon S3) bucket, you upload them from the bucket.

An index can contain documents that are indexed from a data source, documents that are added directly to the index, and FAQs. You can create indexes with the Amazon Kendra console, the AWS CLI, or an AWS SDK. For information about the types of documents that can be indexed, see [Types of documents \(p. 5\)](#).

For information about using a data source with an index, see [Data sources \(p. 6\)](#).

To add documents directly to an index, you call the [BatchPutDocument \(p. 129\)](#) operation. The documents are supplied as plain text, as a binary blob, or using a path to a document stored in an Amazon S3 bucket. For an example, see [Adding documents from an Amazon S3 bucket \(p. 43\)](#).

Index Fields

An index contains fields that you map to the attributes of your document. Fields provide the schema for your index. Amazon Kendra uses the fields to search your documents. After you map your documents' attributes to fields, you can use the information in the field for searching, for display, and to create facets of the search result.

Amazon Kendra has six reserved fields, which you can map to your document attributes:

- `_category` – The category of the document.
- `_created_at` – The date and time that the document was created.
- `_file_type` – The file type of the document (html, pdf, and so on).
- `_last_updated_at` – The date and time that the document was last updated.
- `_version` – The version of the document.
- `_view_count` – The number of times that the document has been viewed.

You can also create custom fields, which you can use like the reserved fields for search and display, and to create facets.

There are four types of custom fields:

- Date
- Number
- String
- String list

You create a custom field using the console or by using the [UpdateIndex \(p. 201\)](#) operation. After you create a custom field, you map it to a document attribute, just as you do with a reserved field. If you added a document to the index with [BatchPutDocument \(p. 129\)](#) operation, you map the attributes with the API. For documents indexed from an Amazon S3 data source, you map the attributes using a metadata file that contains a JSON structure that describes the document attributes. For documents indexed with a database or SharePoint Online data source, you map attributes with the console or the data source configuration. For more information, see [Document attributes \(p. 6\)](#).

Searching Indexes

After you create an index, you can use it for search queries. For more information, see [Searching indexes \(p. 69\)](#).

Documents

Amazon Kendra can index many types of documents. You can also associate attributes with documents to provide information such as the source URI and the author of a document.

Topics

- [Types of documents \(p. 5\)](#)
- [Document attributes \(p. 6\)](#)

Types of documents

An index can include both structured and unstructured text:

- Structured text
 - Frequently asked questions and answers
- Unstructured text
 - HTML files
 - Microsoft PowerPoint presentations
 - Microsoft Word documents
 - Plain text documents
 - PDFs

You can add documents directly to an index by calling the [BatchPutDocument \(p. 129\)](#) operation. You can also add documents from a data source. For information about adding files to a data source, see [Adding documents from a data source \(p. 46\)](#). For an example that shows how to add Microsoft Word documents directly to an index from an Amazon S3 bucket, see [Adding documents from an Amazon S3 bucket \(p. 43\)](#).

An index can contain multiple documents and multiple types of documents.

HTML

HTML format files. You add an HTML file to an index the same way that you add a plain text file.

Plain text

You can add plain text files to an index using the `BatchPutDocument` operation or from a data source. For an example of adding a plain text document directly to an index, see [Adding documents with the API \(p. 42\)](#).

Microsoft Word document

Microsoft Word format files can be added to an index as binary data, from an Amazon S3 bucket, or from an Amazon Kendra data source.

Microsoft PowerPoint document

Microsoft PowerPoint format files can be added to an index as binary data, from an Amazon S3 bucket, or from an Amazon Kendra data source.

Portable document format (PDF)

PDF format files can be added to an index either as binary data, from an Amazon S3 bucket, or from an Amazon Kendra data source.

Frequently asked questions and answers

Frequently asked question and answer format documents are used to answer questions such as *How tall is the Space Needle?* You can specify multiple questions that return the same answer. You specify the questions and answers in a comma-separated values (CSV) file stored in an Amazon S3 bucket.

For an example, see [Adding questions and answers \(p. 45\)](#).

Document attributes

A document has attributes associated with it. You can also add your own custom attributes. Custom attributes are attributes that you can specify for your own needs. For example, if your index searches tax documents, you might specify a custom attribute for the type of tax document (W-2, 1099, and so on).

Before you can use a document attribute in a query, it must be mapped to a database field. For more information, see [Index Fields \(p. 4\)](#).

You can use document attributes to filter responses and to make faceted search suggestions. For example, you can filter a response to only return a specific version of a document, or you can filter searches to only return 1099 tax documents that match the search term. For more information, see [Filtering queries \(p. 72\)](#).

You can also use document attributes to manually tune the query response. For example, you can choose to increase the importance of the title field to increase the weight that Amazon Kendra assigns to the field when determining which documents to return in the response. For more information, see [Manually tuning an index \(p. 91\)](#).

Before you can add an attribute, you must create an index field to map the attribute to. You create index fields using the console or by using the [UpdateIndex \(p. 201\)](#) operation.

If you are adding a document directly to an index, you specify the attributes in the [Document \(p. 232\)](#) input parameter to the [BatchPutDocument \(p. 129\)](#) operation. You specify the custom attribute values in a [DocumentAttribute \(p. 234\)](#) object array. If you are using a data source, the method that you use to add the document attributes depends on the data source. For more information, see [Creating custom document attributes \(p. 62\)](#).

Data sources

A *data source* is a location, such as an Amazon Simple Storage Service (Amazon S3) bucket, where you store the documents for indexing. Data sources can be automatically synchronized with an Amazon Kendra index so that new, updated, or deleted documents in the source repositories are included in searches. Supported data sources are:

- Amazon S3 buckets
- Amazon RDS for MySQL and Amazon RDS for PostgreSQL databases
- Microsoft OneDrive for Business
- Microsoft SharePoint Online
- Salesforce sites
- ServiceNow instances

Supported document formats are: plain text, Microsoft Word, Microsoft PowerPoint, HTML, and PDF. For more information, see [Types of documents \(p. 5\)](#).

Note

To create an index, you don't need a data source. You can add documents directly to an index. For more information, see [Adding documents directly to an index \(p. 42\)](#).

To index documents using a data source.

1. [Create an index \(p. 39\)](#).
2. [Create a data source \(p. 46\)](#).

For a walkthrough with the Amazon Kendra console or with the AWS CLI, see [Getting started \(p. 11\)](#).

Queries

To get answers, users query an index. Users can use natural language in their queries. The response contains information, such as the title, a text excerpt, and the location of documents in the index that provide the best answer.

Amazon Kendra uses all of the information that you provide about your documents, not just the contents of the documents, to determine whether a document is relevant to the query. For example, if your index contains information about when documents were last updated, you can tell Amazon Kendra to assign a higher relevance to documents that were updated more recently.

A query can also contain criteria for how to filter the response so that Amazon Kendra returns only documents that satisfy the filter criteria. For example, if you created an index field called *department*, you can filter the response so that only documents with the department field set to *legal* are returned. For more information, see [Filtering queries \(p. 72\)](#).

You can influence the results of a query by tuning the relevance of individual fields in the index. Tuning changes the importance of a field on the results. For example, if you raise the importance of documents with the file type *pdf*, PDF files are more likely to be included in the response. For more information, see [Manually tuning an index \(p. 91\)](#).

For more information about using queries, see [Searching indexes \(p. 69\)](#).

Tags

Manage your indexes, data sources, and FAQs by assigning metadata to them with *tags*. Use tags to categorize your Amazon Kendra resources in various ways, for example, by purpose, owner, or application, or any combination. Each tag consists of a *key* and a *value*, both of which you define.

Tags help you to:

- Identify and organize your AWS resources. Many AWS services support tagging, so you can assign the same tag to resources in different services to indicate that the resources are related. For example, you can tag an index and the Amazon Lex bot that uses it with the same tag.
- Allocate costs. You activate tags on the AWS Billing and Cost Management dashboard. AWS uses tags to categorize your costs and deliver a monthly cost allocation report to you. For more information, see [Cost Allocation and Tagging](#) in *About AWS Billing and Cost Management*.
- Control access to your resources. You can use tags in AWS Identity and Access Management (IAM) policies that control access to Amazon Kendra resources. You can attach these policies to an IAM role or user to enable tag-based access control. For more information, see [Authorization based on Amazon Kendra tags \(p. 115\)](#).

You can create and manage tags using the AWS Management Console, the AWS Command Line Interface (AWS CLI), or the Amazon Kendra API.

Tagging resources

If you're using the Amazon Kendra console, you can tag resources when you create them or add them later. You can also use the console to update or remove tags.

If you're using the AWS Command Line Interface (AWS CLI) or the Amazon Kendra API, use the following operations to manage tags for your resources:

- [CreateDataSource \(p. 132\)](#) – Apply tags when you create a data source.
- [CreateFaq \(p. 139\)](#) – Apply tags when you create an FAQ.
- [CreateIndex \(p. 142\)](#) – Apply tags when you create an index.
- [ListTagsForResource \(p. 177\)](#) – View the tags associated with a resource.
- [TagResource \(p. 191\)](#) – Add and modify tags for a resource.
- [UntagResource \(p. 193\)](#) – Remove tags from a resource.

Tag restrictions

The following restrictions apply to tags on Amazon Kendra resources:

- Maximum number of tags – 50
- Maximum key length – 128 characters
- Maximum value length – 256 characters
- Valid characters for key and value – a–z, A–Z, space, and the following characters: `_` `:` `/` `=` `+` `-` and `@`
- Keys and values are case sensitive
- Don't use `aws:` as a prefix for keys; it's reserved for AWS use

Setting up Amazon Kendra

Before using Amazon Kendra, you must have an Amazon Web Services (AWS) account. After you have an AWS account, you can access Amazon Kendra through the Amazon Kendra console, the AWS Command Line Interface (AWS CLI), or the AWS SDKs.

This guide includes examples for AWS CLI, Java, and Python.

Topics

- [Sign up for AWS](#) (p. 9)
- [Regions and endpoints](#) (p. 9)
- [Setting up the AWS CLI](#) (p. 9)
- [Setting up the AWS SDKs](#) (p. 10)

Sign up for AWS

When you sign up for Amazon Web Services (AWS), your account is automatically signed up for all services in AWS, including Amazon Kendra. You are charged only for the services that you use.

If you have an AWS account already, skip to the next task. If you don't have an AWS account, use the following procedure to create one.

To sign up for AWS

1. Open <https://aws.amazon.com>, and then choose **Create an AWS Account**.
2. Follow the on-screen instructions to complete the account creation. Note your 12-digit AWS account number. Part of the sign-up procedure involves receiving a phone call and entering a PIN using the phone keypad.
3. Create an AWS Identity and Access Management (IAM) admin user. See [Creating Your First IAM User and Group](#) in the *AWS Identity and Access Management User Guide* for instructions.

Regions and endpoints

An endpoint is a URL that is the entry point for a web service. Each endpoint is associated with a specific AWS region. If you use a combination of the Amazon Kendra console, the AWS CLI, and the Amazon Kendra SDKs, pay attention to their default regions as all Amazon Kendra components of a given campaign (index, query, etc.) must be created in the same region. For the regions and endpoints supported by Amazon Kendra, see [Regions and Endpoints](#).

Setting up the AWS CLI

The AWS Command Line Interface (AWS CLI) is a unified developer tool for managing AWS services, including Amazon Kendra. We recommend that you install it.

1. To install the AWS CLI, follow the instructions in [Installing the AWS Command Line Interface](#) in the *AWS Command Line Interface User Guide*.

2. To configure the AWS CLI and set up a profile to call the AWS CLI, follow the instructions in [Configuring the AWS CLI](#) in the *AWS Command Line Interface User Guide*.
3. To confirm that the AWS CLI profile is configured properly, run the following command:

```
aws configure --profile default
```

If your profile has been configured correctly, you will see output similar to the following:

```
AWS Access Key ID [*****52FQ]:  
AWS Secret Access Key [*****xgyZ]:  
Default region name [us-west-2]:  
Default output format [json]:
```

4. To verify that the AWS CLI is configured for use with Amazon Kendra, run the following commands:

```
aws kendra help
```

If the AWS CLI is configured correctly, you will see a list of the supported AWS CLI commands for Amazon Kendra, Amazon Kendra runtime, and Amazon Kendra events.

Setting up the AWS SDKs

Download and install the AWS SDKs that you want to use. This guide provides examples for Python. For information about other AWS SDKs, see [Tools for Amazon Web Services](#).

Getting started

This section shows how to create a data source and add documents, create an index, get search results, and get code samples to help you integrate Kendra into your application. Instructions are provided for the AWS console, the AWS CLI, a Python program using the AWS SDK for Python (Boto 3), and a Java program using the AWS SDK for Java.

Topics

- [Prerequisites \(p. 11\)](#)
- [Getting started with an S3 bucket \(Console\) \(p. 14\)](#)
- [Getting started \(AWS CLI\) \(p. 15\)](#)
- [Getting started \(AWS SDK for Python \(Boto 3\)\) \(p. 16\)](#)
- [Getting started \(AWS SDK for Java\) \(p. 18\)](#)
- [Getting started with a Microsoft OneDrive for Business data source \(Console\) \(p. 21\)](#)
- [Getting started with an Amazon S3 data source \(Console\) \(p. 22\)](#)
- [Getting started with a Salesforce data source \(Console\) \(p. 23\)](#)
- [Getting started with a ServiceNow data source \(Console\) \(p. 24\)](#)
- [Getting started with a Microsoft SharePoint Online data source \(Console\) \(p. 25\)](#)

Prerequisites

The following steps are prerequisites for the getting started exercises. The steps show you how to set up your account, create an IAM role that gives Amazon Kendra permission to make calls on your behalf, and upload documents for indexing into an Amazon S3 bucket.

1. Create an AWS account and an AWS Identity and Access Management user, as specified in [Sign up for AWS \(p. 9\)](#).
2. If you are using an S3 bucket containing documents to test Amazon Kendra, create an S3 bucket in the same region that you are using Amazon Kendra. For instructions, see [Creating and Configuring an S3 Bucket](#) in the *Amazon Simple Storage Service Console User Guide*.

Upload your documents to your S3 bucket. For instructions, see [Uploading, Downloading, and Managing Objects](#) in the *Amazon Simple Storage Service Console User Guide*.

If you are using another data source, you must have an active site and credentials to connect to the data source.

If you are using the console to get started, do [Getting started with an S3 bucket \(Console\) \(p. 14\)](#).

Prerequisites for the AWS CLI and SDK

If you are using the AWS CLI or the SDK, you need to create IAM roles and policies for Kendra to use to access resources.

To create an IAM policy and role that enables Kendra to access your Amazon CloudWatch Logs.

1. Sign in to the AWS Management Console and open the IAM console at <https://console.aws.amazon.com/iam/>.
2. From the left menu, choose **Policies** and then choose **Create policy**.

3. Choose **JSON** and then replace the default policy with the following:

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": [
        "cloudwatch:PutMetricData"
      ],
      "Resource": "*",
      "Condition": {
        "StringEquals": {
          "cloudwatch:namespace": "AWS/Kendra"
        }
      }
    },
    {
      "Effect": "Allow",
      "Action": [
        "logs:DescribeLogGroups"
      ],
      "Resource": "*"
    },
    {
      "Effect": "Allow",
      "Action": [
        "logs:CreateLogGroup"
      ],
      "Resource": [
        "arn:aws:logs:region:account ID:log-group:/aws/kendra/*"
      ]
    },
    {
      "Effect": "Allow",
      "Action": [
        "logs:DescribeLogStreams",
        "logs:CreateLogStream",
        "logs:PutLogEvents"
      ],
      "Resource": [
        "arn:aws:logs:region:account ID:log-group:/aws/kendra/*:log-stream:*"
      ]
    }
  ]
}
```

4. Choose **Review policy**.
5. Name the policy "KendraPolicyForGettingStartedIndex" and then choose **Create policy**.
6. From the left menu, choose **Roles** and then choose **Create role**.
7. Choose **Another AWS account** and then type your account ID in **Account ID**. Choose **Next: Permissions**.
8. Choose the policy that you created above and then choose **Next: Tags**.
9. Don't add any tags. Choose **Next: Review**.
10. Name the role "KendraRoleForGettingStartedIndex" and then choose **Create role**.
11. Find the role that you just created. Choose the role name to open the summary. Choose **Trust relationships** and then choose **Edit trust relationship**.
12. Replace the existing trust relationship with the following:

```
{
```

```
"Version": "2012-10-17",
"Statement": [
  {
    "Effect": "Allow",
    "Principal": {
      "Service": "kendra.amazonaws.com"
    },
    "Action": "sts:AssumeRole"
  }
]
```

13. Choose **Update trust policy**.

To create an IAM policy and role that enables Kendra to access and index your Amazon S3 bucket.

1. Sign in to the AWS Management Console and open the IAM console at <https://console.aws.amazon.com/iam/>.
2. From the left menu, choose **Policies** and then choose **Create policy**.
3. Choose **JSON** and then replace the default policy with the following:

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Action": [
        "s3:GetObject"
      ],
      "Resource": [
        "arn:aws:s3:::bucket name/*"
      ],
      "Effect": "Allow"
    },
    {
      "Action": [
        "s3:ListBucket"
      ],
      "Resource": [
        "arn:aws:s3:::bucket name"
      ],
      "Effect": "Allow"
    },
    {
      "Effect": "Allow",
      "Action": [
        "kendra:BatchPutDocument",
        "kendra:BatchDeleteDocument"
      ],
      "Resource": "arn:aws:kendra:region:account ID:index/*"
    }
  ]
}
```

4. Choose **Review policy**.
5. Name the policy "KendraPolicyForGettingStartedDataSource" and then choose **Create policy**.
6. From the left menu, choose **Roles** and then choose **Create role**.
7. Choose **Another AWS account** and then type your account ID in **Account ID**. Choose **Next: Permissions**.
8. Choose the policy that you created above and then choose **Next: Tags**

9. Don't add any tags. Choose **Next: Review**.
10. Name the role "KendraRoleForGettingStartedDataSource" and then choose **Create role**.
11. Find the role that you just created. Choose the role name to open the summary. Choose **Trust relationships** and then choose **Edit trust relationship**.
12. Replace the existing trust relationship with the following:

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Principal": {
        "Service": "kendra.amazonaws.com"
      },
      "Action": "sts:AssumeRole"
    }
  ]
}
```

13. Choose **Update trust policy**.

Depending on how you want to use the Amazon Kendra API, do one of the following.

- [Getting started \(AWS CLI\) \(p. 15\)](#)
- [Getting started \(AWS SDK for Java\) \(p. 18\)](#)
- [Getting started \(AWS SDK for Python \(Boto 3\)\) \(p. 16\)](#)

Getting started with an S3 bucket (Console)

The following procedures show how to create and test an Amazon Kendra index by using the AWS console. In the procedures you create an index and a data source for an S3 bucket. Finally, you test your index by making a search request.

Step 1: To create an index (Console)

1. Choose **Create index** to start creating a new index.
2. In **Specify index details**, give your index a name and a description.
3. In **IAM role** choose **Create a new role** and then give the role a name. The IAM role will have the prefix "AmazonKendra-".
4. Leave all of the other fields at their defaults. Choose **Next**.
5. **Provisioning details** page, choose **Developer edition**.
6. Choose **Create** to create your index.
7. Wait for your index to be created. Kendra provisions the hardware for your index. This operation can take some time.

Step 2: To add a data source to an index (Console)

- Create a data source that connects the Amazon Kendra index to your documents. You can choose from one of the following procedures to create a data source.
 - [Getting started with a Microsoft OneDrive for Business data source \(Console\) \(p. 21\)](#)
 - [Getting started with an Amazon S3 data source \(Console\) \(p. 22\)](#)
 - [Getting started with a Salesforce data source \(Console\) \(p. 23\)](#)

- [Getting started with a ServiceNow data source \(Console\) \(p. 24\)](#)
- [Getting started with a Microsoft SharePoint Online data source \(Console\) \(p. 25\)](#)

Step 3: To search an index (Console)

1. In the navigation pane, choose **Search console**
2. Enter a search term that's appropriate for your index. The **top results** and **top document** results are shown.

Getting started (AWS CLI)

The following procedure shows how to create an Amazon Kendra index using the AWS CLI. The procedure creates a data source, index, and runs a query on the index.

To create an Amazon Kendra index (CLI)

1. Do the [Prerequisites \(p. 11\)](#).
2. Enter the following command to create an index

```
aws kendra create-index \  
  --name cli-getting-started-index \  
  --description "Index for CLI getting started guide." \  
  --role-arn arn:aws:iam::account id:role/KendraRoleForGettingStartedIndex
```

3. Wait for Amazon Kendra to create the index. Check the progress using the following command. When the status field is **ACTIVE**, go on to the next step.

```
aws kendra describe-index \  
  --id index id
```

4. At the command prompt, enter the following command to create a data source.

```
aws kendra create-data-source \  
  --index-id index id \  
  --name data source name \  
  --role-arn arn:aws:iam::account id:role/KendraRoleForGettingStartedDataSource \  
  --type S3 \  
  --configuration '{"S3Configuration":{"BucketName":"S3 bucket name"}}'
```

5. It will take Amazon Kendra a while to create the data source. Enter the following command to check the progress. When the status is **ACTIVE**, go on to the next step.

```
aws kendra describe-data-source \  
  --id data source ID \  
  --index-id index ID
```

6. Enter the following command to synchronize the data source.

```
aws kendra start-data-source-sync-job \  
  --id data source ID \  
  --index-id index ID
```

7. Kendra will index your data source. The amount of time that it takes depends on the number of documents. You can check the status of the sync job using the following command. When the status is **ACTIVE**, go on to the next step.

```
aws kendra describe-data-source \  
  --id data source ID \  
  --index-id index ID
```

8. Enter the following command to make a query.

```
aws kendra query \  
  --index-id index ID \  
  --query-text "search term"
```

The results of the search are displayed in JSON format.

Getting started (AWS SDK for Python (Boto 3))

The following program is an example of using Amazon Kendra in a Python program. The program performs the following actions:

1. Creates a new index using the [CreateIndex \(p. 142\)](#) operation.
2. Waits for index creation to complete. It uses the [DescribeIndex \(p. 162\)](#) operation to monitor the status of the index.
3. Once the index is active, it creates a data source using the [CreateDataSource \(p. 132\)](#) operation.
4. Waits for data source creation to complete. It uses the [DescribeDataSource \(p. 151\)](#) operation to monitor the status of the data source.
5. When the data source is active, it synchronizes the index with the contents of the data source using the [StartDataSourceSyncJob \(p. 185\)](#) operation.

```
import boto3  
from botocore.exceptions import ClientError  
import pprint  
import time  
  
kendra = boto3.client("kendra")  
  
print("Create an index")  
  
description = "Getting started index"  
index_name = "python-getting-started-index"  
index_role_arn = "arn:aws:iam::${accountId}:role/KendraRoleForGettingStartedIndex"  
  
try:  
    index_response = kendra.create_index(  
        Description = description,  
        Name = index_name,  
        RoleArn = index_role_arn  
    )  
  
    pprint.pprint(index_response)  
  
    index_id = index_response["Id"]  
  
    print("Wait for Kendra to create the index.")  
  
    while True:  
        # Get index description  
        index_description = kendra.describe_index(  
            Id = index_id
```



```
)
# When status is not CREATING quit.
status = index_description["Status"]
print("    Creating index. Status: "+status)
time.sleep(60)
if status != "CREATING":
    break

print("Create an S3 data source")

data_source_name = "python-getting-started-data-source"
data_source_description = "Getting started data source."
s3_bucket_name = "${bucketName}"
data_source_type = "S3"
data_source_role_arn = "arn:aws:iam::${accountId}:role/
KendraRoleForGettingStartedDataSource"

configuration = {"S3Configuration":
    {
        "BucketName": s3_bucket_name
    }
}

data_source_response=kendra.create_data_source(
    Configuration = configuration,
    Name = data_source_name,
    Description = description,
    RoleArn = data_source_role_arn,
    Type = data_source_type,

    IndexId = index_id
)

pprint.pprint(data_source_response)

data_source_id = data_source_response["Id"]

print("Wait for Kendra to create the data source.")

while True:
    data_source_description = kendra.describe_data_source(
        Id = data_source_id,
        IndexId = index_id
    )
    # When status is not CREATING quit.
    status = data_source_description["Status"]
    print("    Creating data source. Status: "+status)
    time.sleep(60)
    if status != "CREATING":
        break

print("Synchronize the data source.")

sync_response = kendra.start_data_source_sync_job(
    Id = data_source_id,
    IndexId = index_id
)

pprint.pprint(sync_response)

print("Wait for the data source to sync with the index.")

while True:

    jobs = kendra.list_data_source_sync_jobs(
        Id=data_source_id,
```

```
        IndexId=index_id
    )

    # There should be exactly one job item in response
    status = jobs["History"][0]["Status"]

    print("    Syncing data source. Status: "+status)
    if status != "SYNCING":
        break
    time.sleep(60)

except ClientError as e:
    print("%s" % e)

print("Program ends.")
```

Getting started (AWS SDK for Java)

The following program is an example of using Amazon Kendra in a Python program. The program performs the following actions:

1. Creates a new index using the [CreateIndex \(p. 142\)](#) operation.
2. Waits for index creation to complete. It uses the [DescribeIndex \(p. 162\)](#) operation to monitor the status of the index.
3. Once the index is active, it creates a data source using the [CreateDataSource \(p. 132\)](#) operation.
4. Waits for data source creation to complete. It uses the [DescribeDataSource \(p. 151\)](#) operation to monitor the status of the data source.
5. When the data source is active, it synchronizes the index with the contents of the data source using the [StartDataSourceSyncJob \(p. 185\)](#) operation.

```
package com.amazonaws.kendra;

import java.util.concurrent.TimeUnit;
import software.amazon.awssdk.services.kendra.KendraClient;
import software.amazon.awssdk.services.kendra.model.CreateDataSourceRequest;
import software.amazon.awssdk.services.kendra.model.CreateDataSourceResponse;
import software.amazon.awssdk.services.kendra.model.CreateIndexRequest;
import software.amazon.awssdk.services.kendra.model.CreateIndexResponse;
import software.amazon.awssdk.services.kendra.model.DataSourceConfiguration;
import software.amazon.awssdk.services.kendra.model.DataSourceStatus;
import software.amazon.awssdk.services.kendra.model.DataSourceSyncJob;
import software.amazon.awssdk.services.kendra.model.DataSourceSyncJobStatus;
import software.amazon.awssdk.services.kendra.model.DataSourceType;
import software.amazon.awssdk.services.kendra.model.DescribeDataSourceRequest;
import software.amazon.awssdk.services.kendra.model.DescribeDataSourceResponse;
import software.amazon.awssdk.services.kendra.model.DescribeIndexRequest;
import software.amazon.awssdk.services.kendra.model.DescribeIndexResponse;
import software.amazon.awssdk.services.kendra.model.IndexStatus;
import software.amazon.awssdk.services.kendra.model.ListDataSourceSyncJobsRequest;
import software.amazon.awssdk.services.kendra.model.ListDataSourceSyncJobsResponse;
import software.amazon.awssdk.services.kendra.model.S3DataSourceConfiguration;
import software.amazon.awssdk.services.kendra.model.StartDataSourceSyncJobRequest;
import software.amazon.awssdk.services.kendra.model.StartDataSourceSyncJobResponse;

public class CreateIndexAndDataSourceExample {

    public static void main(String[] args) throws InterruptedException {
```

```
System.out.println("Create an index");

String indexDescription = "Getting started index for Kendra";
String indexName = "java-getting-started-index";
String indexRoleArn = "arn:aws:iam::<your AWS account ID>:role/<name of an IAM
role>";

System.out.println(String.format("Creating an index named %s", indexName));
KendraClient kendra = KendraClient.builder().build();

CreateIndexRequest createIndexRequest = CreateIndexRequest
    .builder()
    .description(indexDescription)
    .name(indexName)
    .roleArn(indexRoleArn)
    .build();
CreateIndexResponse createIndexResponse = kendra.createIndex(createIndexRequest);
System.out.println(String.format("Index response %s", createIndexResponse));

String indexId = createIndexResponse.id();

System.out.println(String.format("Waiting until the index with index ID %s is
created", indexId));
while (true) {
    DescribeIndexRequest describeIndexRequest =
DescribeIndexRequest.builder().id(indexId).build();
    DescribeIndexResponse describeIndexResponse =
kendra.describeIndex(describeIndexRequest);
    IndexStatus status = describeIndexResponse.status();
    if (status != IndexStatus.CREATING) {
        break;
    }

    TimeUnit.SECONDS.sleep(60);
}

System.out.println("Creating an S3 data source");
String dataSourceName = "java-getting-started-data-source";
String dataSourceDescription = "Getting started data source";
String s3BucketName = "an-aws-kendra-test-bucket";
String dataSourceRoleArn = "arn:aws:iam::<your aws account ID>:role/<name of an IAM
role>";

CreateDataSourceRequest createDataSourceRequest = CreateDataSourceRequest
    .builder()
    .indexId(indexId)
    .name(dataSourceName)
    .description(dataSourceDescription)
    .roleArn(dataSourceRoleArn)
    .type(DataSourceType.S3)
    .configuration(
        DataSourceConfiguration
            .builder()
            .s3Configuration(
                S3DataSourceConfiguration
                    .builder()
                    .bucketName(s3BucketName)
                    .build()
            ).build()
    ).build();

CreateDataSourceResponse createDataSourceResponse =
kendra.createDataSource(createDataSourceRequest);
System.out.println(String.format("Response of creating data source: %s",
createDataSourceResponse));
```

```
String dataSourceId = createDataSourceResponse.id();
System.out.println(String.format("Waiting for Kendra to create the data source %s",
dataSourceId));
DescribeDataSourceRequest describeDataSourceRequest = DescribeDataSourceRequest
    .builder()
    .indexId(indexId)
    .id(dataSourceId)
    .build();

while (true) {
    DescribeDataSourceResponse describeDataSourceResponse =
kendra.describeDataSource(describeDataSourceRequest);

    DataSourceStatus status = describeDataSourceResponse.status();
    System.out.println(String.format("Creating data source. Status: %s", status));
    if (status != DataSourceStatus.CREATING) {
        break;
    }

    TimeUnit.SECONDS.sleep(60);
}

System.out.println(String.format("Synchronize the data source %s", dataSourceId));
StartDataSourceSyncJobRequest startDataSourceSyncJobRequest =
StartDataSourceSyncJobRequest
    .builder()
    .indexId(indexId)
    .id(dataSourceId)
    .build();
StartDataSourceSyncJobResponse startDataSourceSyncJobResponse =
kendra.startDataSourceSyncJob(startDataSourceSyncJobRequest);
System.out.println(String.format("Waiting for the data source to sync with the
index %s for execution ID %s", indexId, startDataSourceSyncJobResponse.executionId()));

// For this particular list, there should be just one job
ListDataSourceSyncJobsRequest listDataSourceSyncJobsRequest =
ListDataSourceSyncJobsRequest
    .builder()
    .indexId(indexId)
    .id(dataSourceId)
    .build();

while (true) {
    ListDataSourceSyncJobsResponse listDataSourceSyncJobsResponse =
kendra.listDataSourceSyncJobs(listDataSourceSyncJobsRequest);
    DataSourceSyncJob job = listDataSourceSyncJobsResponse.history().get(0);
    System.out.println(String.format("Syncing data source. Status: %s",
job.status()));

    if (job.status() != DataSourceSyncJobStatus.SYNCING) {
        break;
    }

    TimeUnit.SECONDS.sleep(60);
}

System.out.println("Index setup is complete");
}
}
```

Getting started with a Microsoft OneDrive for Business data source (Console)

You can use the Amazon Kendra console to get started indexing a Microsoft OneDrive for Business site. When you use the console you can specify all of the connection information that you need to index the contents of a OneDrive site. For more information, see [Using a Microsoft OneDrive data source \(p. 52\)](#).

Use the following procedure to create a basic OneDrive data source using the default configuration. The procedure assumes that you have already created an index following the steps in step 1 of [Getting started with an S3 bucket \(Console\) \(p. 14\)](#).

Before you can create a OneDrive data source, you must register an Azure Active Directory application. Use the following procedure to register the application.

Step 1: To register an Azure AD application

1. Log into the [Azure Management Portal](#).
2. Choose **Azure Active Directory** and then choose **App registrations**.
3. Choose **New registration**.
4. Create the application with the following values.
 - **Name** Enter the name for your application.
 - **Application type** **Web app/API**
 - **Sign-on URL** Any valid URL. The URL doesn't need to exist.
5. Choose **Register**.
6. Choose the application name in the list of applications to open the application settings.
7. Choose **Overview** and then copy the **Application ID** value.

The application requires permission to access OneDrive data. Use the following procedure to assign permissions to the application. This procedure starts where the previous procedure left off.

Step 2: To assign permissions to the Azure AD application

1. From the application settings, choose **API Permission**.
2. Choose **Add**, and choose the **Microsoft Graph** option, and then choose **Add**.
3. From the **Application permissions**, choose the following.
 - Read files in all site collections (Files.Read.All)
 - Read all users' full profiles (User.Read.All)
 - Read directory data (Directory.Read.All)
 - Read all groups (Group.Read.All)
 - Read items in all site collections (Sites.Read.All)
4. Choose **Save** and then choose **Grant permissions**. Choose **Yes** when prompted.

Next you create an application key that Amazon Kendra uses to identify itself with the OneDrive site. This procedure starts where the previous procedure left off.

Step 3: To create an application key

1. From the **Application settings**, choose **Keys**.
2. Add the following information to the key.

- **Description** Add a description for your key.
 - **Expires** Choose a duration based on your company policy. When the application key is rotated, you must update the key in the Amazon Kendra data source.
3. Choose **Save** to generate the key value. Copy this value to use in the next step.

Once you have created the Azure AD application and generated a secret key for the application, you are ready to create a OneDrive data source.

Step 4: To create a OneDrive data source using the Amazon Kendra console

1. Sign into the AWS Management Console and open the Amazon Kendra console at <https://console.aws.amazon.com/kendra/home>.
2. From the list of indexes, choose the index that you want to add the data source to.
3. Choose **Add data sources**.
4. From the list of data source connectors, choose **OneDrive**.
5. On the **Define attributes** page, give your data source a name and optionally a description. Leave the **Tags** field blank. Choose **Next** to continue.
6. On the **Define targets** page, enter the domain name of your OneDrive site. Enter the name without the protocol ("https://").
7. In the **IAM role** field choose **Create a new role**. Enter a name for the role in the **Role name** field. Choose **Next** to continue.
8. In the **Type of authentication** field, choose **New**. This tells the console to create a new AWS Secrets Manager secret to contain credentials for your OneDrive data source.
9. In the **New secret container name** field enter a name to identify the Secrets Manager secret.
10. In the **Application ID** and **Application password** fields enter the ID and password for the Active Directory application that you created for the data source. Choose **Save authentication** to save the credentials for your OneDrive site.
11. In the **Add OneDrive users** section, choose **Names list**. Add the email addresses of up to 10 users whose documents you want to index.
12. In the **Set sync run schedule** section, choose **Run on demand** in the **Frequency** field.
13. In the **OneDrive field mapping** section, leave the default fields selected and then choose **Next**.
14. On the **Review and create** page review the details of your OneDrive data source. If you want to make changes, choose the **Edit** button next to the item that you want to change. When you are satisfied with your choices, choose **Create** to create your OneDrive data source.

After you choose **Create**, Amazon Kendra starts creating the data source. It can take several minutes for the data source to be created. When it is finished, the status of the data source changes from **Creating** to **Active**.

After creating the data source, you need to sync the Amazon Kendra index with the data source. Choose **Sync now** to start the sync process. It can take several minutes to several hours to synchronize the data source, depending on the number and size of the documents.

Getting started with an Amazon S3 data source (Console)

You can use the Amazon Kendra console to get started using an Amazon S3 bucket as a data store. When you use the console you specify all of the connection information you need to index the contents of the bucket. For more information, see [Using an Amazon S3 data source \(p. 47\)](#).

Use the following procedure to create a basic S3 bucket data source using the default configuration. The procedure assumes that you created an index following the steps in step 1 of [Getting started with an S3 bucket \(Console\) \(p. 14\)](#).

To create an S3 bucket data source using the Amazon Kendra console

1. Sign into the AWS Management Console and open the Amazon Kendra console at <https://console.aws.amazon.com/kendra/home>.
2. From the list of indexes, choose the index that you want to add the data source to.
3. Choose **Add data sources**.
4. From the list of data source connectors, choose **Amazon S3**.
5. On the **Define attributes** page, give your data source a name and optionally a description. Leave the **Tags** field blank. Choose **Next** to continue.
6. In the **Enter the data source location** field, enter the name of the S3 bucket that contains your documents. You can enter the name directly, or you can browse for the name by choosing **Browse**. The bucket must be in the same region as the index.
7. In **IAM role** choose **Create a new role** and then type a role name.
8. In the **Set sync run schedule** section, choose **Run on demand**.
9. Choose **Next** to continue.
10. On the **Review and create** page review the details of your S3 data source. If you want to make changes, choose the **Edit** button next to the item that you want to change. When you are satisfied with your choices, choose **Create** to create your S3 data source.

After you choose **Create**, Amazon Kendra starts creating the data source. It can take several minutes for the data source to be created. When it is finished, the status of the data source changes from **Creating** to **Active**.

After creating the data source, you need to sync the Amazon Kendra index with the data source. Choose **Sync now** to start the sync process. It can take several minutes to several hours to synchronize the data source, depending on the number and size of the documents.

Getting started with a Salesforce data source (Console)

You can use the Amazon Kendra console to get started using a Salesforce data store. When you use the console you specify the connection information you need to index the contents of a Salesforce instance. For more information, see [Using a Salesforce data source \(p. 54\)](#).

Use the following procedure to create a basic Salesforce data source using the default configuration. The procedure assumes you created an index following the steps in step 1 of [Getting started with an S3 bucket \(Console\) \(p. 14\)](#).

To create a Salesforce data source using the Amazon Kendra console

1. Sign into the AWS Management Console and open the Amazon Kendra console at <https://console.aws.amazon.com/kendra/home>.
2. From the list of indexes, choose the index that you want to add the data source to.
3. Choose **Add data sources**.
4. From the list of data source connectors, choose **Salesforce**.
5. On the **Define attributes** page, give your data source a name and optionally a description. Leave the **Tags** field blank. Choose **Next** to continue.

6. On the **Define targets** page, enter the URL of your Salesforce server.
7. In the **IAM role** field choose **Create a new role**. Enter a name for the role in the **Role name** field. Choose **Next** to continue.
8. In the **Type of authentication** field, choose **New**. This tells the console to create a new AWS Secrets Manager secret to contain credentials for your Salesforce data source.
9. In the **New secret container name** field enter a name to identify the Secrets Manager secret.
10. In the **Username** field, enter the user name for your Salesforce account.
11. In the **Password** field, enter the password for your Salesforce account.
12. In the **Security token** field, enter the security token for your Salesforce account.
13. In the **Consumer key** field, enter the consumer key of the Salesforce connected app that you are using.
14. In the **Consumer secret** field, enter the secret associated with the consumer key of the Salesforce connected app that you are using.
15. In the **Authentication URL** field, enter the OAUTH endpoint for Salesforce.
16. Choose **Save authentication** to save your authentication information.
17. In the **Crawl settings** section, under **Standard objects**, choose **Document**.
18. In the **Set sync run schedule** section, choose **Run on demand**.
19. Choose **Next** to continue.
20. On the **Set field mappings - optional** page, leave the defaults and choose **Next**.
21. On the **Review and create** page review the settings for the data source. Use the **Edit** buttons to make any changes that you need to make. When you are satisfied with the settings, choose **Create** to create your Salesforce data source.

After you choose **Create**, Amazon Kendra starts creating the data source. It can take several minutes for the data source to be created. When it is finished, the status changes from **Creating** to **Active**.

After creating the index, you need to sync the Amazon Kendra index with the data source. Choose **Sync now** to start the sync process. It can take several minutes to several hours to synchronize the data source, depending on the number and size of the documents.

Getting started with a ServiceNow data source (Console)

You can use the Amazon Kendra console to get started using a ServiceNow data store. When you use the console you specify the connection information you need to index the contents of a ServiceNow instance. For more information see [Using a ServiceNow data source \(p. 57\)](#).

Use the following procedure to create a basic ServiceNow data source using the default configuration. The procedure assumes you created an index following the steps in [Getting started with an S3 bucket \(Console\) \(p. 14\)](#).

To create a ServiceNow data source using the Amazon Kendra console

1. Sign into the AWS Management Console and open the Amazon Kendra console at <https://console.aws.amazon.com/kendra/home>.
2. From the list of indexes, choose the index that you want to add the data source to.
3. Choose **Add data sources**.
4. From the list of data source connectors, choose **ServiceNow**.

5. On the **Define attributes** page, give your data source a name and optionally a description. Leave the **Tags** field blank. Choose **Next** to continue.
6. For the **ServiceNow host**, enter the URL of your ServiceNow instances.
7. From the **ServiceNow version** dropdown, choose the ServiceNow version that your organization uses. You can choose **London** or all other versions.
8. In the **IAM role** field choose **Create a new role**. Enter a name for the role in the **Role name** field. Choose **Next** to continue.
9. In the **Type of authentication** field, choose **New**. This tells the console to create a new AWS Secrets Manager secret to contain credentials for your ServiceNow data source.
10. In the **New secret container name** field enter a name to identify the Secrets Manager secret.
11. Enter the username and password for your ServiceNow account, and then choose **Save authentication** to save the new secret.
12. In the **ServiceNow configuration** section leave the defaults selected.
13. In the **Set sync run schedule** section, choose **Run on demand** in the **Frequency** field.
14. Choose **Next** to continue.
15. On the **Set field mappings - optional** page, leave the defaults checked and choose **Next** to continue.
16. On the **Review and create** page review the details of your ServiceNow data source. If you want to make changes, choose the **Edit** button next to the item that you want to change. When you are satisfied with your choices, choose **Create** to create your ServiceNow data source.

After you choose **Create**, Amazon Kendra starts creating the data source. It can take several minutes for the data source to be created. When it is finished, the status of the data source changes from **Creating** to **Active**.

After creating the data source, you need to sync the Amazon Kendra index with the data source. Choose **Sync now** to start the sync process. It can take several minutes to several hours to synchronize the data source, depending on the number and size of the documents.

Getting started with a Microsoft SharePoint Online data source (Console)

You can use the Amazon Kendra console to get started indexing a Microsoft SharePoint Online site. When you use the console you specify the connection information you need to index the contents of the SharePoint Online site.

Use the following procedure to create a basic SharePoint Online data source using the default configuration. The procedure assumes you created an index following the steps in step 1 of [Getting started with an S3 bucket \(Console\)](#) (p. 14).

To create a SharePoint Online data source using the Amazon Kendra console

1. Sign into the AWS Management Console and open the Amazon Kendra console at <https://console.aws.amazon.com/kendra/home>.
2. From the list of indexes, choose the index that you want to add the data source to.
3. Choose **Add data sources**.
4. From the list of data source connectors, choose **SharePoint Online**.
5. On the **Define attributes** page, give your data source a name and optionally a description. Leave the **Tags** field blank. Choose **Next** to continue.
6. In the **Add SharePoint URLs** section, add the URLs of the SharePoint Online site to index.

7. In **IAM role** choose **Create a new role** and then type a role name.
8. In the **Type of authentication** field, choose **New**. This tells the console to create a new AWS Secrets Manager secret to contain credentials for your SharePoint Online data source.
9. In the **New secret container name** field enter a name to identify the Secrets Manager secret.
10. Enter the username and password for your SharePoint Online account, and then choose **Save authentication** to save the new secret.
11. In the **Set sync run schedule** section, choose **Run on demand** in the **Frequency** field.
12. Choose **Next** to continue.
13. On the **Set field mappings - optional** page, leave the defaults checked and choose **Next** to continue.
14. On the **Review and create** page review the details of your SharePoint Online data source. If you want to make changes, choose the **Edit** button next to the item that you want to change. When you are satisfied with your choices, choose **Create** to create your SharePoint Online data source.

After you choose **Create**, Amazon Kendra starts creating the data source. It can take several minutes for the data source to be created. When it is finished, the status of the data source changes from **Creating** to **Active**.

After creating the data source, you need to sync the Amazon Kendra index with the data source. Choose **Sync now** to start the sync process. It can take several minutes to several hours to synchronize the data source, depending on the number and size of the documents.

IAM access roles for Amazon Kendra

When you create an index, data source, or an FAQ, Amazon Kendra needs access to the AWS resources required to create the Amazon Kendra resource. You must create a AWS Identity and Access Management (IAM) policy before you create the Amazon Kendra resource. When you call the operation, you provide the Amazon Resource Name (ARN) of the role with the policy attached. For example, if you are calling the [BatchPutDocument](#) (p. 129) operation to add documents from an Amazon S3 bucket, you provide Amazon Kendra with a role with a policy that has access to the bucket.

The Amazon Kendra console enables you to create a new IAM role or to choose an IAM existing role to use. The console displays roles that have the string "kendra" or "Kendra" in the role name.

The following topics provide details for the required policies. If you create IAM roles using the Amazon Kendra console these policies are created for you.

Topics

- [IAM roles for indexes](#) (p. 27)
- [IAM Roles for the BatchPutDocument operation](#) (p. 28)
- [IAM roles for data sources](#) (p. 29)
- [IAM roles for frequently asked questions](#) (p. 37)

IAM roles for indexes

When you create an index, you must provide an IAM role with permission to write to an Amazon CloudWatch Logs. You must also provide a trust policy that allows Amazon Kendra to assume the role. The following are the policies that must be provided.

A role policy to enable Amazon Kendra to access a CloudWatch log.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": "cloudwatch:PutMetricData",
      "Resource": "*",
      "Condition": {
        "StringEquals": {
          "cloudwatch:namespace": "Kendra"
        }
      }
    },
    {
      "Effect": "Allow",
      "Action": "logs:DescribeLogGroups",
      "Resource": "*"
    },
    {
      "Effect": "Allow",
      "Action": "logs:CreateLogGroup",
      "Resource": "arn:aws:logs:region:account ID:log-group:/aws/kendra/*"
    }
  ]
}
```

```
    },
    {
      "Effect": "Allow",
      "Action": [
        "logs:DescribeLogStreams",
        "logs:CreateLogStream",
        "logs:PutLogEvents"
      ],
      "Resource": "arn:aws:logs:region:account ID:log-group:/aws/kendra/*:log-
stream:*"
    }
  ]
}
```

A trust policy to enable Amazon Kendra to assume a role.

```
{
  "Version": "2012-10-17",
  "Statement": {
    "Effect": "Allow",
    "Principal": {
      "Service": "kendra.amazonaws.com"
    },
    "Action": "sts:AssumeRole"
  }
}
```

IAM Roles for the BatchPutDocument operation

When you use the [BatchPutDocument](#) (p. 129) operation to index documents in an Amazon S3 bucket, you must provide Amazon Kendra with an IAM role with access to the bucket. You must also provide a trust policy that enables Amazon Kendra to assume the role. If the documents in the bucket are encrypted, you must provide permission to use the AWS KMS customer master key (CMK) to decrypt the documents.

A required role policy to enable Amazon Kendra to access an Amazon S3 bucket.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": [
        "s3:GetObject"
      ],
      "Resource": [
        "arn:aws:s3:::bucket name/*"
      ]
    }
  ]
}
```

A required trust policy to enable Amazon Kendra to assume a role.

```
{
  "Version": "2012-10-17",
  "Statement": {
    "Sid": "AllowKendraToAssumeAttachedRole",

```

```
    "Effect": "Allow",
    "Principal": {
      "Service": "kendra.amazonaws.com"
    },
    "Action": "sts:AssumeRole"
  }
}
```

An optional role policy to enable Amazon Kendra to use an AWS KMS customer master key (CMK) to decrypt documents in an Amazon S3 bucket.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": [
        "kms:Decrypt"
      ],
      "Resource": [
        "arn:aws:kms:region:account ID:key/key ID"
      ]
    }
  ]
}
```

IAM roles for data sources

When you use the [CreateDataSource](#) (p. 132) operation, you must give Amazon Kendra an IAM role that has permission to access the database resources. The specific permissions required depend on the data source.

Topics

- [IAM roles for Amazon S3 data sources](#) (p. 29)
- [IAM roles for database data sources](#) (p. 30)
- [IAM roles for Microsoft OneDrive data sources](#) (p. 32)
- [IAM role for Salesforce data sources](#) (p. 34)
- [IAM role for ServiceNow data sources](#) (p. 35)
- [IAM roles for Microsoft SharePoint Online data sources](#) (p. 36)

IAM roles for Amazon S3 data sources

When you use an Amazon S3 bucket as a data source, you supply a role that has permission to access the bucket, and to use the `BatchPutDocument` and `BatchDeleteDocument` operations. If the documents in the Amazon S3 bucket are encrypted, you must provide permission to use the AWS KMS customer master key (CMK) to decrypt the documents.

A required role policy to enable Amazon Kendra to use an Amazon S3 bucket as a data source.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
```

```
        "Action": [
            "s3:GetObject"
        ],
        "Resource": [
            "arn:aws:s3:::bucket name/*"
        ],
        "Effect": "Allow"
    },
    {
        "Action": [
            "s3:ListBucket"
        ],
        "Resource": [
            "arn:aws:s3:::bucket name"
        ],
        "Effect": "Allow"
    },
    {
        "Effect": "Allow",
        "Action": [
            "kendra:BatchPutDocument",
            "kendra:BatchDeleteDocument"
        ],
        "Resource": [
            "arn:aws:kendra:region:account ID:index/index ID"
        ]
    }
]
```

An optional role policy to enable Amazon Kendra to use an AWS KMS customer master key (CMK) to decrypt documents in an Amazon S3 bucket.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": [
        "kms:Decrypt"
      ],
      "Resource": [
        "arn:aws:kms:region:account ID:key/key ID"
      ]
    }
  ]
}
```

IAM roles for database data sources

When you use a database as a data source, you provide Amazon Kendra with a role that has the permissions necessary for connecting to the database. These include:

- Permission to access the AWS Secrets Manager secret that contains the user name and password for the database site. For more information about the contents of the secret, see [Using a database data source \(p. 50\)](#).
- Permission to use the AWS KMS customer master key (CMK) to decrypt the user name and password secret stored by Secrets Manager.
- Permission to use the `BatchPutDocument` and `BatchDeleteDocument` operations to update the index.

- Permission to access the Amazon S3 bucket that contains the SSL certificate used to communicate with the database site.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": [
        "secretsmanager:GetSecretValue"
      ],
      "Resource": [
        "arn:aws:secretsmanager:region:account ID:secret:secret ID"
      ]
    },
    {
      "Effect": "Allow",
      "Action": [
        "kms:Decrypt"
      ],
      "Resource": [
        "arn:aws:kms:region:account ID:key/key ID"
      ]
    },
    {
      "Effect": "Allow",
      "Action": [
        "kendra:BatchPutDocument",
        "kendra:BatchDeleteDocument"
      ],
      "Resource": [
        "arn:aws:kendra:region:account ID:index/index ID"
      ],
      "Condition": {
        "StringLike": {
          "kms:ViaService": [
            "kendra.amazonaws.com"
          ]
        }
      }
    },
    {
      "Effect": "Allow",
      "Action": [
        "s3:GetObject"
      ],
      "Resource": [
        "arn:aws:s3:::bucket name/*"
      ]
    }
  ]
}
```

There are two optional policies that you might use with a database data source.

If you have encrypted the Amazon S3 bucket that contains the SSL certificate used to communicate with the database, provide a policy to give Amazon Kendra access to the key.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": [
```

```
        "kms:Decrypt"
      ],
      "Resource": [
        "arn:aws:kms:region:account ID:key/key ID"
      ]
    }
  ]
}
```

If you are using a VPC, provide a policy that gives Amazon Kendra access to the required resources.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": [
        "ec2:CreateNetworkInterface",
        "ec2:DescribeNetworkInterfaces",
        "ec2>DeleteNetworkInterface"
      ],
      "Resource": "*"
    },
    {
      "Effect": "Allow",
      "Action": [
        "ec2:CreateNetworkInterfacePermission"
      ],
      "Resource": "*",
      "Condition": {
        "StringEquals": {
          "ec2:AuthorizedService": "kendra.amazonaws.com"
        },
        "ArnEquals": {
          "ec2:Subnet": [
            "arn:aws:ec2:region:account ID:subnet/subnet IDs"
          ]
        }
      }
    },
    {
      "Effect": "Allow",
      "Action": [
        "ec2:DescribeSubnets"
      ],
      "Resource": "*"
    }
  ]
}
```

IAM roles for Microsoft OneDrive data sources

When you use a Microsoft OneDrive data source, you provide Amazon Kendra with a role that has the permissions necessary for connecting to the site. These include:

- Permission to get and decrypt the AWS Secrets Manager secret that contains the application ID and secret key necessary to connect to the OneDrive site. For more information about the contents of the secret, see [Using a Microsoft OneDrive data source \(p. 52\)](#).
- Permission to use the [BatchPutDocument \(p. 129\)](#) and [BatchDeleteDocument \(p. 126\)](#) operations.

The following IAM policy provides the necessary permissions:


```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": [
        "secretsmanager:GetSecretValue"
      ],
      "Resource": [
        "arn:aws:secretsmanager:region:account ID:secret:secret ID"
      ]
    },
    {
      "Effect": "Allow",
      "Action": [
        "kms:Decrypt"
      ],
      "Resource": [
        "arn:aws:kms:region:account ID:key/key ID"
      ],
      "Condition": {
        "StringLike": {
          "kms:ViaService": [
            "secretsmanager.*.amazonaws.com"
          ]
        }
      }
    },
    {
      "Effect": "Allow",
      "Action": [
        "kendra:BatchPutDocument",
        "kendra:BatchDeleteDocument"
      ],
      "Resource": "arn:aws:kendra:region:account ID:index/index ID"
    }
  ]
}
```

If you are storing the list of users to index in an S3 bucket, you must also provide permission to use the S3 GetObject operation. The following IAM policy provides the necessary permissions:

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": [
        "secretsmanager:GetSecretValue"
      ],
      "Resource": [
        "arn:aws:secretsmanager:region:account ID:secret:secret ID"
      ]
    },
    {
      "Action": [
        "s3:GetObject"
      ],
      "Resource": [
        "arn:aws:s3:::input_bucket_name/*"
      ],
      "Effect": "Allow"
    },
    {
      "Effect": "Allow",

```

```
"Action": [
  "kms:Decrypt"
],
"Resource": [
  "arn:aws:kms:region:account ID:key/[key IDs]"
],
"Condition": {
  "StringLike": {
    "kms:ViaService": [
      "secretsmanager.*.amazonaws.com",
      "s3.*.amazonaws.com"
    ]
  }
},
{
  "Effect": "Allow",
  "Action": [
    "kendra:BatchPutDocument",
    "kendra:BatchDeleteDocument"
  ],
  "Resource": "arn:aws:kendra:region:account ID:index/index ID"
}]
}
```

IAM role for Salesforce data sources

When you use a Salesforce as a data source, you provide a role with the following policies:

- Permission to access the AWS Secrets Manager secret that contains the user name and password for the Salesforce site. For more information about the contents of the secret, see [Using a Salesforce data source \(p. 54\)](#).
- Permission to use the AWS KMS customer master key (CMK) to decrypt the user name and password secret stored by Secrets Manager.
- Permission to use the BatchPutDocument and BatchDeleteDocument operations to update the index.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": [
        "secretsmanager:GetSecretValue"
      ],
      "Resource": [
        "arn:aws:secretsmanager:region:account ID:secret:secret ID"
      ]
    },
    {
      "Effect": "Allow",
      "Action": [
        "kms:Decrypt"
      ],
      "Resource": [
        "arn:aws:kms:region:account ID:key/key ID"
      ],
      "Condition": {
        "StringLike": {
          "kms:ViaService": [
            "secretsmanager.*.amazonaws.com"
          ]
        }
      }
    }
  ]
}
```

```
    ]
  }
},
{
  "Effect": "Allow",
  "Action": [
    "kendra:BatchPutDocument",
    "kendra:BatchDeleteDocument"
  ],
  "Resource": "arn:aws:kendra:region:account ID:index/index ID"
}]
}
```

IAM role for ServiceNow data sources

When you use a ServiceNow as a data source, you provide a role with the following policies:

- Permission to access the AWS Secrets Manager secret that contains the user name and password for the ServiceNow site. For more information about the contents of the secret, see [Using a ServiceNow data source \(p. 57\)](#).
- Permission to use the AWS KMS customer master key (CMK) to decrypt the user name and password secret stored by Secrets Manager.
- Permission to use the BatchPutDocument and BatchDeleteDocument operations to update the index.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": [
        "secretsmanager:GetSecretValue"
      ],
      "Resource": [
        "arn:aws:secretsmanager:region:account ID:secret:secret ID"
      ]
    },
    {
      "Effect": "Allow",
      "Action": [
        "kms:Decrypt"
      ],
      "Resource": [
        "arn:aws:kms:region:account ID:key/key ID"
      ],
      "Condition": {
        "StringLike": {
          "kms:ViaService": [
            "secretsmanager.*.amazonaws.com"
          ]
        }
      }
    }
  ],
  {
    "Effect": "Allow",
    "Action": [
      "kendra:BatchPutDocument",
      "kendra:BatchDeleteDocument"
    ],
    "Resource": "arn:aws:kendra:region:account ID:index/index ID"
  }
}
```

```
}]  
}
```

IAM roles for Microsoft SharePoint Online data sources

For a Microsoft SharePoint Online data source, you provide a role with the following policies.

- Permission to access the AWS Secrets Manager secret that contains the user name and password for the SharePoint site. For more information about the contents of the secret, see [Using a Microsoft SharePoint data source \(p. 60\)](#).
- Permission to use the AWS KMS customer master key (CMK) to decrypt the user name and password secret stored by Secrets Manager.
- Permission to use the `BatchPutDocument` and `BatchDeleteDocument` operations to update the index.
- Permission to access the Amazon S3 bucket that contains the SSL certificate used to communicate with the SharePoint site.

You must also attach a trust policy that enables Amazon Kendra to assume the role.

```
{  
  "Version": "2012-10-17",  
  "Statement": [  
    {  
      "Effect": "Allow",  
      "Action": [  
        "secretsmanager:GetSecretValue"  
      ],  
      "Resource": [  
        "arn:aws:secretsmanager:region:account ID:secret:secret ID"  
      ]  
    },  
    {  
      "Effect": "Allow",  
      "Action": [  
        "kms:Decrypt"  
      ],  
      "Resource": [  
        "arn:aws:kms:region:account ID:key/key ID"  
      ]  
    },  
    {  
      "Effect": "Allow",  
      "Action": [  
        "kendra:BatchPutDocument",  
        "kendra:BatchDeleteDocument"  
      ],  
      "Resource": [  
        "arn:aws:kendra:region:account ID:index/index ID"  
      ],  
      "Condition": {  
        "StringLike": {  
          "kms:ViaService": [  
            "kendra.amazonaws.com"  
          ]  
        }  
      }  
    }  
  ],  
  {  
    "Effect": "Allow",  
    "Action": [  
      "kendra:BatchPutDocument",  
      "kendra:BatchDeleteDocument"  
    ],  
    "Resource": [  
      "arn:aws:kendra:region:account ID:index/index ID"  
    ],  
    "Condition": {  
      "StringLike": {  
        "kms:ViaService": [  
          "kendra.amazonaws.com"  
        ]  
      }  
    }  
  }  
]
```

```
        "Effect": "Allow",
        "Action": [
            "s3:GetObject"
        ],
        "Resource": [
            "arn:aws:s3:::bucket name/*"
        ]
    }
]
```

You must apply the following trust policy to the role.

```
{
  "Version": "2012-10-17",
  "Statement": {
    "Effect": "Allow",
    "Principal": {
      "Service": "kendra.amazonaws.com"
    },
    "Action": "sts:AssumeRole"
  }
}
```

If you have encrypted the Amazon S3 bucket that contains the SSL certificate used to communicate with the Sharepoint site, provide a policy to give Amazon Kendra access to the key.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": [
        "kms:Decrypt"
      ],
      "Resource": [
        "arn:aws:kms:region:account ID:key/key ID"
      ]
    }
  ]
}
```

IAM roles for frequently asked questions

When you use the [CreateFaq \(p. 139\)](#) operation to load questions and answers into an index, you must provide Amazon Kendra with an IAM role with access to the Amazon S3 bucket that contains the source files. If the source files are encrypted, you must provide permission to use the AWS KMS customer master key (CMK) to decrypt the files.

A required role policy to enable Amazon Kendra to access an Amazon S3 bucket.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": [
        "s3:GetObject"
      ],

```

```
        "Resource": [
            "arn:aws:s3:::bucket name/*"
        ]
    }
}
}
```

A required trust policy to enable Amazon Kendra to assume a role.

```
{
  "Version": "2012-10-17",
  "Statement": {
    "Sid": "AllowKendraToAssumeAttachedRole",
    "Effect": "Allow",
    "Principal": {
      "Service": "kendra.amazonaws.com"
    },
    "Action": "sts:AssumeRole"
  }
}
```

An optional role policy to enable Amazon Kendra to use an AWS KMS customer master key (CMK) to decrypt files in an Amazon S3 bucket.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": [
        "kms:Decrypt"
      ],
      "Resource": [
        "arn:aws:kms:region:account ID:key/key ID"
      ],
      "Condition": {
        "StringLike": {
          "kms:ViaService": [
            "kendra.amazonaws.com"
          ]
        }
      }
    }
  ]
}
```

Creating an index

You can create an index using the console, the AWS Command Line Interface, or by calling the [CreateIndex \(p. 142\)](#) API operation. The following procedures show how to create an index. Once you have created your index, you can add documents directly to your index or you can add them from a data source.

To create an index, you need to provide the Amazon Resource Name (ARN) of an IAM role that has permissions to any Amazon S3 bucket that you use and to perform actions on your behalf.

To create an index (Console)

1. Sign into the AWS Management Console and open the Amazon Kendra console at <https://console.aws.amazon.com/kendra/>.
2. Choose **Create index**.
3. In **Specify index details**, give your index a name and a description.
4. In **IAM role** provide an IAM role. You can either choose from roles in your account that contain the word "kendra" or you can type the name of another role. For more information about the permissions that the role requires, see [IAM roles for indexes \(p. 27\)](#).
5. Choose **Create**.
6. Creating an index can take some time. Check the list of indexes to watch the progress of creating your index. When the status of the index is **ACTIVE**, your index is ready to use.

To create an index (CLI)

1. Use the following command to create an index. The role-arn should be the Amazon Resource Name (ARN) of a role that can execute Amazon Kendra actions. For more information, see [IAM access roles for Amazon Kendra \(p. 27\)](#).

```
aws kendra create-index \
  --index-name index name \
  --description "index description" \
  --role-arn arn:aws:iam::account ID:role/role name
```

2. Creating an index can take some time. Use the index ID returned by `create-index` with the following command to check the state of your index. When the status of the index is **ACTIVE**, your index is ready to use.

```
aws kendra describe-index \
  --index-id index ID
```

To create an index (SDK)

1. You need to provide values for the following variables:
 - `description` – A description of the index that you are creating.
 - `index_name` – The name of the index that you are creating.

- `role_arn` – The Amazon Resource Name (ARN) of a role that can execute Amazon Kendra actions. For more information, see [IAM access roles for Amazon Kendra \(p. 27\)](#).
2. The following examples create an index with Amazon Kendra.

Python

```
import boto3
from botocore.exceptions import ClientError
import pprint
import time

kendra = boto3.client("kendra")

print("Create an index")

description = "index description"
index_name = "index-name"
role_arn = "arn:aws:iam::${account id}:role/${role name}"

try:
    index_response = kendra.create_index(
        Description = description,
        Name = index_name,
        RoleArn = role_arn
    )

    pprint.pprint(index_response)

    index_id = index_response["IndexId"]

    print("Wait for Kendra to create the index.")

    while True:
        # Get index description
        index_description = kendra.describe_index(
            Id = index_id
        )
        # If status is not CREATING quit
        status = index_description["Status"]
        print("    Creating index. Status: "+status)
        if status != "CREATING":
            break
        time.sleep(60)

except ClientError as e:
    print("%s" % e)

print("Program ends.")
```

Java

```
package com.amazonaws.kendra;

import java.util.concurrent.TimeUnit;
import software.amazon.awssdk.services.kendra.KendraClient;
import software.amazon.awssdk.services.kendra.model.CreateIndexRequest;
import software.amazon.awssdk.services.kendra.model.CreateIndexResponse;
import software.amazon.awssdk.services.kendra.model.DescribeIndexRequest;
import software.amazon.awssdk.services.kendra.model.DescribeIndexResponse;
import software.amazon.awssdk.services.kendra.model.IndexStatus;

public class CreateIndexExample {
```



```

public static void main(String[] args) throws InterruptedException {

    String indexDescription = "Getting started index for Kendra";
    String indexName = "java-getting-started-index";
    String indexRoleArn = "arn:aws:iam::<your AWS account ID>:role/
KendraRoleForGettingStartedIndex";

    System.out.println(String.format("Creating an index named %s", indexName));
    CreateIndexRequest createIndexRequest = CreateIndexRequest
        .builder()
        .description(indexDescription)
        .name(indexName)
        .roleArn(indexRoleArn)
        .build();
    KendraClient kendra = KendraClient.builder().build();
    CreateIndexResponse createIndexResponse =
kendra.createIndex(createIndexRequest);
    System.out.println(String.format("Index response %s",
createIndexResponse));

    String indexId = createIndexResponse.id();

    System.out.println(String.format("Waiting until the index with ID %s is
created.", indexId));
    while (true) {
        DescribeIndexRequest describeIndexRequest =
DescribeIndexRequest.builder().id(indexId).build();
        DescribeIndexResponse describeIndexResponse =
kendra.describeIndex(describeIndexRequest);
        IndexStatus status = describeIndexResponse.status();
        if (status != IndexStatus.CREATING) {
            break;
        }

        TimeUnit.SECONDS.sleep(60);
    }

    System.out.println("Index creation is complete.");
}
}

```

Once you have created your index, you add documents to it. You can either add them directly or you can create a data source that automatically updates your index on a regular schedule.

Topics

- [Adding documents directly to an index \(p. 42\)](#)
- [Adding documents from a data source \(p. 46\)](#)
- [Deleting Data Sources \(p. 61\)](#)
- [Creating custom document attributes \(p. 62\)](#)
- [Mapping data source fields \(p. 63\)](#)
- [Configuring Amazon Kendra to use a VPC \(p. 65\)](#)

Adding documents directly to an index

You can add documents directly to an index using the [BatchPutDocument \(p. 129\)](#) operation. You can't add documents directly using the console. When you are using the console, you use a data source to add documents.

You can add only the following types of documents with the `BatchPutDocuments` operation.

- Plain text
- HTML
- PDF
- Microsoft PowerPoint
- Microsoft Word

Documents can be added from an Amazon S3 bucket or supplied as binary data. The following examples show how to add documents directly to an index.

Topics

- [Adding documents with the API \(p. 42\)](#)
- [Adding documents from an Amazon S3 bucket \(p. 43\)](#)
- [Adding questions and answers \(p. 45\)](#)

Adding documents with the API

The following example adds text to an index by calling the [BatchPutDocument \(p. 129\)](#) operation.

You can use the `BatchPutDocument` operation to add documents in the following formats:

- DOC
- HTML
- PDF
- Plain text
- PPT

Files added to the index must be in a UTF-8 encoded byte stream. The following example adds UTF-8 encoded text to the index.

Python

```
import boto3

kendra = boto3.client('kendra')

index_id = '${indexID}'

title = 'Information about Amazon.com'
text = 'Amazon.com is an online retailer.'

document = {
    "Id": "1",
    "Blob": text,
    "ContentType": "PLAIN_TEXT",
```

```
        "Title": title
    }

    documents = [
        document
    ]

    result = kendra.batch_put_document(
        IndexId = index_id,
        Documents = documents
    )

    print(result)
```

Java

```
package com.amazonaws.kendra;

import software.amazon.awssdk.core.SdkBytes;
import software.amazon.awssdk.services.kendra.KendraClient;
import software.amazon.awssdk.services.kendra.model.BatchPutDocumentRequest;
import software.amazon.awssdk.services.kendra.model.BatchPutDocumentResponse;
import software.amazon.awssdk.services.kendra.model.ContentType;
import software.amazon.awssdk.services.kendra.model.Document;

public class AddDocumentsViaAPIExample {
    public static void main(String[] args) {
        KendraClient kendra = KendraClient.builder().build();

        String indexId = "yourIndexId";

        Document testDoc = Document
            .builder()
            .title("The title of your document")
            .id("a_doc_id")
            .blob(SdkBytes.fromUtf8String("your text content"))
            .contentType(ContentType.PLAIN_TEXT)
            .build();

        BatchPutDocumentRequest batchPutDocumentRequest = BatchPutDocumentRequest
            .builder()
            .indexId(indexId)
            .documents(testDoc)
            .build();

        BatchPutDocumentResponse result =
            kendra.batchPutDocument(batchPutDocumentRequest);

        System.out.println(String.format("BatchPutDocument Result: %s", result));
    }
}
```

Adding documents from an Amazon S3 bucket

You can add documents directly to your index from an Amazon S3 bucket. You can add up to 10 documents in the same call. When you use an S3 bucket, you must provide an IAM role with permission to access the bucket containing your documents. You specify the role in the `RoleArn` parameter.

Using the [BatchPutDocument](#) (p. 129) operation to add documents from an Amazon S3 bucket is a one-time operation. To keep an index synchronized with the contents of a bucket, create an S3 data source. For more information, see [Using an Amazon S3 data source](#) (p. 47).

The following example adds two Microsoft Word documents to the index using the `BatchPutDocument` operation.

Python

```
import boto3

kendra = boto3.client('kendra')

index_id = '${indexId}'
role_arn = 'arn:aws:iam:${accountID}:policy/${roleName}'

doc1_s3_file_data = {
    'Bucket': '${bucketName}',
    'Key': 'document1.docx'
}

doc1_document = {
    'S3Path': doc1_s3_file_data,
    'Title': 'Document 1 title',
    'Id': 'doc_1'
}

doc2_s3_file_data = {
    'Bucket': '${bucketName}',
    'Key': 'document2.docx'
}

doc2_document = {
    'S3Path': doc2_s3_file_data,
    'Title': 'Document 2 title',
    'Id': 'doc_2'
}

documents = [
    doc1_document,
    doc2_document
]

result = kendra.batch_put_document(
    Documents = documents,
    IndexId = index_id,
    RoleArn = role_arn
)

print(result)
```

Java

```
package com.amazonaws.kendra;

import software.amazon.awssdk.services.kendra.KendraClient;
import software.amazon.awssdk.services.kendra.model.BatchPutDocumentRequest;
import software.amazon.awssdk.services.kendra.model.BatchPutDocumentResponse;
import software.amazon.awssdk.services.kendra.model.Document;
import software.amazon.awssdk.services.kendra.model.S3Path;

public class AddFilesFromS3Example {
    public static void main(String[] args) {
        KendraClient kendra = KendraClient.builder().build();

        String indexId = "yourIndexId";
        String roleArn = "yourIndexRoleArn";
    }
}
```

```
Document pollyDoc = Document
    .builder()
    .s3Path(
        S3Path.builder()
            .bucket("an-aws-kendra-test-bucket")
            .key("What is Amazon Polly.docx")
            .build()
    )
    .title("What is Amazon Polly")
    .id("polly_doc_1")
    .build();

Document rekognitionDoc = Document
    .builder()
    .s3Path(
        S3Path.builder()
            .bucket("an-aws-kendra-test-bucket")
            .key("What is Amazon Rekognition.docx")
            .build()
    )
    .title("What is Amazon rekognition")
    .id("rekognition_doc_1")
    .build();

BatchPutDocumentRequest batchPutDocumentRequest = BatchPutDocumentRequest
    .builder()
    .indexId(indexId)
    .roleArn(roleArn)
    .documents(pollyDoc, rekognitionDoc)
    .build();

BatchPutDocumentResponse result =
kendra.batchPutDocument(batchPutDocumentRequest);

System.out.println(String.format("BatchPutDocument result: %s", result));
    }
}
```

Adding questions and answers

You add questions and answers (FAQs) to your index using the console or the [CreateFaq \(p. 139\)](#) operation. The data for the FAQ comes from a comma-separated values (csv) document stored in an Amazon S3 bucket. The file contains a list of questions, answers, and optionally the URI of a document that contains more information. You can provide one or more questions for each answer.

The following is a csv file that provides answers to questions about the height of buildings in Seattle.

```
What is the height of the Space Needle?, 605 feet, https://www.spaceneedle.com/
How tall is the Space Needle?, 605 feet, https://www.spaceneedle.com/
What is the height of the Smith Tower?, 484 feet, https://www.smithtower.com
How tall is the Smith Tower, 484 feet, https://www.smithtower.com/
```

Once you store your FAQ input file in an Amazon S3 bucket, you use the console or the [CreateFaq](#) operation to put the questions and answers into your index. You must provide an IAM role that has access to the bucket containing your source files. You specify the role in the console or in the `RoleArn` parameter. The following is a program that adds an FAQ file to an index.

Python

```
import boto3
```

```
kendra = boto3.client('kendra')

index_id = '${indexId}'
role_arn = 'arn:aws:iam::${accountId}:role/${roleName}'

faq_path = {
    'Bucket': '${bucketName}',
    'Key': 'SeattleBuildings.csv'
}

response = kendra.create_faq(
    S3Path = faq_path,
    Name = 'SeattleBuildings',
    IndexId = index_id,
    RoleArn = role_arn
)

print(response)
```

Java

```
package com.amazonaws.kendra;

import software.amazon.awssdk.services.kendra.KendraClient;
import software.amazon.awssdk.services.kendra.model.CreateFaqRequest;
import software.amazon.awssdk.services.kendra.model.CreateFaqResponse;
import software.amazon.awssdk.services.kendra.model.S3Path;

public class AddFaqExample {
    public static void main(String[] args) {
        KendraClient kendra = KendraClient.builder().build();

        String indexId = "yourIndexId";
        String roleArn = "your role for accessing S3 files";

        CreateFaqRequest createFaqRequest = CreateFaqRequest
            .builder()
            .indexId(indexId)
            .name("SeattleBuildings")
            .roleArn(roleArn)
            .s3Path(
                S3Path
                    .builder()
                    .bucket("an-aws-kendra-test-bucket")
                    .key("SeattleBuildings.csv")
                    .build()
            )
            .build();

        CreateFaqResponse response = kendra.createFaq(createFaqRequest);

        System.out.println(String.format("The result of creating FAQ: %s", response));
    }
}
```

Adding documents from a data source

When you create a data source you give Amazon Kendra the location of documents that it should index. Unlike adding documents directly to an index, you can periodically scan the data source to update the index.

For example, say that you have a repository of tax instruction stored in an Amazon S3 bucket. Existing documents are changed and new documents are added to the repository from time to time. If you add the repository to Amazon Kendra as a data source, you can keep your index up to date by periodically updating your index.

You can update the index manually using console or the [StartDataSourceSyncJob](#) (p. 185) operation, or you can set up a schedule to update the index.

An index can have more than one data source. Each data source can have its own update schedule. For example, you might update the index of your working documents daily, or even hourly, while updating your archived documents manually whenever the archive changes.

Setting an update schedule

Configure your data source to periodically update with the console or by using the `Schedule` parameter when you create or update a data source. The content of the parameter is a string that holds either a cron-format schedule string or an empty string to indicate that the index should be updated on demand. For the format of a cron expression, see [Schedule Expressions for Rules](#) in the *Amazon CloudWatch Events User Guide*. Amazon Kendra only supports cron expressions, it does not support rate expressions.

- [Using an Amazon S3 data source](#) (p. 47)
- [Using a database data source](#) (p. 50)
- [Using a Microsoft OneDrive data source](#) (p. 52)
- [Using a Salesforce data source](#) (p. 54)
- [Using a ServiceNow data source](#) (p. 57)
- [Using a Microsoft SharePoint data source](#) (p. 60)

Using an Amazon S3 data source

Use an S3 data source when your document repository is an Amazon S3 bucket.

You must create an index before you create a data source. You provide the index identifier as a parameter to the [CreateDataSource](#) (p. 132) operation.

Amazon Kendra must have permission to access the Amazon S3 bucket that contains your documents. You provide the Amazon Resource Name (ARN) of a role that has access when you create the data source using the `RoleARN` parameter.

The following examples demonstrate creating an S3 data source. The examples assume that you have already created an index and an IAM role with permission to read the data from the index. For more information about the IAM role, see [IAM roles for Amazon S3 data sources](#) (p. 29). For more information about creating an index, see [Creating an index](#) (p. 39).

CLI

```
aws kendra create-data-source \
  --index-id index ID \
  --name example-data-source \
  --type S3 \
  --configuration '{"S3n-Configuration":{"BucketName":"bucket name"}}'
  --role-arn 'arn:aws:iam::account id:role:role name'
```

Python

The following snippet of Python code creates an S3 data source. For the complete example, see [Getting started \(AWS SDK for Python \(Boto 3\)\)](#) (p. 16).

```
print("Create an S3 data source")

name = "getting-started-data-source"
description = "Getting started data source."
s3_bucket_name = "${bucketName}"
type = "S3"
role_arn = "arn:aws:iam::${accountID}:role/${roleName}"

configuration = {"S3DataSourceConfiguration":
    {
        "BucketName": s3_bucket_name
    }
}

data_source_response=kendra.create_data_source(
    Configuration = configuration,
    Name = name,
    Description = description,
    RoleArn = role_arn,
    Type = type,

    IndexId = index_id
)
```

It can take some time to create your data source. You can monitor the progress by using the [DescribeDataSource \(p. 151\)](#) operation. When the data source status is **ACTIVE** the data source is ready to use.

The following examples demonstrate getting the status of a data source.

CLI

```
aws kendra describe-data-source \
  --index-id index ID \
  --id data source ID
```

Python

The following snippet of Python code gets information about an S3 data source. For the complete example, see [Getting started \(AWS SDK for Python \(Boto 3\)\) \(p. 16\)](#).

```
print("    Wait for Kendra to create the data source.")

while True:
    data_source_description = kendra.describe_data_source(
        Id = "data source ID",
        IndexId = "index ID"
    )
    status = data_source_description["Status"]
    print("Creating data source. Status: "+status)
    time.sleep(60)
    if status != "CREATING":
        break
```

This data source doesn't have a schedule, so it will not run automatically. To index the data source you call the [StartDataSourceSyncJob \(p. 185\)](#) operation to synchronize the index with the data source.

The following examples demonstrate synchronizing a data source.

CLI

```
aws kendra start-data-source-sync-job \  
  --index-id index ID \  
  --id data source ID
```

Python

The following snippet of Python code synchronizes an S3 data source. For the complete example, see [Getting started \(AWS SDK for Python \(Boto 3\)\)](#) (p. 16).

```
print("Synchronize the data source.")  
  
sync_response = kendra.start_data_source_sync_job(  
    Id = "data source ID",  
    IndexId = "index ID"  
)
```

S3 document metadata

You can add metadata, additional information about a document, to documents in an Amazon S3 bucket using a metadata file. Each metadata file is associated with an indexed document.

Your metadata files must be stored in the same bucket as your indexed files. You can specify a location within the bucket for your metadata files using the console or the `S3Prefix` field of the `DocumentsMetadataConfiguration` parameter when you create an S3 data source. If you don't specify an S3 prefix, your metadata files must be stored in the same location as your indexed documents.

If you specify an S3 prefix for your metadata files, they live in a directory structure parallel to your indexed documents. Amazon Kendra only looks in the specified directory for your metadata. If the metadata isn't read, check that the directory location matches the location of your metadata.

The following examples show how the indexed document location maps to the metadata file location. Note that the document's S3 key is appended to the metadata's S3 prefix and then suffixed with `.metatdata.json` to form the metadata file's S3 path.

```
Bucket name:  
  s3://bucketName  
Document path:  
  documents  
Metadata path:  
  none  
File mapping  
  s3://bucketName/documents/file.txt ->  
    s3://bucketName/documents/file.txt.metatdata.json
```

```
Bucket name:  
  s3://bucketName  
Document path:  
  documents/legal  
Metadata path:  
  metadata  
File mapping  
  s3://bucketName/documents/legal/file.txt ->  
    s3://bucketName/metadata/documents/legal/file.txt.metatdata.json
```

Your document metadata is defined in a JSON file. The file must be a UTF-8 text file without a BOM marker. The file name of the JSON file should be `document.extension.metadata.json`, where "document" is the name of the document that the metadata applies to and "extension" is the file extension for the document.

The content of the JSON file follows this template. All of the attributes are optional. If you don't specify the `_source_uri`, then the links returned by Amazon Kendra in search results point to the S3 bucket that contains the document.

```
{
  "DocumentId": "document ID",
  "Attributes": {
    "_category": "document category",
    "_created_at": "ISO 8601 encoded string",
    "_last_updated_at": "ISO 8601 encoded string",
    "_source_uri": "document URI",
    "_version": "file version",
    "_view_count": "number of times document has been viewed",
    "custom attribute key": "custom attribute value",
    additional custom attributes
  },
  "AccessControlList": [
    {
      "Name": "user name",
      "Type": "GROUP | USER",
      "Access": "ALLOW | DENY"
    }
  ],
  "Title": "document title",
  "ContentType": "HTML | MS_WORD | PDF | PLAIN_TEXT | PPT"
}
```

The `_created_at` and `_last_updated_at` metadata fields are ISO 8601 encoded dates. For example, "2019-09-24T01:04:41Z".

You can add additional information to the `Attributes` field about a document that you use to filter queries or to group query responses. For more information, see [Creating custom document attributes \(p. 62\)](#).

The `AccessControlList` field enables you to filter the response from a query so that only certain users and groups have access to documents. For more information, see [Filtering on user context \(p. 74\)](#).

Using a database data source

You can index documents stored in a database using a database data source. After you provided connection information for the database Amazon Kendra will connect and index documents. Amazon Kendra supports the following databases:

- Amazon Aurora MySQL
- Amazon Aurora PostgreSQL
- Amazon RDS for MySQL
- Amazon RDS for PostgreSQL

Before you create a database data source, you need to create an index and create custom fields in the index for the data from the database. For more information, see [Creating an index \(p. 39\)](#) and [Mapping data source fields \(p. 63\)](#).

To use a database data source, you need to identify the following:

- Connection information such as credentials for the database stored in AWS Secrets Manager, the host name, port, and name of the data table that contains the document data.
- Column information such as the names of the columns in the data table that contain the document data and document ID, one to five columns to detect if a document has changed, and optional data table columns that map to custom index fields. You can map any of the Amazon Kendra reserved field names to a table column.
- Optionally, VPC information to connect to the database server. For more information about using a VPC, see [Configuring Amazon Kendra to use a VPC \(p. 65\)](#). If you are using a database data source with a VPC, the subnets provided in the VPC configuration must be in one of the following availability zone IDs:
 - US West (Oregon) – usw2-az1, usw2-az2, usw2-az3
 - US East (N. Virginia) – use1-az1, use1-az2, use1-az4
 - EU (Ireland) – euw1-az1, uew1-az2, euw1-az3

Database configuration provides the information required to connect to your database server. The host and port tell Amazon Kendra where to find the database server on the internet, the database name and table name tell Amazon Kendra where to find the document data on the database server.

To enable Amazon Kendra to access your documents, you must specify a user that has read access to the table that contains the documents. Amazon Kendra requires credentials for the user to access the database. You provide these credentials using AWS Secrets Manager. Once you have created the secret, you provide the Amazon Resource Name (ARN) of the secret to Amazon Kendra. The secret must contain the user name and password that Amazon Kendra uses to access the database in a JSON structure. The secret may contain additional information, but Amazon Kendra uses only the user name and password. The following is the minimum JSON structure that must be in the secret:

```
{
  "username": "user name",
  "password": "password"
}
```

The secret can contain more information, however, Amazon Kendra ignores other fields. For more information, see [What Is AWS Secrets Manager](#) in the *AWS Secrets Manager User Guide*.

The following example shows a database configuration.

```
"DatabaseConfiguration": {
  "ConnectionConfiguration": {
    "DatabaseHost": "host.subdomain.domain.tld",
    "DatabaseName": "DocumentDatabase",
    "DatabasePort": 3306,
    "SecretArn": "arn:aws:secretmanager:region:account ID:secret/secret name",
    "TableName": "DocumentTable"
  }
}
```

Note

The `DatabaseHost` field should be the RDS instance endpoint for the database. Don't use the cluster endpoint.

The SQL queries that Amazon Kendra uses to index your database use the database name, table name, and column names exactly as set in the database data source configuration. Amazon Kendra does not change the case of the names or enclose them with quotes.

A PostgreSQL database always changes unquoted table and column names to lower case. For example, if Amazon Kendra is configured to use the table name **SAMPLE_TABLE**, PostgreSQL converts it internally to

sample_table. If a table or column name contains upper case letters, the SQL query won't match the correct columns or table because PostgreSQL internally changes them to lower case.

You add document table information to an index by mapping table columns to index fields. There are two types of information that you add. The first is one to five columns that Amazon Kendra uses to determine if a document has changed since the last time that an index update was run. For example, if you have columns in your table named `LastUpdateDate` and `LastUpdateTime` you can tell Amazon Kendra to use them to determine if a document was updated.

The second type of information about columns is to map some or all of the columns in your table to index fields. For example, you can map a column that contains the document abstract to an index field. If you mark the field searchable, Amazon Kendra will use the contents of the field when determining if a document matches the query. For more information about the attributes that you can assign a custom field, see [Mapping data source fields \(p. 63\)](#).

After you map the columns you can also use the index fields as custom attributes to filter the results of a query. For more information, see [Filtering queries \(p. 72\)](#).

You must specify the `DocumentDataColumnName` and `DocumentIdColumnName` fields. The column mapped to the `DocumentIdColumnName` field must be an integer column.

The following example shows a simple column configuration for a database data source.

```
"ColumnConfiguration": {
  "ChangeDetectingColumns": [
    "LastUpdateDate",
    "LastUpdateTime"
  ],
  "DocumentDataColumnName": "TextColumn",
  "DocumentIdColumnName": "IdentifierColumn",
  "DocoumentTitleColumnName": "TitleColumn",
  "FieldMappings": [
    {
      "DataSourceFieldName": "AbstractColumn",
      "IndexFieldName": "Abstract"
    }
  ]
}
```

Using a Microsoft OneDrive data source

Amazon Kendra can use a data source to connect to Microsoft OneDrive sites to index the documents that your users create. When you use a OneDrive data source to connect Amazon Kendra to your OneDrive site, you choose the users whose documents are indexed. You can optionally provide inclusion and exclusion patterns to specify the documents to index.

To create a OneDrive data source, you must first create an Azure Active Directory (AD) application that Amazon Kendra connects to. You must grant the application the following permissions on the Microsoft Graph option:

- Read files in all site collections (File.Read.All)
- Read all users full profile (User.Read.All)
- Read directory data (Directory.Read.All)
- Read all groups (Group.Read.All)
- Read items in all site collections (Site.Read.All)

When you create the active directory application, it is assigned an application ID. You must use the active directory site to register a secret key for the application. Amazon Kendra uses the ID and key as

credentials to authenticate when it connects to the OneDrive site. You store the ID and key in an AWS Secrets Manager secret. If you are using the console to create your OneDrive data source, you can enter the credentials there to create a Secrets Manager secret or you can choose an existing Secrets Manager secret. If you are using the API, you must provide the Amazon Resource Name (ARN) of an existing secret.

The secret must contain the application ID and secret key that Amazon Kendra uses to access the site in a JSON structure. The following is the minimum JSON structure that must be stored in the secret:

```
{
  "username": "application ID",
  "password": "secret key"
}
```

The data source AWS Identity and Access Management (IAM) role must have permission to access the secret. For more information, see [IAM roles for Microsoft OneDrive data sources \(p. 32\)](#).

The secret can contain more information, but Amazon Kendra ignores it. For more information, see [What is AWS Secrets Manager](#) in the *AWS Secrets Manager User Guide*.

You must create an index before you create the OneDrive data source. For information, see [Creating an index \(p. 39\)](#). You provide the ID of the index when you create the data source.

You specify connection and other information in the console or using an instance of the [OneDriveConfiguration \(p. 248\)](#) data type. You must provide the following information:

- The credentials required to log in to the OneDrive site.
- The tenant domain that contains the OneDrive site.
- A list of users whose documents should be indexed. You can provide a list of user names, or you can provide the user names in a file stored in an Amazon Simple Storage Service (Amazon S3) bucket. If you store the list of user names in an S3 bucket, the IAM policy for the data source must provide access to the bucket and access to the key that the bucket was encrypted with, if any.

After you create a data source, you can't change the list of users for that data source. You can't:

- Change from a list of users to a list stored in an S3 bucket
- Change the S3 bucket location of a list of users
- Change the members of a list of users. You can change the contents of a list of users stored in an S3 bucket.

After the first successful synchronization of a data source, you can't change the contents of a user list file stored in an S3 bucket.

To change the list of users for a data source, delete the data source and recreate it.

You can optionally provide the following information:

- A list of inclusion and exclusion regular expressions that filter the documents that are included in the index. The regular expressions are applied to the file name of the document.
- Field mappings that map fields from your OneDrive site to Amazon Kendra index fields. For information, see [Mapping data source fields \(p. 63\)](#).

After you sync the data source, you can't change the inclusion and exclusion patterns or the remove field mapping. You can map additional fields.

You can map OneDrive properties to Amazon Kendra index fields. The following table shows the OneDrive properties that can be mapped and a suggested Amazon Kendra index field.

OneDrive field name	Suggested Amazon Kendra field name
body	_document_body
createdDateTime	_created_at
name	_document_title
webUrl	_document_id
createdBy.displayName	od_createdBy_displayName
createdBy.id	od_createdBy_id
createdBy.email	oc_createdBy_email
cTag	od_ctag
eTag	od_etag
fileSystemInfo.createdDateTime	od_fileSystemInfo_createdDateTime
fileSystemInfo.lastAccessedDateTime	od_fileSystemInfo_lastAccessedDateTime
fileSystemInfo.lastModifiedDateTime	od_fileSystemInfo_lastModifiedDateTime
file.mimeType	od_file_mimeType
lastModifiedDateTime	_last_updated_at
lastModifiedBy.displayName	od_lastModifiedBy_displayName
lastModifiedBy.id	od_lastModifiedBy_id
lastModifiedBy.email	od_lastModifiedBy_email
size	od_size
webDavUrl	oe_webDavUrl

Using a Salesforce data source

Amazon Kendra can connect to your Salesforce server to index your customer relationship information. When you use Amazon Kendra to index your Salesforce server, you can choose to index up to 17 of the standard Salesforce objects. You can also index knowledge articles, chatter feeds, and attachments.

Amazon Kendra uses the Salesforce API version 48. The Salesforce API limits the number of requests that you can make per day. If Amazon Kendra exceeds those requests, it will retry until it is able to continue.

Before you can connect Amazon Kendra to your Salesforce server, you must create a Salesforce connected app with OAuth enabled so that Amazon Kendra can connect. When you create an app, it is assigned a consumer key and a consumer secret that Amazon Kendra needs to connect to the app.

You must provide Amazon Kendra with credentials to access your Salesforce server. These credentials identify the user making the connection and the Salesforce connected app that Amazon Kendra connects to.

The credentials should be for a user with read only access to Salesforce. To create permissions for the user, clone the ReadOnly profile and then add the "View All Data" and "Manage Articles" permissions.

You store the credentials in an AWS Secrets Manager secret. If you are using the console to create your data source, you can create the secret there, or you can use an existing Secrets Manager secret. If you are using the API, you must provide the Amazon Resource Name (ARN) of an existing secret.

The secret must contain the following information:

- `authenticationUrl` – The URL of the OAuth authentication server used to authenticate with Salesforce. Typically, this is `https://login.salesforce.com/services/oauth2/token`.
- `consumerKey` – The consumer key, also called the client ID, of the Salesforce Connected App that is used to index the server. The app must have permission that allows access to the REST API.
- `consumerSecret` – The consumer secret, also called the client secret, of the Salesforce Connected App used to index the server.
- `securityToken` – The Salesforce security token associated with the account used to connect to Salesforce.
- `password` – The password associated with the account used to connect to Salesforce.
- `username` – The user name of the account used to connect to Salesforce. The account must have read access to the objects and fields that you want to index.

The credentials are stored as a JSON string in the Secrets Manager secret. The following is the minimum JSON structure that must be in the secret:

```
{
  "username": "user name",
  "password": "password",
  "securityToken": "token",
  "consumerKey": "key",
  "consumerSecret": "secret",
  "authenticationUrl": "https://login.salesforce.com/services/oauth2/token"
}
```

The data source IAM role must have permission to access the secret. For more information, see [IAM role for Salesforce data sources \(p. 34\)](#).

The secret can contain more information, however, Amazon Kendra ignores other fields. For more information, see [What is AWS Secrets Manager](#) in the *AWS Secrets Manager User Guide*.

You must create an index before you create the Salesforce data source. For more information, see [Creating an index \(p. 39\)](#). You provide the ID of the index when you create the data source.

You specify connection and other information in the console or using an instance of the [SalesforceConfiguration \(p. 262\)](#) data type. You must provide the following information:

- The URL of the Salesforce server that contains the information to index.
- The credentials required to connect to the Salesforce server.

You must provide configuration information for indexing at least one of the following:

- Salesforce objects
- Salesforce knowledge articles
- Salesforce chatter feeds

You can optionally:

- Provide configuration information for indexing attachments.

- Indicate whether Amazon Kendra should gather access control information for user context filtering.

Standard objects

Salesforce provides an extensive list of standard objects that contain information about your customer relations. You can choose to index any of these standard objects:

- Account
- Campaign
- Case
- Contact
- Contract
- Chatter
- Document
- Group
- Idea
- Lead
- Opportunity
- Partner
- Pricebook
- Product
- Profile
- Solution
- Task
- User

For each object, you must map an object field to the Amazon Kendra built-in `_body` field so that Amazon Kendra knows where to find the object content to index. You can map additional object fields to custom Amazon Kendra fields.

Salesforce enables you to add custom fields to standard objects. To use the custom field with Amazon Kendra, you must use the internal Salesforce field name. The internal name is the name of the field followed by `__c` (two underscores and the character c). For example, if you have a custom field named `AccountOriginalOwner`, the internal name is `AccountOriginalOwner__c`.

You can map fields from multiple objects to a single Amazon Kendra field. For example, you can map the `Account` object `Name` field and the `Partner` object `Name` field to the same Amazon Kendra custom field.

Once you save the mapping between an Amazon Kendra field and a Salesforce object field, you can't change the mapping. However, you can add more mappings between Amazon Kendra and Salesforce.

For more information, see [Mapping data source fields \(p. 63\)](#).

Knowledge articles

You can use Amazon Kendra to index the contents of standard knowledge articles or custom knowledge articles.

When you index standard knowledge articles, Amazon Kendra will index every article on your server, including the standard fields of custom knowledge articles. If you index custom knowledge articles,

Amazon Kendra will only index articles of that type, it won't index the contents of standard knowledge articles.

You configure indexing of knowledge articles using the console or the [SalesforceKnowledgeArticleConfiguration](#) (p. 267) object. You can indicate the status of the articles that you want to index, you can tell Amazon Kendra to index draft, published, or archived articles.

For custom knowledge articles, you must specify the name of the custom article type. You must specify the internal name of the article type, which is the name of the type plus "__kav" (two underscores followed by the characters kav). For example, if you have a customer article type called `CustomKnowledgeArticleForTech` the internal name is `CustomKnowledgeArticleForTech__kav`. You can specify up to 10 article types.

For both custom and standard knowledge articles, you must specify the name of the field that contains the content of the article. You can optionally specify the field that contains the title. You can map additional article fields to custom Amazon Kendra fields using the console or the [DataSourceToIndexFieldMapping](#) (p. 230) object.

Chatter feeds

You can index the contents of your Salesforce chatter feeds. You configure indexing using the console or the [SalesforceChatterFeedConfiguration](#) (p. 260) object.

You must specify the field in the Salesforce FeedItem table that contains the content of the item. Typically this is the "Body" column. You have the option of specifying the title of the item, typically this is the "Title" column of the FeedItem table. You can map additional fields to custom Amazon Kendra fields using the console or the [DataSourceToIndexFieldMapping](#) (p. 230) object.

By default, Amazon Kendra indexes all items on the chatter feed. You can use the console or the `IncludeFilterType` field of the [SalesforceChatterFeedConfiguration](#) object to limit indexing to only those items that are from standard Salesforce users or from active user accounts.

You can map additional fields to custom Amazon Kendra fields using the console or the [DataSourceToIndexFieldMapping](#) (p. 230) object.

Attachments

You can choose to have Amazon Kendra index attachments to standard objects, knowledge articles, and chatter feeds. You can use the console or the `CrawlAttachments` option on the [SalesforceConfiguration](#) (p. 262) structure to indicate whether attachments should be indexed.

By default, Amazon Kendra indexes all attachments. You can use the console or the API to filter attachments from the list that is indexed. To filter an attachment, you use a regular expression that is evaluated against the file name of the attachment. For example, to remove JSON files from the list of indexed files, use a regular expression that filters out files that end with ".json".

You can also restrict indexed documents by specifying the attachments to index. For example, to index only Microsoft Word files, specify a regular expression that selects files that end with ".doc" or ".docx".

Using a ServiceNow data source

Amazon Kendra can connect to your ServiceNow instance that contains a public knowledge base and service catalog to provide an index of its contents. For a walkthrough of creating a ServiceNow data source, see [Getting started with a ServiceNow data source \(Console\)](#) (p. 24).

When you use Amazon Kendra to index a ServiceNow instance, you choose the instance to index and whether to index a public knowledge base, a public service catalog, or both. You can optionally provide inclusion and exclusion patterns for document attachments stored in the knowledge base or service catalog.

For public knowledge bases in your ServiceNow instance, Amazon Kendra indexes only public articles. A knowledge base must have the public role under **Can Read** and **Cannot Read** must be null or not set

You also must provide Amazon Kendra with the credentials for an administrative user for your ServiceNow instance.

The user name and password for the ServiceNow instance must be stored in an AWS Secrets Manager secret. If you are using the Amazon Kendra console, you can enter the ServiceNow credentials there to create a new secret, or you can choose an existing secret. If you are using the Amazon Kendra API, you must provide the Amazon Resource Name (ARN) of an existing secret that contains your ServiceNow user name and password.

The secret must contain the username and password of the ServiceNow account that you want Amazon Kendra to use to access ServiceNow. The following is the minimum JSON structure that must be stored in the secret.

```
{
  "username": "user-name",
  "password": "password"
}
```

The secret can contain other information;but Amazon Kendra ignores it. For more information, see [What is Secrets Manager](#) in the *AWS Secrets Manager User Guide*.

When you create the ServiceNow data source you specify an IAM role that grants Amazon Kendra permission to access resources required to index your ServiceNow instance. The data source IAM role must have permission to access the secret and to use the AWS Key Management Service (AWS KMS) key that was used to decrypt it. For more information, see [IAM role for ServiceNow data sources \(p. 35\)](#).

You provide ServiceNow connection information in the Amazon Kendra console or by using an instance of the [ServiceNowConfiguration \(p. 274\)](#) data type. You must provide the following information:

- The ARN of the Secrets Manager secret that contains the credentials required to access the ServiceNow instance.
- The version of the ServiceNow instance. For Amazon Kendra, this is `LONDON` for the London version and `OTHERS` for all other ServiceNow versions.
- The ServiceNow instance host. For example, if the URL of the instance is `https://your-domain.service-now.com`, the host is `your-domain.service-now.com`.
- Whether to index knowledge articles, service catalogs, or both. You must also provide the name of the ServiceNow field that contains the document body.

You can optionally provide the following information:

- The name of the ServiceNow field that contains document titles. This is typically the `ServiceNow title` field. If you don't specify a document title field, Amazon Kendra uses the document ID as the title.
- Whether Amazon Kendra should index attachments to knowledge base or catalog items. If you choose to index attachments, you can specify the file type of attachments to exclude from the index.
- Field mappings that map fields in your ServiceNow instance to fields in your Amazon Kendra index. For more information, see [Mapping data source fields \(p. 63\)](#).
- An inclusion pattern to specify the file type of document attachments to include in the index. If you specify an inclusion pattern, any attachment that doesn't match the pattern isn't indexed. If the pattern includes file types that Amazon Kendra does not support, those files won't be included. For a list of supported file types, see [Types of documents \(p. 5\)](#).
- An exclusion pattern to specify the file type of document attachments to exclude from the index. If you specify an exclusion pattern, any attachment that doesn't match the pattern is indexed.

If you specify both an inclusion and an exclusion pattern, attachments that match the exclusion pattern won't be indexed even if they match the inclusion pattern.

After you sync the data source for the first time, the inclusion and exclusion patterns can't be changed.

You can map ServiceNow properties to Amazon Kendra index fields. The following table shows the ServiceNow knowledge article properties that can be mapped and a suggested Amazon Kendra index field. You can also create custom ServiceNow fields that you map to Amazon Kendra index fields.

ServiceNow field name	Suggested Amazon Kendra field name
content	_document_body
displayUrl	sn_display_url
first_name	sn_ka_first_name
kb_category	sn_ka_category
kb_catagory_name	_category
kb_knowledge_base	sn_ka_knowledge_base
last_name	sn_ka_last_name
number	sn_kb_number
published	sn_ka_publish_date
replItemType	sn_replItemType
short_description	_document_title
sys_created_by	sn_createdBy
sys_created_on	_created_at
sys_id	sn_sys_id
sys_updated_by	sn_updatedBy
sys_updated_on	_last_updated_at
url	sn_url
user_name	sn_ka_user_name
valid_to	sn_ka_valid_to
workflow_state	sn_ka_workflow_state

The following table shows the ServiceNow catalog properties that can be mapped and a suggested Amazon Kendra field.

ServiceNow field name	Suggested Amazon Kendra field name
category	sn_sc_category
category_full_name	sn_sc_category_full_name

ServiceNow field name	Suggested Amazon Kendra field name
category_name	_category
description	_document_body
displayUrl	sn_display_url
replItemType	sn_replItemType
sc_catalogs	sn_sc_catalogs
sc_catalogs_name	sn_sc_catalogs_name
short_description	_document_body
sys_created_by	sn_createdBy
sys_created_on	_created_at
sys_id	sn_sys_id
sys_updated_by	sn_updatedBy
sys_updated_on	_last_updated_at
title	_document_title
url	sn_url

Using a Microsoft SharePoint data source

You can use your Microsoft SharePoint Online site as a data source for Amazon Kendra. When you use Amazon Kendra to index your site, you choose which SharePoint URLs to include in the index, and you specify inclusion and exclusion patterns for the documents stored on those URLs.

Amazon Kendra requires credentials to access the SharePoint site. The SharePoint user must have administrative permission to the SharePoint sites that you want to index. If you are using the console to create your data source, you can enter the credentials there or you can choose an existing AWS Secrets Manager secret. If you are using the API, you must provide the Amazon Resource Name (ARN) of an existing secret.

The secret must contain the user name and password that Amazon Kendra uses to access the SharePoint site in a JSON structure. The following is the minimum JSON structure that must be in the secret:

```
{
  "username": "user name",
  "password": "password"
}
```

The data source IAM role must have permission to access the secret. For more information, see [IAM roles for Microsoft SharePoint Online data sources \(p. 36\)](#).

The secret can contain more information, however, Amazon Kendra ignores other fields. For more information, see [What Is AWS Secrets Manager](#) in the *AWS Secrets Manager User Guide*.

You must create an index before you create the SharePoint data source. For more information, see [Creating an index \(p. 39\)](#). You provide the ID of the index when you create the data source.

You specify connection and other information in the console or using an instance of the [SharePointConfiguration](#) (p. 280) data type. You must provide the following information:

- The credentials required to log in to the SharePoint site.
- The URLs of the SharePoint site, SharePoint site collection, or SharePoint list to index.
- The ARN of an IAM role that has permission to run Amazon Kendra commands. For the required permissions, see [IAM roles for Microsoft SharePoint Online data sources](#) (p. 36).

You can optionally provide the following information:

- Whether Amazon Kendra should index the contents of attachments to SharePoint list items.
- An inclusion pattern to specify the documents that should be included in the index. If you specify an inclusion pattern, any document that does not match the pattern will not be indexed.
- An exclusion pattern to specify the documents that should be excluded from the index. If you specify an exclusion pattern, any document that does not match the pattern will be indexed. If you specify both an inclusion and an exclusion pattern, documents that match the exclusion pattern will not be indexed even if they match the inclusion pattern.
- Whether Amazon Kendra should use the SharePoint change log mechanism to determine if a document needs to be updated in the index. You should use the change log if you don't want Amazon Kendra to scan all of the documents in the site to update the index. If your change log is large, it may take Amazon Kendra less time to scan the site than to process the change log.
- Field mappings that map attributes from your SharePoint site to Amazon Kendra index fields. For more information, see [Mapping data source fields](#) (p. 63).

Deleting Data Sources

You delete a data source when you want to remove the information contained in the data source from your Amazon Kendra index. For example, delete a data source when:

- A data source is incorrectly configured. Delete the data source, wait for the data source to finish deleting, and then recreate it.
- You migrated documents from one data source to another. Delete the original data source and recreate it in the new location.
- You have reached the limit of data sources for an index. Delete one of the existing data sources and add a new one. For more information about the number of data sources that you can create, see [Quotas](#) (p. 123).

To delete a data source, use the console, the AWS Command Line Interface, or the `DeleteDataSource` operation. Deleting a data source removes all of the information about the data source from the index. If you only want to stop synchronizing the data source, change the synchronization schedule for the data source to "run on demand".

To delete a data source (console)

1. Sign in to the AWS Management Console and open the Amazon Kendra console at <https://console.aws.amazon.com/kendra/>.
2. From the left menu, choose **Indexes**, and then choose the index that contains the data source to delete.
3. From the left menu, choose **Data sources**.
4. Choose the data source to remove.
5. Choose **Delete** to delete the data source.

To delete a data source (CLI)

- In the AWS Command Line Interface, use the following command. The command is formatted for Linux and macOS. If you are using Windows, replace the Unix line continuation character (\) with a caret (^).

```
aws kendra delete-data-source \  
  --id data-source-id \  
  --index-id index-id
```

When you delete a data source, Amazon Kendra removes all of the stored information about the data source. Amazon Kendra removes all of the document data stored in the index, as well as all run histories and metrics associated with the data source. Deleting a data source does not remove the original documents from your storage.

Deleting a data source is an asynchronous operation. When you start deleting a data source, the data source status is changed to `DELETING`. It remains in the `DELETING` state until the information related to the data source is removed. After the delete is complete, the data source no longer appears in the results of a call to the [ListDataSources](#) (p. 166) operation. If you call the [DescribeDataSource](#) (p. 151) operation with the deleted data source's identifier, you receive a `ResourceNotFound` exception.

Amazon Kendra releases the resources for a data source as soon as you call the `DeleteDataSource` operation or choose to delete the data source in the console. If you are deleting the data source to reduce the number of data sources below your limit, you can create a new data source right away.

If you are deleting a data source and then creating another data source to the document data, you should wait for the first data source to complete deleting before you sync the new data source.

You can delete a data source that is in the process of syncing with Amazon Kendra. The sync is stopped and the data source is removed. If you attempt to start a sync when the data source is being deleted you will get a `ConflictException` exception in response.

You can't delete a data source if the associated index is in the `DELETING` state. Deleting an index deletes all of the data sources for the index. You can start deleting an index while a data source for that index is in the `DELETING` state.

If you have two data sources pointing to the same documents, such as two data sources pointing to the same S3 bucket, you can get inconsistent data in the index when one of the data sources is deleted. When two data sources reference the same documents, only one copy of the document data is stored in the index. Removing one data source removes the index data for the documents. The other data source is not aware that the documents have been removed, so it won't correctly re-index the documents the next time it syncs. When you have two data sources pointing to the same document location, you should delete both data sources and then recreate one.

Creating custom document attributes

When the source of your data is an S3 bucket, you can apply custom attributes to your documents. For example, you could create a custom attribute called "Department" with values "HR", "Sales", and "Manufacturing". You can apply these attributes to your documents so that you can limit the response to documents in the "HR" department, for example.

For other data sources, you use field mapping for the same purpose. For more information, see [Mapping data source fields](#) (p. 63).

Amazon Kendra has six reserved attributes that you can use. The attributes are:

- `_category` (String)
- `_created_at` (ISO 8601 encoded string)
- `_last_updated_at` (ISO 8601 encoded string)
- `_source_uri` (String)
- `_version` (String)
- `_view_count` (Long)

After you have created a custom attribute, you can use the attribute when you call the `Query` operation. You can use it for faceted search, use it to filter the response, and choose whether or not the attribute should be returned in the response. For more information, see [Filtering queries \(p. 72\)](#).

Adding custom attributes with the BatchPutDocument operation

When you use the [BatchPutDocument \(p. 129\)](#) operation to add a document to your index, you specify custom attributes as part of the `Attributes` structure. You can add multiple attributes when you call the operation. The following example is a `Attributes` structure that adds "Department" and "_category" attributes to a document.

```
"Attributes": {
  "Department": "HR",
  "_category": "Vacation policy"
}
```

Adding custom attributes to an S3 data source

When you use an Amazon S3 bucket as a data source for your index, you add metadata to the documents with companion metadata files. You place the metadata JSON files in a directory structure that is parallel to your documents. For more information, see [S3 document metadata \(p. 49\)](#).

You specify custom attributes in the `Attributes` JSON structure. You can add a list of attributes. For example, the following example is an `Attributes` structure that defines three custom attributes and one reserved attribute.

```
"Attributes": {
  "brand": "Amazon Basics",
  "price": 1595,
  "_category": "sports",
  "subcategories": ["outdoors", "electronics"]
}
```

Mapping data source fields

When the source of your data is a data source other than Amazon S3 you can map data source fields to fields in your index. For example if you have a field that contains department information for a document, you can map it to an index field called "Department" so that you can use the field in queries.

You can create field mappings for the following data sources:

- Database

- Microsoft OneDrive
- Microsoft SharePoint
- Salesforce
- ServiceNow

If you are storing your documents in an Amazon S3 bucket, or if you are using an Amazon S3 data source, you use custom attributes to map index fields. For more information, see [Creating custom document attributes \(p. 62\)](#).

Mapping your data source fields to an index field is a three step process:

1. Create an index. For more information, see [Creating an index \(p. 39\)](#).
2. Update the index to add custom fields.
3. Create a data source that maps data source fields to the index fields.

To update the index to add custom fields, you use the console or the [UpdateIndex \(p. 201\)](#) operation.

When you are using the console, you can choose to map a data source field to one of the seven reserved field names, or you can choose to create a new index field that maps to the field. For database data sources, if the name of the database column matches the name of a reserved field, the field and column are automatically mapped.

With the API, you add custom and reserved fields using the `DocumentMetadataConfigurationUpdates` parameter.

The following JSON example is a `DocumentMetadataConfigurationUpdates` structure that adds a field called "Department" to the index.

```
"DocumentMetadataConfigurationUpdates": [
  {
    "Name": "Department",
    "Type": "STRING_VALUE"
  }
]
```

When you create the field you have the option of setting how the field should be used in searches. You can choose from the following:

- **Displayable** – determines whether the field is returned in the query response. The default is `true`.
- **Facetable** – indicates that the field can be used to create facets. The default is `false`.
- **Searchable** – Determines whether the field is used in the search. The default is `true` for string fields and `false` for number and date fields.

The following JSON example is a `DocumentMetadataConfigurationUpdates` structure that adds a field called "Department" to the index and marks it as facetable.

```
"DocumentMetadataConfigurationUpdates": [
  {
    "Name": "Department",
    "Type": "STRING_VALUE",
    "Search": {
      "Facetable": "true"
    }
  }
]
```



```
] ]
```

Amazon Kendra has six reserved fields that you can map to data source fields. The fields are:

- `_category` (String)
- `_created_at` (ISO 8601 encoded string)
- `_file_type` (String)
- `_last_updated_at` (ISO 8601 encoded string)
- `_source_uri` (String)
- `_view_count` (Long)

Once you have created the index fields, you can map the data source fields to the index fields. If you are using the console, you can create index fields and map data source fields using the **Custom field mappings** editor. If you are using the API, you can add field mappings using the [CreateDataSource](#) (p. 132) or [UpdateDataSource](#) (p. 195) operations.

Configuring Amazon Kendra to use a VPC

Amazon Kendra connects to your Amazon virtual private cloud (VPC) to index information stored in databases running in your private cloud. When you create the database data source you provide security group and subnet identifiers for the subnet that contains your database. Amazon Kendra uses this information to create an elastic network interface that it uses to securely communicate with your database.

If your database isn't running on an Amazon VPC, you can connect your database to your Amazon VPC using a Virtual Private Network (VPN). You get a default VPC when you create your Amazon account. For information on setting up a VPN, see the [AWS Virtual Private Network Documentation](#).

To use a VPC, you must tell Amazon Kendra the identifier of the subnet that the database belongs to and the identifiers of any security groups that Amazon Kendra must use to access the subnet. For example, if you are using the default port for a MySQL database, the security groups must enable Amazon Kendra to access port 3306 on the host that runs the database.

The identifiers for subnets and security groups are configured in the Amazon VPC control panel. To see the identifiers, open the Amazon VPC console as follows:

To view subnet identifiers

1. Sign in to the AWS Management Console and open the Amazon VPC console at <https://console.aws.amazon.com/vpc/>.
2. From the left menu, choose **Subnets**.
3. From the subnet list, choose the subnet that contains your database server.
4. In the description tab, the identifier is in the **Subnet ID** field.

To view security group identifiers

1. Sign in to the AWS Management Console and open the Amazon VPC console at <https://console.aws.amazon.com/vpc/>.
2. From the left menu, choose **Security Groups**.
3. From the security group list, choose the group that you want the identifier for.

4. In the description tab, the identifier is in the **Group ID** field.

You can provide more than one subnet if Amazon Kendra must route the connection between two or more subnets. For example, if the subnet that contains your database server is out of IP addresses, Amazon Kendra can connect to a subnet with free IP addresses and route the connection to the first subnet. If you list multiple subnets, the subnets must be able to communicate with each other. Each subnet should be associated with a route table that provides outbound internet access using a network address translator (NAT) device.

You can also provide multiple security groups. The combined effect of the security groups should allow Amazon Kendra to access the database server that you have specified in the connection configuration for the data source.

Connecting to a database in a VPC

The following example shows how to connect a database data source to a MySQL database running in a VPC. The example assumes that you are starting with your default VPC and that you need to create a MySQL database. If you already have a VPC, make sure that it is configured as shown. If you have a MySQL database, you can use that instead of creating a new one.

Topics

- [Step 1: Configure a VPC \(p. 66\)](#)
- [Step 2: Configure security \(p. 67\)](#)
- [Step 3: Create a database \(p. 67\)](#)
- [Step 4: Create a MySQL data source \(p. 68\)](#)

Step 1: Configure a VPC

First, configure your VPC so that you have a private subnet and a security group that enables Amazon Kendra to access a MySQL database running in the subnet.

To configure a VPC

1. Sign in to the AWS Management Console and open the Amazon VPC console at <https://console.aws.amazon.com/vpc/>.
2. From the left menu, choose **Route tables** then choose **Create route table**.
3. For the **Name tag** field, enter **Private subnet route table**. In the **VPC field**, choose your VPC, and then choose **Create**. Choose **Close** to return to the list of route tables.
4. From the left menu, choose **NAT Gateways** then choose **Create NAT Gateway**.
5. In the **Subnet** field, choose the subnet that should be the public subnet and note the subnet ID.
6. If you don't have an available Elastic IP, choose **Create New EIP** and then choose **Create a NAT Gateway** and then choose **Close**.
7. From the left menu, choose **Route Tables**.
8. From the route table list, choose **Private subnet route table** that you created in step 3. From the **Actions** menu, choose **Edit Routes**.
9. Choose **Add route**. Add the destination 0.0.0.0/0 to allow all outgoing traffic to the Internet. In the **Target**, choose **NAT Gateway** and then choose the gateway that you created in step 4. Finally, choose **Save routes** and then choose **Close**.
10. From **Actions**, choose **Edit subnet associations**.
11. Select the subnets that you want to be private. Do not select the subnet with the NAT gateway that you noted above.

Step 2: Configure security

Next, configure security groups for your database.

To create security groups

1. Sign in to the AWS Management Console and open the Amazon VPC console at <https://console.aws.amazon.com/vpc/>.
2. From the description of your VPC, note the IPv4 CIDR.
3. From the left menu, choose **Security Groups** and then choose **Create security group**.
4. In **Security group name** type **MySQLSecurityGroup**. Provide a description, then choose your VPC from the drop down list. Choose **Create** and then choose **Close**.
5. Choose the **Inbound** tab.
6. Choose **Edit rules** and then choose **Add Rule**
7. For a MySQL database, type 3306 for the **Port Range**. For the **Source**, type the CIDR of your VPC. Choose **Save rules** and then choose **Close**.

The security group allows anyone within the VPC to connect to the database, and it allows outbound connections to the Internet.

Step 3: Create a database

Now create a database to hold your documents. If you already have a database, you can use that instead.

To create a MySQL database

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. From the left menu, choose **Subnet groups** and then choose **Create DB Subnet Group**.
3. Name the group and choose your VPC.
4. Add your VPC's private subnets. Your private subnets are the ones that are not connected to your NAT. Choose **Create**.
5. From the left menu, choose **Databases** and then choose **Create database**.
6. Use the following parameters to create the database. Leave all of the other parameters at their defaults.
 - **Engine options** – MySQL
 - **Templates** – Free tier
 - **Credential Settings** – Enter and confirm a master password
 - Under **Connectivity**, choose **Additional connectivity configuration**
 - **Subnet group** – Choose the subnet group that you created in step 4
 - **VPC security group** – Choose the group that you created in the VPC (**MySQLSecurityGroup**).
 - Under **Additional configuration**, set the **Initial database name** to **content**.
7. Choose **Create database**.
8. From the list of databases, choose your new database. Make a note of the database endpoint.
9. After you create your database, you must create a table to hold your documents. Creating a table is outside the scope of these instructions. When you create your table, note the following:
 - Database name – **content**
 - Table name – **documents**
 - Columns – **ID**, **Title**, **Body**, and **LastUpdate**. You can include additional columns if you want.

Step 4: Create a MySQL data source

Now that you have configured your VPC and created your database, you can create a data source for the database.

To create a MySQL data source

1. Sign in to the AWS Management Console and open the Amazon Kendra console at <https://console.aws.amazon.com/kendra/>.
2. From the left menu, choose **Indexes** and then choose your index.
3. Choose **Add data sources** and then choose **Amazon RDS**.
4. Type a name and description for the data source and then choose **Next**.
5. Choose **MySQL**.
6. Under **Connection access** enter the following information:
 - **Endpoint** – The endpoint of the database that you created earlier.
 - **Port** – The port number for the database. For MySQL, the default is 3306.
 - **Type of authentication** – Choose **New**.
 - **New secret container name** – type a name for the Secrets Manager container for the database credentials.
 - **Username** – Type the name of a user with administrative access to the database.
 - **Password** – Type the password for the user and then choose **Save authentication**.
 - **Database name** – type **content**.
 - **Table name** – type **documents**.
 - **IAM role** – Choose **Create a new role** and then type a name for the role.
7. In **Column configuration** provide the following:
 - **Document ID column name** – type **ID**.
 - **Document title column name** – type **Title**.
 - **Document data column name** – type **Body**.
8. In **Column change detection** provide the following:
 - **Change detecting columns** – type **LastUpdate**.
9. In **Configure VPC & security group** provide the following:
 - In **Virtual Private Cloud (VPC)** choose your VPC.
 - In **Subnets** choose the private subnets that you created in your VPC.
 - In **VPC security groups** select the security group that you created in your VPC for MySQL databases (**MySQLSecurityGroup**).
10. In **Set sync run schedule** choose **Run on demand** and then choose **Next**.
11. In **Data source field mapping**, choose **Next**.
12. Review the configuration of your data source to make sure that it is correct. When you're satisfied that everything is correct, choose **Create**.

Searching indexes

To search an Amazon Kendra index you use the [Query \(p. 179\)](#) operation. The `query` operation returns information about the indexed documents that you use in your application. This section shows you how to make a query, perform filters, and interpret the response that you get from the `Query` operation. This section also shows how to submit feedback about the quality of a search result.

Topics

- [Querying an index \(p. 69\)](#)
- [Filtering queries \(p. 72\)](#)
- [Filtering on user context \(p. 74\)](#)
- [Query responses \(p. 77\)](#)
- [Types of response \(p. 79\)](#)
- [Submitting feedback \(p. 82\)](#)

Querying an index

When you search your index, Amazon Kendra uses all of the information that you provided about your documents to determine the documents most relevant to the search terms entered. Some of the items that Amazon Kendra considers are:

- The text of the document.
- The title of the document.
- Custom text fields that you have marked searchable.
- The date field that you have indicated should be used to determine the "freshness" of a document.

Once a set of relevant documents has been selected from the index, Amazon Kendra filters the response based on any attribute filters that you have requested for the search. For example, if you have a custom attribute called "department," you can filter the response so that it only contains documents from a department called "legal." For information about creating custom attributes, see [Creating custom document attributes \(p. 62\)](#).

After finding the relevant documents and then filtering based on the attributes that you set, Amazon Kendra returns the results. The results are sorted by the relevance that Amazon Kendra determined for each doc. The results are paginated so that you can show a page at a time to your user.

The following Python example shows how to search an index by using the [the section called "Query" \(p. 179\)](#) operation. The example determines the type of the search result (answer, document, question/answer) and displays the answer text.

For information about the query responses, see [Query responses \(p. 77\)](#).

Note

You can use this code to filter document attributes. The topic [Filtering queries \(p. 72\)](#) contains examples that you can use to replace the following code.

```
response=kendra.query(  
    QueryText = query,  
    IndexId = index)
```

Prerequisites

To run this example, you need to:

- Set up permissions. For more information, see [IAM access roles for Amazon Kendra](#) (p. 27).
- Set up the AWS CLI. For more information, see [Setting up the AWS CLI](#) (p. 9).
- Create a data source and index. For more information, see [Getting started with an S3 bucket \(Console\)](#) (p. 14).

Searching an index (Console)

You can use the Amazon Kendra search console to test your index. You can make queries and see the results.

To search an index with the console

1. Sign in to the AWS Management Console and open the Amazon Kendra console at <http://console.aws.amazon.com/kendra/>.
2. From the left menu, choose **Indexes**.
3. Choose your index.
4. From the left menu, choose **Search console**.
5. Enter a query in the text box and then press enter or choose the magnifying glass icon.
6. Amazon Kendra returns the results of the search.

Searching an index (SDK)

To search an index with Python or Java

- The following example searches an index. Change the value of `query` to your search query and `index_id` or `indexId` to the index identifier of the index that you want to search.

Python

```
import boto3
import pprint

kendra = boto3.client('kendra')

query='${searchString}'
index_id='${indexID}'

response=kendra.query(
    QueryText = query,
    IndexId = index_id)

print ('\nSearch results for query: ' + query + '\n')

for query_result in response['ResultItems']:

    print('-----')
    print('Type: ' + str(query_result['Type']))

    if query_result['Type']=='ANSWER':
        answer_text = query_result['DocumentExcerpt']['Text']
        print(answer_text)
```

```
if query_result['Type']=='DOCUMENT':
    if 'DocumentTitle' in query_result:
        document_title = query_result['DocumentTitle']['Text']
        print('Title: ' + document_title)
    document_text = query_result['DocumentExcerpt']['Text']
    print(document_text)

print ('-----\n\n')
```

Java

```
package com.amazonaws.kendra;

import software.amazon.awssdk.services.kendra.KendraClient;
import software.amazon.awssdk.services.kendra.model.QueryRequest;
import software.amazon.awssdk.services.kendra.model.QueryResponse;
import software.amazon.awssdk.services.kendra.model.QueryResultItem;

public class SearchIndexExample {
    public static void main(String[] args) {
        KendraClient kendra = KendraClient.builder().build();

        String query = "some queries";
        String indexId = "anIndexId";

        QueryRequest queryRequest = QueryRequest
            .builder()
            .queryText(query)
            .indexId(indexId)
            .build();

        QueryResponse queryResponse = kendra.query(queryRequest);

        System.out.println(String.format("\nSearch results for query: %s", query));
        for(QueryResultItem item: queryResponse.resultItems()) {
            System.out.println("-----");
            System.out.println(String.format("Type: %s", item.type()));

            switch(item.type()) {
                case QUESTION_ANSWER:
                case ANSWER:
                    String answerText = item.documentExcerpt().text();
                    System.out.println(answerText);
                    break;
                case DOCUMENT:
                    String documentTitle = item.documentTitle().text();
                    System.out.println(String.format("Title: %s", documentTitle));
                    String documentExcerpt = item.documentExcerpt().text();
                    System.out.println(String.format("Excerpt: %s",
documentExcerpt));
                    break;
                default:
                    System.out.println(String.format("Unknown query result type:
%s", item.type()));
            }

            System.out.println("-----\n");
        }
    }
}
```

Filtering queries

You can improve the response from the [Query \(p. 179\)](#) operation by using filters. Filters restrict the documents in the response to ones that directly apply to the query. You can create faceted search suggestions, use boolean logic to filter out documents that don't match a specified criteria, or filter out specific document attributes from the response.

Facets

Facets are scoped views of a set of search results. For example, if your index provides search results for cities across the world, you can use facets to offer your users search results filtered to a specific city. Or, you can create facets for search results that are written by a specific author. The following example shows how to get facet information for the `_category` custom attribute.

```
response=kendra.query(  
    QueryText = query,  
    IndexId = index,  
    Facets = [  
        {  
            "DocumentAttributeKey" : "_category"  
        }  
    ]  
)
```

Facet information is returned in the `FacetResults` response array. You use the contents to display faceted search suggestions in your application. For example, if the document attribute `Category` contains the city that a search could apply to, use that information to display a list of city searches. Users can select a city to get faceted search results scoped to the city. To make the faceted search, call [Query \(p. 179\)](#) and use the selected document attribute to filter the results. For an example, see [Using document attributes to filter queries \(p. 73\)](#).

The following sample JSON response shows facets scoped to the `_category` document attribute. The response includes the count of documents that include the value of the document attribute.

```
{  
  'FacetResults': [  
    {  
      'DocumentAttributeKey': '_category',  
      'DocumentAttributeValueCountPairs': [  
        {  
          'Count': 3,  
          'DocumentAttributeValue': {  
            'StringValue': 'Dubai'  
          }  
        },  
        {  
          'Count': 3,  
          'DocumentAttributeValue': {  
            'StringValue': 'Seattle'  
          }  
        },  
        {  
          'Count': 1,  
          'DocumentAttributeValue': {  
            'StringValue': 'Paris'  
          }  
        }  
      ]  
    }  
  ]  
}
```



```
]
```

When you use a string list field to create facets, the facet results returned are based on the contents of the string list. For example, if you have a string list field that contains two items, one with the value "corgi","dachshund," and one with the value "husky," you will get a `FacetsResults` structure with three facets.

For more information, see [Query responses \(p. 77\)](#).

Using document attributes to filter queries

By default, `Query` returns all search results. To filter responses, you can perform logical operations on the document attributes. For example, if you only want documents for a specific city, you can filter on the `City` and `State` custom document attributes. You use the [AttributeFilter \(p. 210\)](#) input parameter to create a boolean operation on filters that you supply.

The following example shows how to perform a logical AND operation by filtering on a specific city, *Seattle*, and state, *Washington*.

```
response=kendra.query(  
    QueryText = query,  
    IndexId = index,  
    AttributeFilter = {'AndAllFilters':  
        [  
            {"EqualsTo": {"Key": "City", "Value": {"StringValue": "Seattle"}}},  
            {"EqualsTo": {"Key": "State", "Value": {"StringValue": "Washington"}}}  
        ]  
    }  
)
```

The following example shows how to perform a logical OR operation for when any of the `Fileformat`, `Author`, or `SourceURI` keys match the specified values.

```
response=kendra.query(  
    QueryText = query,  
    IndexId = index,  
    AttributeFilter = {'OrAllFilters':  
        [  
            {"EqualsTo": {"Key": "Fileformat", "Value": {"StringValue":  
"AUTO_DETECT"}}},  
            {"EqualsTo": {"Key": "Author", "Value": {"StringValue": "Ana Carolina"}}},  
            {"EqualsTo": {"Key": "SourceURI", "Value": {"StringValue": "https://  
aws.amazonaws.com/234234242342"}}}  
        ]  
    }  
)
```

For `StringList` fields, use the `ContainsAny` or `ContainsAll` attribute filters to return documents with the specified string. The following example shows how to return all documents that have the values "Seattle" or "Portland" in their `Locations` custom attribute.

```
response=kendra.query(  
    QueryText = query,  
    IndexId = index,  
    AttributeFilter = {  
        "ContainsAny": { "Key": "Locations", "Value": { "StringListValue":  
[ "Seattle", "Portland" ] }}  
    }  
)
```

Filtering a document's attributes

By default, all document attributes for a document are returned in the `DocumentAttributes` response field. You can choose to only include specific document attributes by using the `IncludeDocumentAttributes` input parameter. In the following example, only the `SourceURI` and `Author` document attributes are included in the response for a document.

```
response=kendra.query(  
    QueryText = query,  
    IndexId = index,  
    IncludeDocumentAttributes = ["SourceURI", "Author"]  
)
```

Filtering on user context

Amazon Kendra enables you to filter a query response to remove documents from the result based on a user name and group membership. User context filtering depends on two things:

- Setting attributes on the indexed documents.
- Including user and group information when querying the index.

Warning

User context filtering requires you to provide user and group information. If no user and group information is included in the query, Amazon Kendra returns all documents.

You can add a user or a group to either an allow list or a deny list for the document. If a user or group is added to the deny list, the document is filtered out of any query that contains the user or group.

You specify the user and group information when you query the index using the built-in attributes `_user_id` and `_group_ids`. You can set up to 100 group identifiers.

Amazon Kendra does not match users to groups. When you use user-context filtering, you must include all of the groups that a user belongs to when you make the query. For example, if a user belongs to two groups, "HR" and "IT," you must include both groups in the request.

The following example shows a request that filters the query response based on the user ID and groups. The query will return any document that has the user or the "HR" or "IT" groups in the allow list. If the user or either group is in the deny list for a document, the document will not be returned.

```
response = kendra.query(  
    QueryText = query,  
    IndexId = index,  
    AttributeFilter = {  
        "OrAllFilters": [  
            {  
                "EqualsTo": {  
                    "Key": "_user_id",  
                    "Value": {  
                        "StringValue": "user1"  
                    }  
                }  
            },  
            {  
                "EqualsTo": {  
                    "Key": "_group_ids",
```

```
        "Value": {  
            "StringListValue": ["HR", "IT"]  
        }  
    }  
]  
}  
)
```

Note

User context filtering isn't an authentication or authorization control for your content. It doesn't do user authentication on the user and groups sent to the `Query` operation. It is up to your application to ensure that the user and group information sent to `Query` operation is authenticated and authorized.

There is an implementation of user context filtering for each data source. The following section describes each implementation.

Topics

- [User context filtering for documents added directly to an index \(p. 75\)](#)
- [User context filtering for frequently asked questions \(p. 75\)](#)
- [User context filtering for database data sources \(p. 75\)](#)
- [User context filtering for Microsoft OneDrive data sources \(p. 76\)](#)
- [User context filtering for Amazon S3 data sources \(p. 76\)](#)
- [User context filtering for Salesforce data sources \(p. 77\)](#)
- [User context filtering for ServiceNow data sources \(p. 77\)](#)
- [User context filtering for Microsoft SharePoint data sources \(p. 77\)](#)

User context filtering for documents added directly to an index

To specify user context filters for a document that you add directly to an index, you provide the filter in the `AccessControlList` field of the document. You provide three pieces of information:

- The access that the entity should have. You can say `ALLOW` or `DENY`.
- The type of entity. You can say `USER` or `GROUP`.
- The name of the entity.

You can add up to 200 entries in the `AccessControlList` field.

User context filtering for frequently asked questions

You can't add user context filtering to frequently asked question (FAQ) index entries.

User context filtering for database data sources

For a database data source, information for user context filtering comes from a column in the source table. You specify the column name in the console or when you create the data source using the `AclConfiguration` field.

A database data source has the following limitations:

- You can only specify an allow list for a database data source. You can't specify a deny list.
- You can only specify groups. You can't specify individual users for the allow list.
- The database column should be string containing a semi-colon delimited lists of groups.

User context filtering for Microsoft OneDrive data sources

Amazon Kendra retrieves user and group information from Microsoft OneDrive when it indexes the documents on the site. The user and group information is taken from the underlying Microsoft SharePoint site that hosts OneDrive.

When you use a OneDrive user or group for user context filtering, calculate the ID as follows:

- Get the site name. For example, `https://host.onmicrosoft.com/sites/siteName`.
- Take the MD5 hash of the site name. For example, `430a6b90503eef95c89295c8999c7981`.
- Create the user email or group ID by concatenating the MD5 hash with a vertical bar (|) and the ID. For example, if a group name is "site owners", the group ID would be

`"430a6b90503eef95c89295c8999c7981|site owners"`

For the user name "someone@host.onmicrosoft.com" the user ID would be

`"430a6b90503eef95c89295c8999c7981|someone@host.onmicrosoft.com"`

Send the user or group ID to Amazon Kendra as the `_user_id` or `_group_ids` attribute when you call the [Query](#) (p. 179) operation. For example, the AWS CLI command that uses a group to filter the query response looks like this:

```
aws kendra query \
  --index-id index ID
  --query-text "query text"
  --attribute-filter '{
    "EqualsTo":{
      "Key": "_group_ids",
      "Value": {"StringValue": "430a6b90503eef95c89295c8999c7981|site
owners"}}'
```

User context filtering for Amazon S3 data sources

You add user context filtering to a document in an Amazon S3 data source using a metadata file associated with the document. You add the information to the `AccessControlList` field in the JSON document. For more information about adding metadata to the documents indexed from an Amazon S3 data source, see [S3 document metadata](#) (p. 49).

You provide three pieces of information:

- The access that the entity should have. You can say `ALLOW` or `DENY`.
- The type of entity. You can say `USER` or `GROUP`.
- The name of the entity.

You can add up to 200 entries in the `AccessControlList` field.

User context filtering for Salesforce data sources

A Salesforce data source returns user and group information from Salesforce access control list (ACL) entities. The ProfileSet, PermissionSet, Group, and Role entities are mapped to the Amazon Kendra `_group_ids` index field. The Salesforce User entity is mapped to the Amazon Kendra `_user_id` index field.

The `_group_id` for a Salesforce chatter feed has two possible values. If the feed is visible only to a Chatter or Collaboration group, that group is the Amazon Kendra `_group_id` value. If not, Amazon Kendra creates a group named after the user that posted the chatter item. All users that subscribed to the item are part of that group.

User context filtering for ServiceNow data sources

User context filtering isn't currently supported for ServiceNow.

User context filtering for Microsoft SharePoint data sources

Amazon Kendra retrieves user and group information from Microsoft SharePoint when it indexes the documents on the site. To filter your documents, provide user and group information when you call the `Query` operation.

To filter using a user name, use the user's email address. For example, `johnstiles@example.com`.

When you use a SharePoint group for user context filtering, calculate the group ID as follows:

- Get the site name. For example, `https://host.onmicrosoft.com/sites/siteName`.
- Take the MD5 hash of the site name. For example, `430a6b90503eef95c89295c8999c7981`.
- Create the group ID by concatenating the MD5 hash with a vertical bar (|) and the group name. For example, if the group name is "site owners", the group ID would be

```
"430a6b90503eef95c89295c8999c7981|site owners"
```

Send the group ID to Amazon Kendra as the `_group_ids` attribute when you call the [Query \(p. 179\)](#) operation. For example, the AWS CLI command looks like this:

```
aws kendra query \
  --index-id index ID
  --query-text "query text"
  --attribute-filter '{
    "EqualsTo":{
      "Key": "_group_ids",
      "Value": {"StringValue": "430a6b90503eef95c89295c8999c7981|site
owners"}}
  }'
```

Query responses

A call to [Query \(p. 179\)](#) returns information about the results of a search. The results are in an array of [QueryResultItem \(p. 252\)](#) objects (`ResultItems`). Each `QueryResultItem` includes a summary of the result. Document attributes associated with the query result are included.

Summary information

The summary information varies depending on the type of result. In each case it includes document text that matches the search term. It also includes highlight information that you can use to highlight the search text in your application's output. For example, if the search term is *what is the height of the Space Needle?*, the summary information includes text location for the words *height* and *space needle*. For information about response types, see [Types of response \(p. 79\)](#).

Document attributes

Each result contains document attributes for the document that matches a query. Some of the attributes are predefined, such as `DocumentId`, `DocumentTitle` and `DocumentUri`. Others are custom attributes that you define. You can use document attributes to filter the response from the `Query` operation. For example, you might want only the documents written by a specific author or a specific version of a document. For more information, see [Filtering queries \(p. 72\)](#). You specify document attributes when you add documents to an index. For more information, see [Creating custom document attributes \(p. 62\)](#).

The following is sample JSON code for a query result. Note the document attributes in `DocumentAttributes` and `AdditionalAttributes`.

```
{
  "QueryId": "query-id",
  "ResultItems": [
    {
      "Id": "result-id",
      "Type": "ANSWER",
      "AdditionalAttributes": [
        {
          "Key": "AnswerText",
          "ValueType": "TEXT_WITH_HIGHLIGHTS_VALUE",
          "Value": {
            "TextWithHighlightsValue": {
              "Text": "text",
              "Highlights": [
                {
                  "BeginOffset": 55,
                  "EndOffset": 90,
                  "TopAnswer": false
                }
              ]
            }
          }
        }
      ],
      "DocumentId": "document-id",
      "DocumentTitle": {
        "Text": "title"
      },
      "DocumentExcerpt": {
        "Text": "text",
        "Highlights": [
          {
            "BeginOffset": 0,
            "EndOffset": 300,
            "TopAnswer": false
          }
        ]
      },
      "DocumentURI": "uri",
      "DocumentAttributes": []
    },
    {
      "Id": "result-id",
      "Type": "DOCUMENT",
      "AdditionalAttributes": [],
      "DocumentId": "document-id",
```

```
    "DocumentTitle": {
      "Text": "title",
      "Highlights": []
    },
    "DocumentExcerpt": {
      "Text": "text",
      "Highlights": [
        {
          "BeginOffset": 74,
          "EndOffset": 77,
          "TopAnswer": false
        }
      ]
    },
    "DocumentURI": "uri",
    "DocumentAttributes": [
      {
        "Key": "_source_uri",
        "Value": {
          "StringValue": "uri"
        }
      }
    ]
  },
  "FacetResults": [],
  "TotalNumberOfResults": number
}
```

Types of response

Amazon Kendra returns three types of query response.

- Answer
- Document
- Question and answer

The type of the response is returned in the Type response field of [QueryResultItem \(p. 252\)](#).

Answer

Amazon Kendra detected one or more question answers in the response. A factoid is the response to a who, what, when, or where question such as *What is the height of the space needle?* Amazon Kendra returns text in the index that best matches the query. The text is in the `AnswerText` field and contains highlight information for the search term within the response text.

```
{
  'AnswerText': {
    'Highlights': [
      {
        'BeginOffset': 271,
        'EndOffset': 279,
        'TopAnswer': False
      },
      {
        'BeginOffset': 481,
        'EndOffset': 489,
        'TopAnswer': False
      }
    ]
  }
}
```

```

        {
            'BeginOffset': 547,
            'EndOffset': 555,
            'TopAnswer': False
        },
        {
            'BeginOffset': 764,
            'EndOffset': 772,
            'TopAnswer': False
        }
    ],
    'Text': 'Asynchronousoperationscan\n''alsoprocess
\n''documentsthatareinPDF''format.UsingPDFformatfilesallowsyoutoprocess''multi-
page\n''documents.\n''Forinformationabouthow''AmazonTextractrepresents
\n''documentsasBlockobjects,
    'seeDocumentsandBlockObjects.\n''\n''\n''\n''Forinformationaboutdocument''limits,
    seeLimitsinAmazonTextract.
\n''\n''\n''\n''TheAmazonTextractsynchronous''operationscanprocessdocumentsstoredinanAmazon
\n''S3Bucketoryoucanpass''base64encodedimagebytes.\n''Formoreinformation,

    see''CallingAmazonTextractSynchronousOperations.''Asynchronousoperationsrequireinputdocuments
\n''tobesuppliedinanAmazon''S3Bucket.'
},
    'Excerpt': {
        'Highlights': [
            {
                'BeginOffset': 0,
                'EndOffset': 300,
                'TopAnswer': False
            }
        ],
        'Text': 'Asynchronousoperationscan\n''alsoprocess
\n''documentsthatareinPDF''format.UsingPDFformatfilesallowsyoutoprocess''multi-page
\n''documents.\n''ForinformationabouthowAmazon''Textractrepresents\n''
    },
    'Type': 'ANSWER'
}

```

Document

Amazon Kendra returns ranked documents for those that match the search term. The ranking is based on the confidence that Amazon Kendra has in the accuracy of the search result. Information about the matching document is returned in the [QueryResultItem](#) (p. 252). It includes the title of the document, The excerpt includes highlight information for search text and the section of matching text in the document. The URI for matching documents is in the `SourceURI` document attribute. The following sample JSON show the document summary for a matching document.

```

{
    'DocumentTitle': {
        'Highlights': [
            {
                'BeginOffset': 7,
                'EndOffset': 15,
                'TopAnswer': False
            },
            {
                'BeginOffset': 97,
                'EndOffset': 105,
                'TopAnswer': False
            }
        ],
        'Text': 'AmazonTextractAPIPermissions: Actions,
\n''Permissions,

```



```

        andResourcesReference-'AmazonTextract'
    },
    'Excerpt': {
        'Highlights': [
            {
                'BeginOffset': 68,
                'EndOffset': 76,
                'TopAnswer': False
            },
            {
                'BeginOffset': 121,
                'EndOffset': 129,
                'TopAnswer': False
            }
        ],
        'Text': '...LoggingandMonitoring\tMonitoring
\n'\tCloudWatchMetricsforAmazonTextract
\n'\tLoggingAmazonTextractAPICallswithAWSCloudTrail\n'\tAPIReference\tActions
\tAnalyzeDocument\n'\tDetectDocumentText\n'\tGetDocumentAnalysis...'
    },
    'Type': 'DOCUMENT'
}

```

Question and answer

A question and answer response is returned when Amazon Kendra matches a question with one of the frequently asked questions in your index. The response includes the matching question and answer in the [QueryResultItem \(p. 252\)](#) field. It also includes highlight information for query terms detected in query string. The following JSON shows a question and answer response. Note that the response includes the question text

```

{
    'AnswerText': {
        'Highlights': [

        ],
        'Text': '605feet'
    },
    'Excerpt': {
        'Highlights': [
            {
                'BeginOffset': 0,
                'EndOffset': 8,
                'TopAnswer': False
            }
        ],
        'Text': '605feet'
    },
    'Type': 'QUESTION_ANSWER',
    'QuestionText': {
        'Highlights': [
            {
                'BeginOffset': 12,
                'EndOffset': 18,
                'TopAnswer': False
            },
            {
                'BeginOffset': 26,
                'EndOffset': 31,
                'TopAnswer': False
            },
            {
                'BeginOffset': 32,

```

```
        'EndOffset': 38,  
        'TopAnswer': False  
    }  
],  
  'Text': 'whatistheheightoftheSpaceNeedle?'  
}  
}
```

For information about adding question and answer text to an index, see [Adding questions and answers \(p. 45\)](#)

Submitting feedback

You can provide feedback about query results returned from an Amazon Kendra index. Amazon Kendra uses this information to improve the underlying search model and provide better search results over time. To submit feedback, you use the [SubmitFeedback \(p. 189\)](#) operation. To identify the query, you supply the `IndexID` of the index that the query applies to, and the `QueryId` returned in the response from the [Query \(p. 179\)](#) operation. Two types of feedback are accepted.

- **Clicks** - Information about which query results were chosen by the end user. The feedback includes the result ID (`ResultId`) and the Unix timestamp of the date and time that the search result was chosen. Your application will need to collect click information from the activities of your end-users.
- **Relevance** - Information about the relevance of a search result. This is typically information provided by the end-user. The feedback contains the result ID and an relevance indicator (`RELEVANT` or `NOT_RELEVANT`). Relevance information is determined by your end user. To submit relevance feedback, your application needs to provide a feedback mechanism that allows the end-user to choose the appropriate relevance for a query result.

The following example shows how to submit click and relevance feedback. You can submit multiple sets of feedback through the `ClickFeedbackItems` and `RelevanceFeedbackItems` arrays. This example submits a single click and a single relevance feedback item. The current timestamp is used for the timing of the feedback submittal.

To submit feedback for a search

1. Use the following code. Change:
 - a. `index_id` to the ID of the index that the query applies to.
 - b. `query_id` to the query that you want to provide feedback on.
 - c. `result_id` to the ID of the query result that you want to provide feedback on. The result ID is returned in the query response.
 - d. `relevance_value` to either `RELEVANT` (the query result is relevant) or `NOT_RELEVANT` (the query result is not relevant).

Python

```
import boto3  
import time  
  
kendra = boto3.client('kendra')  
  
index_id = '${indexID}'  
query_id = '${queryID}'  
result_id = '${resultID}'  
feedback_item = {'ClickTime': int(time.time())},
```

```
        'ResultId':result_id}

relevance_value = 'RELEVANT'
relevance_item = {'RelevanceValue': relevance_value,
                  'ResultId':result_id
                  }

response=kendra.submit_feedback(
    QueryId = query_id,
    IndexId = index_id,
    ClickFeedbackItems = [feedback_item],
    RelevanceFeedbackItems = [relevance_item]
)

print ('Submitted feedback for query: ' + query_id)
```

Java

```
package com.amazonaws.kendra;

import java.time.Instant;
import software.amazon.awssdk.services.kendra.KendraClient;
import software.amazon.awssdk.services.kendra.model.ClickFeedback;
import software.amazon.awssdk.services.kendra.model.RelevanceFeedback;
import software.amazon.awssdk.services.kendra.model.RelevanceType;
import software.amazon.awssdk.services.kendra.model.SubmitFeedbackRequest;
import software.amazon.awssdk.services.kendra.model.SubmitFeedbackResponse;

public class SubmitFeedbackExample {
    public static void main(String[] args) {
        KendraClient kendra = KendraClient.builder().build();

        SubmitFeedbackRequest submitFeedbackRequest = SubmitFeedbackRequest
            .builder()
            .indexId("anIndexId")
            .queryId("aQueryId")
            .clickFeedbackItems(
                ClickFeedback
                    .builder()
                    .clickTime(Instant.now())
                    .resultId("aResultId")
                    .build()
            )
            .relevanceFeedbackItems(
                RelevanceFeedback
                    .builder()
                    .relevanceValue(RelevanceType.RELEVANT)
                    .resultId("aResultId")
                    .build()
            )
            .build();

        SubmitFeedbackResponse response =
            kendra.submitFeedback(submitFeedbackRequest);

        System.out.println("Feedback is submitted");
    }
}
```

```
import boto3
import time

kendra = boto3.client('kendra')

index_id = '${indexID}'
query_id = '${queryID}'
result_id = '${resultID}'
feedback_item = {'ClickTime': int(time.time()),
                 'ResultId':result_id}

relevance_value = 'RELEVANT'
relevance_item = {'RelevanceValue': relevance_value,
                 'ResultId':result_id
                 }

response=kendra.submit_feedback(
    QueryId = query_id,
    IndexId = index_id,
    ClickFeedbackItems = [feedback_item],
    RelevanceFeedbackItems = [relevance_item]
)

print ('Submitted feedback for query: ' + query_id)
```

2. Run the code. A message is displayed after the feedback has been submitted.

Deploying Amazon Kendra

When it comes time to deploy Amazon Kendra search to your Web site, we provide source code that you can use with React to get a head start on your application. The source code is provided free of charge under a modified MIT license so that you can use it as is or change it for your own needs. You can download the sample from <https://kendrasamples.s3.amazonaws.com/kendrasamples.zip>.

The example application is modeled after the search page of the Amazon Kendra console. It has the same features for searching and displaying search results. You can use the whole example application, or you can choose just one of the features for your own use.

To see the three components of the search page in the console, choose the code icon (</>) from the right menu. Hover your pointer over each section to see a brief description of the component and to get the URL of the component's source.

Topics

- [Overview \(p. 85\)](#)
- [Prerequisites \(p. 85\)](#)
- [Setting up the example \(p. 86\)](#)
- [Main search page \(p. 86\)](#)
- [Search component \(p. 86\)](#)
- [Results component \(p. 86\)](#)
- [Facets component \(p. 87\)](#)
- [Pagination component \(p. 87\)](#)

Overview

You add the example code to an existing React application to enable search. The search files and components are structured as follows:

- Main search page – this is the main page that contains all of the components. This is where you will integrate your application with the Amazon Kendra API.
- Search bar – this is the component where a user enters a search term and that calls the search function.
- Results – this is the component that displays the results from Amazon Kendra. It has three components: Suggested answers, FAQ results, and recommended documents.
- Facets – This is the component that shows the facets in the search results and enables you to choose a facet to limit the search.
- Pagination – this is the component that paginates the response from Amazon Kendra.

Prerequisites

Before you begin you need the following:

- An existing React Web application.
- A development environment configured with the correct libraries.

- The SDK for Java or AWS SDK for JavaScript.

Information about the required libraries and AWS SDKs is in the Readme file in the `kendrasamples.zip` file.

Setting up the example

The following is an overview of adding Amazon Kendra search to a React page.

To install the example application

1. Download the source files from the repository. <https://kendrasamples.s3.amazonaws.com/kendrasamples.zip>
2. Create a new Web page for search, or you can choose an existing page to add the search bar to.
3. Add the `Search.tsx` component to the page.
4. Build the change into your package.

A complete procedure for adding Amazon Kendra search to a React application is in the Readme included in the `kendrasamples.zip` file.

Main search page

The main search page contains all of the example search components. It includes the search bar component for output, the results components to display the response from the [Query \(p. 179\)](#) operation, and a pagination component that enables you to page through the response.

Search component

The search component provides a text box to enter query text. The `onSearch` function is a hook that calls the main function in `Search.tsx` to make the Amazon Kendra [Query \(p. 179\)](#) operation call.

Results component

The results component shows the response from the `Query` operation. The results are shown in three separate areas.

- Suggested answers – These are the top results returned by the `Query` operation. It contains up to three suggested answers. In the response, they have the result type `ANSWER`.
- FAQ answers – These are the frequently asked questions results returned by the response. FAQs are added to the index separately. In the response, they have the type `QUESTION_ANSWER`. For more information, see [Adding questions and answers \(p. 45\)](#).
- Recommended documents – These are additional documents that Amazon Kendra returns in the response. In the response from the `Query` operation, they have the type `DOCUMENT`.

The results components share a set of components for features like highlighting, titles, links, etc. The shared components must be present for the result components to work.

Facets component

The facets component lists the facets available in the search results. Each facet classifies the response along a specific dimension, such as author. You can refine the search to a specific facet by choosing one from the list.

After you select a facet, the component calls the `Query` operation with an attribute filter that restricts the search to only those documents that match the facet.

Pagination component

The pagination component enables you to display the search results from the `Query` operation in multiple pages. It calls the `Query` operation with the `PageSize` and `PageNumber` parameters to get a specific page of results.

Troubleshooting

When you synchronize your Amazon Kendra index with a data source, you may run into issues that prevent the documents from being indexed. Indexing is a two step process. First, there is the data source level process of crawling the data source to find the new and updated documents to index, and to find any documents to remove from the index. Second, there is the document level process where each document is accessed and indexed.

An error can occur in either of these steps. Data source level errors are reported in the console in the **Sync run history** section of the data source details page. The status of the synchronization job can be **Succeeded**, **Incomplete**, or **Failed**. You can also see the number of documents indexed and deleted during the job. If the status is **Failed**, a message is shown in the **Details** column.

Document level errors are reported in Amazon CloudWatch Logs. You can see the errors using the CloudWatch console.

My synchronization job failed

A synchronization job typically fails when there is a configuration error in the index or the data source. The error message in the Details column of the data source gives information about what went wrong. The problem is usually that the index or the data source does not have the proper IAM permissions. The error message describes the permissions that are missing. Here are some of the error messages that you can receive:

Failed to create log group for job. Please make sure that the IAM role provided has sufficient permissions.

If your index role does not have permission to use CloudWatch, the data source will not be able to create a CloudWatch log. If you get this error you need to add CloudWatch permissions to the index role.

Failed to access S3 file prefix (*bucket name*) while trying to crawl your metadata files. Please make sure the IAM Role (*role ARN*) provided has sufficient permissions.

When you are using an Amazon S3 data source, Amazon Kendra must have permission to access the bucket that contains the documents. You need to add permission for Amazon Kendra to read the bucket to the data source IAM role.

The provided IAM Role (*role ARN*) could not be assumed. Please make sure Amazon Kendra is a trusted entity that is allowed to assume the role.

Amazon Kendra needs permission to assume the index and data source IAM roles. You need to add a trust policy to the roles with permission for the `sts:AssumeRole` action.

For the IAM policies that Amazon Kendra needs to index a data source, see [IAM access roles for Amazon Kendra](#) (p. 27).

My synchronization job is incomplete

Jobs are generally incomplete when they have completed the data source level process but have had some error during the document level process. When a job is incomplete, some of the documents may have been successfully indexed. For an Amazon S3 data source, an incomplete job is typically caused by:

- The metadata for one or more documents was invalid.
- When there are documents to submit for indexing, at least one document was not submitted.
- When there are documents to submit for deleting from the index, at least one document was not submitted.

To troubleshoot an incomplete synchronization job, look first to your CloudWatch logs.

1. From the details column, choose **View details in CloudWatch**.
2. Review the error messages to see what caused the document to fail.

My synchronization job succeeded but there are no indexed documents

Occasionally you will have a index synchronization job run that is marked as **Succeeded** but there are no new or updated documents indexed when you expect there to be. Here are some reasons:

- Check CloudWatch `DocumentsSubmittedForIndexingFailed` metric to see if there were any documents that failed to synchronize. Check your CloudWatch logs for details.
- For an Amazon S3 data source, you may have given Amazon Kendra the wrong bucket name or prefix. Make sure that the bucket that Amazon Kendra is using is the one that contains the documents to index.
- If you are re-indexing a document that failed to be indexed in an earlier job, Amazon Kendra won't index it unless you make a change to the document or its associated metadata file.

General troubleshooting

Amazon Kendra uses CloudWatch metrics and logs to provide you with insight into synchronizing your data sources. You can use the metrics and logs to determine what went wrong with a synchronization run and to determine what you need to do to fix it.

For general troubleshooting, start with your CloudWatch metrics.

- Check the `DocumentsCrawled` metric to see how many documents your data source checked. For an Amazon S3 bucket, if the number is less than you expect, check to be sure that your data source is pointing to the right bucket.
- Check the `DocumentsSkippedNoChange` metric to see how many documents were skipped because they haven't changed since the last synchronization. If the number does not match what you expect, check to make sure that your repository was updated correctly.
- Check the `DocumentsSkippedInvalidMetadata` metric to see how many documents had invalid metadata. Check your CloudWatch logs to see the specific errors that occurred.
- Check the `DocumentsSubmittedForIndexingFailed` metric to see the number of documents that were sent from the data source to the index but failed to be indexed for some reason. For example, if you use a metadata attribute in an Amazon S3 data source that hasn't been defined as a custom index field the document will not be indexed. Check your CloudWatch logs to see the specific errors that occurred.
- Check the `DocumentsSubmittedForDeletionFailed` metric to see how many documents that the data source attempted to remove from the index failed to be deleted from the index. Check your CloudWatch logs to see the specific errors that occurred.

You can look at the CloudWatch logs for a particular synchronization run to get details of the errors that occurred during the run. For more information about CloudWatch logs with Amazon Kendra, see [Monitoring Amazon Kendra with Amazon CloudWatch Logs \(p. 104\)](#).

Manually tuning an index

Amazon Kendra queries produce search results ranked by their relevance. The searchable fields in the index all contribute to this ranking.

You can modify the effect of field or attribute on the relevance of a document through *relevance tuning*. When you use relevance tuning, a result is given a boost in the response when the query includes terms that match the field or attribute. You also specify how much of a boost the document receives when there is a match. Relevance tuning doesn't cause Amazon Kendra to include a document in the query response, it is only one of the factors that Amazon Kendra uses to determine the relevance of a document.

You can boost certain fields in your index to assign more importance to specific responses. Amazon Kendra enables you to tune for specific data sources or document freshness. For example when searching for "When is re:Invent?" you boost the relevance of document freshness so that the latest date is the suggested answer. Or, in an index of research reports, you could boost a more reputable data source.

Amazon Kendra also supports boosting documents based on votes or view counts which is common in forums and other support knowledge bases. You can combine boosts to, for example, boost documents that are not only viewed more but that are also more recent, like trending news or updates.

You set the amount of boost that a document receives using the `Importance` parameter. The larger the number in this parameter, the more the field or attribute will boost the relevance of the document. When you are tuning your index, you should increase the value of the `Importance` parameter in small increments until you get the effect that you want. Query your index and compare the results to previous queries to determine if you are getting improved search results.

You can specify date, number, or string fields and attributes to tune an index. Each type of field or attribute has specific criteria for when it boosts a result.

- **Date fields and attributes** – There are three specific criteria for date fields, `Duration`, `Freshness` and rank order.

`Duration` sets the time period that the boost applies to. For example, if you set the time period to 86400 seconds (1 day) the boost begins to drop off after one day. The higher the importance, the faster the boost effect drops off.

`Freshness` indicates that the field or attribute should be used to determine how "fresh" a document is. For example, if document 1 was created on November 14, and document 2 was created on November 5, document 1 is "fresher" than document 2. The fresher the document, the more this boost is applied. You can only have one `Freshness` field in your index.

Rank order applies the boost in either ascending or descending order. When you use `ASCENDING`, later dates have precedence. If you specify `DESCENDING`, earlier dates have precedence.

- **Number fields and attributes** – For number fields or attributes you can specify the rank order that Amazon Kendra should use when determining the relevance of the field or attribute. If you specify `ASCENDING` larger numbers are given precedence. If you specify `DESCENDING`, lower numbers have precedence.

- **String fields and attributes** – For string fields or attributes you can create subcategories of a field to give a different boost to different values in the field or attribute. For example, if you are boosting a field or attribute called "Department," you can give a different boost to documents from "HR" to those from "Legal".

You tune the relevance of a field or attribute using the console or the [UpdateIndex \(p. 201\)](#) operation. For example, the following example marks the "_last_updated_at" field as the freshness field for a document.

```
"DocumentMetadataConfigurationUpdates" : [  
  "Name": "_last_updated_at",  
  "Type": "DATE_VALUE",  
  "Relevance": {  
    "Freshness": TRUE,  
    "Importance": 2  
  }  
]
```

The following example applies a different importance to values in the "department" field.

```
"DocumentMetadataConfigurationUpdates" : [  
  "Name": "department",  
  "Type": "STRING_VALUE",  
  "Relevance": {  
    "Importance": 2,  
    "ValueImportanceMap": {  
      "HR": 3,  
      "Legal": 1  
    }  
  }  
]
```

Adjusting capacity

Amazon Kendra provides resources for your index in *capacity units*. Each capacity unit provides additional resources for your index. There are separate capacity units for storage and for queries. You can only add capacity units to Amazon Kendra Enterprise indexes. You can't add capacity to a Developer edition index.

A document storage capacity unit provides the following additional storage for your index.

- 500,000 documents or 150 GB of storage

A query capacity unit provides the following additional queries for your index.

- 0.5 queries per second or approximately 40,000 queries per day

Each index comes with a base capacity equal to 1 capacity unit. There is an additional cost for each additional capacity unit. For details, see [Amazon Kendra pricing](#).

You can adjust capacity units up to 5 times per day to tune the capacity of your index to the expected usage. You can't reduce document storage capacity below the number of documents stored in your index. For example, if you are storing 600,000 documents, you can't reduce the storage capacity below 1 additional unit.

You can add up to 20 capacity units to your storage and query resources. If you need additional resources, [contact AWS support](#).

You can use the Amazon Kendra console or the [DescribeIndex \(p. 162\)](#) operation to view the resources that your index is using. Amazon Kendra also returns exceptions when you exceed the capacity of an index. You get a `ServiceQuotaExceededException` exception when you exceed your storage capacity and a `ThrottlingException` exception when you exceed your queries per second capacity.

Viewing capacity

View the resources that your index is using with the Amazon Kendra console. The console provides graphs that you can use to determine how much storage and query capacity your index uses. You can use this information to help you plan when to add additional capacity.

To view document storage and query use (Console)

1. Sign into the AWS Management Console and open the Amazon Kendra console at <https://console.aws.amazon.com/kendra/home>.
2. From the list of indexes, choose the index to view capacity.
3. The **Document count** and **Storage used** fields of the **Index details** show you the amount of storage that your index uses. Scroll to the **Queries per second** chart to see a graph of the queries per second for your index.

Adding and removing capacity

If you need additional capacity for your index, you can add it using the console or the Amazon Kendra API. When you add capacity, don't add more than you need to help reduce costs.

To add or remove storage or query capacity (Console)

1. Sign into the AWS Management Console and open the Amazon Kendra console at <https://console.aws.amazon.com/kendra/home>.
2. From the list of indexes, choose the index that you want to add or remove capacity.
3. From the **Actions** menu, choose **Edit**.
4. On the **Specify index details** page, choose **Next**.
5. On the **Add additional capacity** page, choose the new query per second and document storage capacity units that you want to use for the index.
6. Choose **Update** to save your changes and update your index to the new capacity.

After you update the capacity of your index, it can take up to 60 minutes for the changes to take effect.

To add or remove capacity using the Amazon Kendra API, use the `CapacityUnits` parameter [UpdateIndex](#) (p. 201) operation.

Monitoring and logging for Amazon Kendra

Topics

- [Monitoring your index \(Console\) \(p. 95\)](#)
- [Logging Amazon Kendra API calls with AWS CloudTrail logs \(p. 98\)](#)
- [Monitoring Amazon Kendra with Amazon CloudWatch \(p. 100\)](#)
- [Monitoring Amazon Kendra with Amazon CloudWatch Logs \(p. 104\)](#)

Monitoring your index (Console)

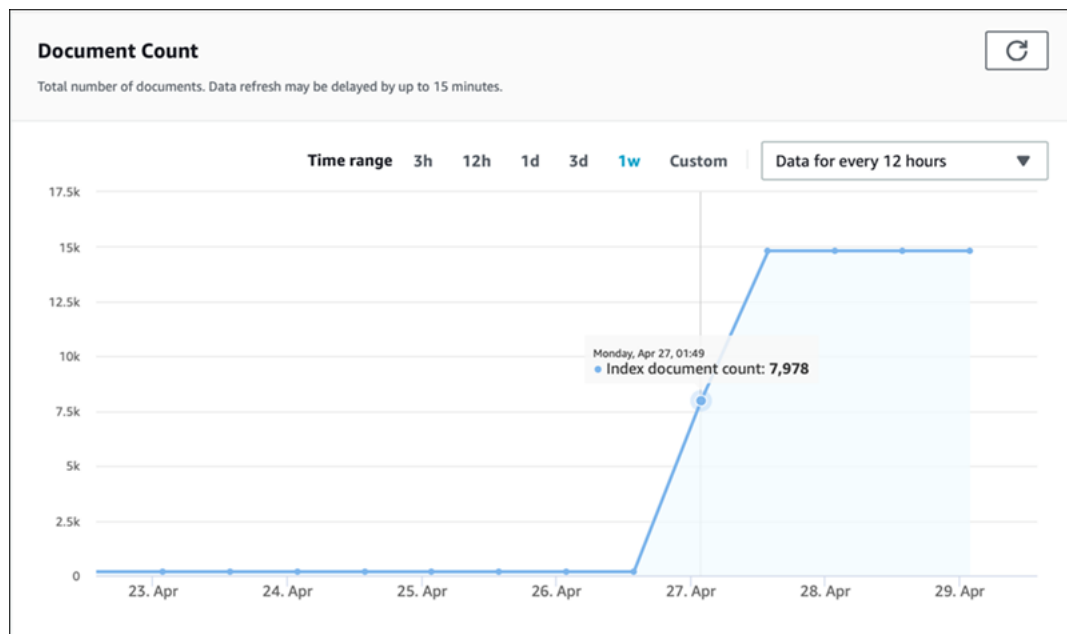
Use the Amazon Kendra console to monitor the state of indexes and data sources. You can use this information to track the size and storage requirements of your index and to monitor the progress and success of synchronization between your index and data sources.

To view index metrics (Console)

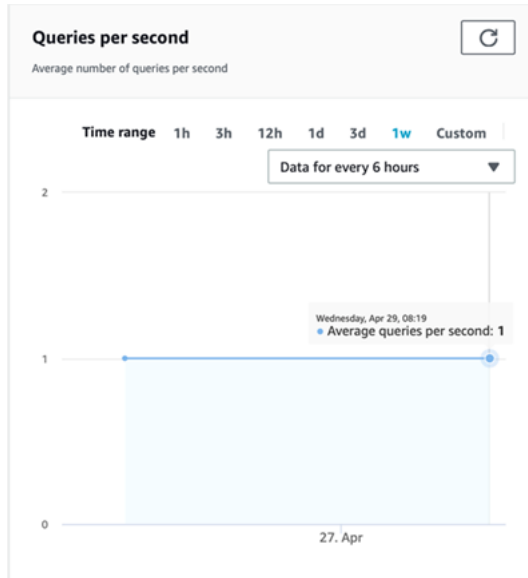
1. Sign into the AWS Management Console and open the Amazon Kendra console at <https://console.aws.amazon.com/kendra/home>.
2. From the list of indexes, choose the index to view.
3. Scroll the screen to see the index metrics.

You can see the following metrics about your index.

- **Document count** – The total number of documents indexed. This includes all documents from all data sources. Use this metric to determine if you need to purchase more or fewer storage units for your index.



- **Queries per second** – The number of index queries that are requested each second. Use this metric to determine if you need to purchase more or fewer query units for your index.



To monitor the progress and success of synchronization between your index and a data source, use the Amazon Kendra console. Use this information to help determine the health of your data source.

To view synchronization metrics (Console)

1. Sign into the AWS Management Console and open the Amazon Kendra console at <https://console.aws.amazon.com/kendra/home>.
2. From the list of indexes, choose the index to view synchronization metrics for.
3. From the left menu, choose **Data sources**.
4. From the list of data sources, choose the data source to view.
5. Scroll the screen to see the sync run metrics.

You can see the following information.


- **Sync run history** – Statistics about the synchronization run, including the start and end time, the number of documents added, deleted, and failed. If the sync run fails, there is a link to CloudWatch Logs with more information. Choose the settings icon in the upper left to change the columns that are displayed in the history. Use this information to determine the general health of your data source.

Sync run history (5)


< 1 > ⚙

Status / Summary	Start time	End time	Added / Modified	Deleted	Failed	Details
⌵ Syncing - indexing	Apr 29, 2020, 9:53 AM PDT	Apr 29, 2020, 9:54 AM PDT	⌵	⌵	⌵	View in CloudWatch
✓ Succeeded	Apr 28, 2020, 1:35 PM PDT	Apr 28, 2020, 1:37 PM PDT	1484	0	2	Service is operating normally ↗
✓ Succeeded	Apr 28, 2020, 1:32 PM PDT	Apr 28, 2020, 1:32 PM PDT	0	0	0	Service is operating normally ↗
✓ Succeeded	Apr 28, 2020, 1:05 PM PDT	Apr 28, 2020, 1:06 PM PDT	5	0	0	Service is operating normally ↗
✓ Succeeded	Apr 28, 2020, 1:05 PM PDT	Apr 28, 2020, 1:05 PM PDT	298	0	1	Service is operating normally ↗


- **Document count** – The total number of documents indexed from this data source. This is the total of all documents added to the data source minus the total of all documents deleted from the data source. Use this information to determine how many documents from this data source are included in the index.



- **Document scans** – The total number of documents scanned during the sync run. This includes all documents in the data source, including those added, updated, deleted, or unchanged. Use this information to determine if Amazon Kendra is scanning all of the documents in the data source. The number of documents scanned affects the amount charged for the service.



- **Average sync run time in minutes** – The average length of time that it takes for a sync run to complete. The time that it takes to sync a data source affects the amount charged for the service.



Logging Amazon Kendra API calls with AWS CloudTrail logs

Amazon Kendra is integrated with AWS CloudTrail, a service that provides a record of actions taken by a user, role, or an AWS service in Amazon Kendra. CloudTrail captures all API calls from Amazon Kendra as events, including calls from the Amazon Kendra console and from code calls to the Amazon Kendra APIs. If you create a trail, you can enable continuous deliver of CloudTrail events to an Amazon S3 bucket, including events for Amazon Kendra. If you don't configure a trail, you can still view the most recent events in the CloudTrail console in **Event history**. Using the information collected by CloudTrail, you can determine the request that was made to Amazon Kendra, the IP address from which the request was made, who made the request, when it was made, and additional details.

To learn more about CloudTrail, including how to configure and enable it, see the [AWS CloudTrail User Guide](#).

Amazon Kendra Information in CloudTrail

CloudTrail is enabled on your AWS account when you create the account. When activity occurs in Amazon Kendra, that activity is recorded in a CloudTrail event along with other AWS service events in the CloudTrail **Event history**. You can view, search, and download recent events in your AWS account. For more information, see [Viewing Events with CloudTrail Event History](#).

For an ongoing record of events in your AWS account, including events for Amazon Kendra, create a trail. A *trail* is a configuration that enables CloudTrail to deliver events as log files to a specified S3 bucket. By default, when you create a trail in the console, the trail applies to all AWS Regions. The trail logs events from all Regions in the AWS partition and delivers the log files to the S3 bucket that you specify. Additionally, you can configure other AWS services to further analyze and act upon the event data collected in CloudTrail logs. For more information, see:

- [Overview for Creating a Trail](#)
- [CloudTrail Supported Services and Integrations](#)
- [Configuring Amazon SNS Notifications for CloudTrail](#)
- [Receiving CloudTrail Log Files from Multiple Regions](#) and [Receiving CloudTrail Log Files from Multiple Accounts](#)

CloudTrail logs all Amazon Kendra actions, which are documented in the [API Reference \(p. 125\)](#). For example, calls to the `CreateIndex`, `CreateDataSource`, and `Query` operations generate entries in the CloudTrail log files.

Every event or log entry contains information about who generated the request. For more information, see the [CloudTrail `userIdentity` Element](#).

Example: Amazon Kendra log file Entries

A *trail* is a configuration that enables delivery of events as log files to a specified S3 bucket. CloudTrail log files contain one or more log entries. An *event* represents a single request from any source and includes information about the requested action, the date and time of the action, request parameters, and so on. CloudTrail log files aren't an ordered stack trace of the public API calls, so they don't appear in any specific order.

Calls to the `Query` operation creates the following entry.

```
{
  "eventVersion": "1.05",
  "userIdentity": {
    "type": "AssumedRole | FederatedUser | IAMUser | Root | SAMLUser |  
WebIdentityUser",
    "principalId": "principal ID",
    "arn": "ARN",
    "accountId": "account ID",
    "accessKeyId": "access key ID",
    "sessionContext": {
      "sessionIssuer": {
        "type": "Role",
        "principalId": "principal ID",
        "arn": "ARN",
        "accountId": "account ID",
        "userName": "user name"
      },
      "webIdFederationData": {
      },
    },
  },
}
```

```
        "attributes": {
            "mfaAuthenticated": "false",
            "creationDate": "timestamp"
        }
    },
    "eventTime": "timestamp",
    "eventSource": "kendra.amazonaws.com",
    "eventName": "Query",
    "awsRegion": "region",
    "sourceIPAddress": "source IP address",
    "userAgent": "user agent",
    "requestParameters": {
        "indexId": "index ID"
    },
    "responseElements": null,
    "requestID": "request ID",
    "eventID": "event ID",
    "eventType": "AwsApiCall",
    "recipientAccountId": "account ID"
},
```

Monitoring Amazon Kendra with Amazon CloudWatch

To track the health of your indexes, use Amazon CloudWatch. With CloudWatch, you can get metrics for document synchronization for your index. You can also set up CloudWatch alarms to be notified when one or more metrics exceeds a threshold that you define. For example, you can monitor the number of documents submitted to be indexed or the number of documents that failed to be indexed.

You must have the appropriate CloudWatch permissions to monitor Amazon Kendra with CloudWatch. For more information, see [Authentication and Access Control for Amazon CloudWatch](#) in the *Amazon CloudWatch User Guide*.

Viewing Amazon Kendra metrics

View Amazon Kendra metrics using the CloudWatch console.

To view metrics (CloudWatch console)

1. Sign in to the AWS Management Console and open the CloudWatch console at <https://console.aws.amazon.com/cloudwatch/>.
2. Choose **Metrics**, choose **All Metrics** and then choose **Kendra**.
3. Choose the dimension, choose a metric name, then choose **Add to graph**.
4. Choose a value for the date range. The metric count for the selected date range is displayed in the graph.

Creating an alarm

A CloudWatch alarm watches a single metric over a specified time period and performs one or more actions: sending a notification to an Amazon Simple Notification Service (Amazon SNS) topic or Auto Scaling policy. The actions or actions are based on the value of the metric relative to a given threshold over a number of time periods that you specify. CloudWatch can also send you an Amazon SNS message when the alarm changes state.

CloudWatch alarms invoke actions only when the state changes and has persisted for the period that you specify.

To set an alarm

1. Sign in to the AWS Management Console and open the CloudWatch console at <https://console.aws.amazon.com/cloudwatch/>.
2. Choose **Alarms** and then choose **Create Alarm**.
3. Choose **Kendra metrics** and then choose a metric.
4. For **Time Range**, choose a time range to monitor, and then choose **Next**.
5. Enter a **Name** and **Description**.
6. For **Whenever**, choose **>=**, and type a maximum value.
7. If you want CloudWatch to send an email when the alarm state is reached, in the **Actions** section, for **Whenever this alarm**, choose **State is ALARM**. For **Send notification to**, choose a mailing list or choose **New list** and create a new mailing list.
8. Preview the alarm in the **Alarm Preview** section. If you are satisfied with the alarm, choose **Create Alarm**.

CloudWatch Metrics for index synchronization Jobs

The following table describes the Amazon Kendra metrics for data source synchronization jobs.

Metric	Description
DocumentsCrawled	<p>The number of documents that the synchronization job scanned or discovered during the run.</p> <p>Dimensions:</p> <ul style="list-style-type: none">• IndexId• DataSourceId <p>Unit: Count</p>
DocumentsSubmittedForIndexing	<p>The number of documents that the synchronization job submitted to the index.</p> <p>Dimensions:</p> <ul style="list-style-type: none">• IndexId• DataSourceId <p>Unit: Count</p>
DocumentsSubmittedForIndexingFailed	<p>The number of documents that failed indexing. Check the contents of the CloudWatch log for the synchronization job for details.</p> <p>Dimensions:</p> <ul style="list-style-type: none">• IndexId• DataSourceId

Metric	Description
	Unit: Count
DocumentsSubmittedForDeletion	<p>The number of documents that the synchronization job asked to be removed from the index.</p> <p>Dimensions:</p> <ul style="list-style-type: none">IndexIdDataSourceId <p>Unit: Count</p>
DocumentsSubmittedForDeletionFailed	<p>The number of documents that failed to be deleted. Check the contents of the CloudWatch log for the synchronization job for details.</p> <p>Dimensions:</p> <ul style="list-style-type: none">IndexIdDataSourceId <p>Unit: Count</p>

Metrics for Amazon Kendra data sources

The following table describes the Amazon Kendra metrics for data source synchronization jobs. Metrics marked with an asterisk (*) are used only for Amazon S3 data sources.

Metric	Description
DocumentsSkippedNoChange *	<p>The number of documents examined and found not to have changed so they weren't submitted for indexing.</p> <p>Dimensions:</p> <ul style="list-style-type: none">IndexIdDataSourceId <p>Unit: Count</p>
DocumentsSkippedInvalidMetadata *	<p>The number of documents skipped because there was a problem with the associated metadata file. Check the contents of the CloudWatch log for the synchronization run for details.</p> <p>Dimensions:</p> <ul style="list-style-type: none">IndexIdDataSourceId

Metric	Description
	Unit: Count
DocumentsCrawled	<p>The number of document files examined.</p> <p>Dimensions:</p> <ul style="list-style-type: none">• IndexId• DataSourceId <p>Unit: Count</p>
DocumentsSubmittedForDeletion	<p>The number of documents examined that were deleted from the data source and submitted for deletion.</p> <p>Dimensions:</p> <ul style="list-style-type: none">• IndexId• DataSourceId <p>Unit: Count</p>
DocumentsSubmittedForDeletionFailed	<p>The number of documents that failed deletion from a data source.</p> <p>Dimensions:</p> <ul style="list-style-type: none">• IndexId• DataSourceId <p>Unit: Count</p>
DocumentsSubmittedForIndexing	<p>The number of documents examined and submitted for indexing.</p> <p>Dimensions:</p> <ul style="list-style-type: none">• IndexId• DataSourceId <p>Unit: Count</p>
DocumentsSubmittedForIndexingFailed	<p>The number of documents submitted for indexing that couldn't be indexed.</p> <p>Dimensions:</p> <ul style="list-style-type: none">• IndexId• DataSourceId <p>Unit: Count</p>

Metrics for indexed documents

The following table describes the Amazon Kendra metrics for indexed documents. For documents that are indexed using the [BatchPutDocument \(p. 129\)](#) operation, only the `IndexId` dimension is supported.

Metric	Description
<code>DocumentsIndexed</code>	<p>The number of documents indexed.</p> <p>Dimensions:</p> <ul style="list-style-type: none"><code>IndexId</code><code>DataSourceId</code> <p>Unit: Count</p>
<code>DocumentsFailedToIndex</code>	<p>The number of documents that could not be indexed. Check the contents of the CloudWatch log for details.</p> <p>Dimensions:</p> <ul style="list-style-type: none"><code>IndexId</code><code>DataSourceId</code> <p>Unit: Count</p>

Monitoring Amazon Kendra with Amazon CloudWatch Logs

Amazon Kendra uses Amazon CloudWatch Logs to give you insight into the operation of your data sources. Amazon Kendra logs process details for the documents that as they are indexed. It logs errors from your data source that occur while your documents are being indexed. You use CloudWatch Logs to monitor, store and access the log files.

CloudWatch Logs stores log events in a log stream that is part of a log group. Amazon Kendra uses these features as follows:

- **Log groups** – Amazon Kendra stores all of your log streams in a single log group for each index. Amazon Kendra creates the log group when the index is created. The log group identifier always begins with `aws/kendra/`.
- **Log stream** – creates a new data source log stream in the log group for each index synchronization job that you run. It also creates a new document log stream when a stream reaches approximately 500 entries.
- **Log entries** – Amazon Kendra creates a log entry in the log stream as it indexes documents. Each entry provides information about processing the document or any errors that are encountered.

For more information about using CloudWatch Logs, see [What Is Amazon Cloud Watch Logs](#) in the *Amazon Cloud Watch Logs User Guide*.

Amazon Kendra creates two types of log streams:

- [data source log streams](#) (p. 105)
- [Document log streams](#) (p. 106)

data source log streams

Data source log streams publish entries about your index synchronization jobs. Each synchronization job creates a new log stream that it uses to publish entries. The log stream name is:

```
data source id/YYYY-MM-DD-HH/data source sync job ID
```

A new log stream is created for each synchronization job run.

There are three types of log messages published to a data source log stream:

- A log message for a document that failed to be sent for indexing. The following is an example of this message for a document in an S3 data source:

```
{
  "DocumentId": "document ID",
  "S3Path": "s3://bucket/prefix/object",
  "Message": "Failed to ingest document via BatchPutDocument.",
  "ErrorCode": "InvalidRequest",
  "ErrorMessage": "No document metadata configuration found for document attribute key
city."
}
```

- A log message for a document that failed to be sent for deletion. The following is an example of this message:

```
{
  "DocumentId": "document ID",
  "Message": "Failed to delete document via BatchDeleteDocument.",
  "ErrorCode": "InvalidRequest",
  "ErrorMessage": "Document can't be deleted because it doesn't exist."
}
```

- A log message when an invalid metadata file for a document in an Amazon S3 bucket is found. The following is an example of this message.

```
{
  "Message": "Found invalid metadata
file bucket/prefix/filename.extension.metadata.json."
}
```

- For SharePoint and database connectors, Amazon Kendra only writes messages to the log stream if a document can't be indexed. The following is an example of the error message that Amazon Kendra logs.

```
{
  "DocumentID": "document ID",
  "IndexID": "index ID",
  "SourceURI": "",
  "CrawlStatus": "FAILED",
  "ErrorCode": "403",
  "ErrorMessage": "Access Denied",
  "DataSourceErrorCode": "403"
}
```

```
}
```

Document log streams

Amazon Kendra logs information about processing documents while they are being indexed. It logs a set of messages for documents stored in an Amazon S3 data source. It logs errors only for documents stored in a Microsoft SharePoint or a database data source.

If the documents were added to the index using the [BatchPutDocument \(p. 129\)](#) operation, the log stream is named as follows:

```
YYYY-MM-DD-HH/UUID
```

If the documents were added to the index using a datasource, the log stream is named as follows:

```
dataSourceId/YYYY-MM-DD-HH/UUID
```

Each log stream contains up to 500 messages.

If indexing a document fails, this message is output to the log stream:

```
{
  "DocumentId": "document ID",
  "IndexName": "index name",
  "IndexId": "index ID"
  "SourceURI": "source URI"
  "IndexingStatus": "DocumentFailedToIndex",
  "ErrorCode": "400 | 500",
  "ErrorMessage": "message"
}
```

Security in Amazon Kendra

Cloud security at AWS is the highest priority. As an AWS customer, you benefit from a data center and network architecture that is built to meet the requirements of the most security-sensitive organizations.

Security is a shared responsibility between AWS and you. The [shared responsibility model](#) describes this as security *of* the cloud and security *in* the cloud:

- **Security of the cloud** – AWS is responsible for protecting the infrastructure that runs AWS services in the AWS Cloud. AWS also provides you with services that you can use securely. Third-party auditors regularly test and verify the effectiveness of our security as part of the [AWS Compliance Programs](#). To learn about the compliance programs that apply to Amazon Kendra, see [AWS Services in Scope by Compliance Program](#).
- **Security in the cloud** – Your responsibility is determined by the AWS service that you use. You are also responsible for other factors including the sensitivity of your data, your company's requirements, and applicable laws and regulations.

This documentation helps you understand how to apply the shared responsibility model when using Amazon Kendra. The following topics show you how to configure Amazon Kendra to meet your security and compliance objectives. You also learn how to use other AWS services that help you to monitor and secure your Amazon Kendra resources.

Topics

- [Data protection in Amazon Kendra \(p. 107\)](#)
- [Identity and access management for Amazon Kendra \(p. 109\)](#)
- [Logging and monitoring in Amazon Kendra \(p. 121\)](#)
- [Compliance validation for Amazon Kendra \(p. 121\)](#)
- [Resilience in Amazon Kendra \(p. 122\)](#)
- [Infrastructure security in Amazon Kendra \(p. 122\)](#)

Data protection in Amazon Kendra

Amazon Kendra conforms to the AWS [shared responsibility model](#), which includes regulations and guidelines for data protection. AWS is responsible for protecting the global infrastructure that runs all the AWS services. AWS maintains control over data hosted on this infrastructure, including the security configuration controls for handling customer content and personal data. AWS customers and APN partners, acting either as data controllers or data processors, are responsible for any personal data that they put in the AWS Cloud.

For data protection purposes, we recommend that you protect AWS account credentials and set up individual user accounts with AWS Identity and Access Management (IAM), so that each user is given only the permissions necessary to fulfill their job duties. We also recommend that you secure your data in the following ways:

- Use multi-factor authentication (MFA) with each account.
- Use SSL/TLS to communicate with AWS resources.
- Set up API and user activity logging with AWS CloudTrail.
- Use AWS encryption solutions, along with all default security controls within AWS services.
- Use advanced managed security services such as Amazon Macie, which assists in discovering and securing personal data that is stored in Amazon S3.

We strongly recommend that you never put sensitive identifying information, such as your customers' account numbers, into free-form fields such as a **Name** field. This includes when you work with Amazon Kendra or other AWS services using the console, API, AWS CLI, or AWS SDKs. Any data that you enter into Amazon Kendra or other services might get picked up for inclusion in diagnostic logs. When you provide a URL to an external server, don't include credentials information in the URL to validate your request to that server.

For more information about data protection, see the [AWS Shared Responsibility Model and GDPR](#) blog post on the *AWS Security Blog*.

Encryption at rest

Amazon Kendra encrypts your data at rest with your choice of an encryption key. You can choose one of the following:

- An AWS owned customer master key (CMK). If you don't specify an encryption key your data is encrypted with this key by default.
- An AWS managed CMK in your account. This key is created, managed, and used on your behalf by Amazon Kendra. The key name is `aws/kendra`.
- A customer managed CMK. You can provide the ARN of an encryption key that you created in your account. When you use a customer managed CMK, you must give the key a key policy that enables Amazon Kendra to use the key. Select a symmetric customer managed CMK, Amazon Kendra does not support asymmetric CMKs. For more information, see [Key management \(p. 108\)](#).

Encryption in transit

Amazon Kendra uses the HTTPS protocol to communicate with your client application. It uses HTTPS and AWS signatures to communicate with other services on your application's behalf.

Key management

Amazon Kendra encrypts the contents of your index using one of three types of keys. You can choose one of the following:

- An AWS owned customer master key (CMK). This is the default.
- An AWS managed CMK. This key is created in your account and is managed and used on your behalf by Amazon Kendra.
- A customer managed CMK. You can create the key when you are creating an Amazon Kendra index or data source, or you can create the key using the AWS KMS console. Select a symmetric customer managed CMK, Amazon Kendra does not support asymmetric CMKs. For more information, see [Using Symmetric and Asymmetric Keys](#) in the *AWS Key Management Service Developer Guide*.

When you create a key using the AWS KMS console, you must give the key the following policy that enables Amazon Kendra to use the key. If you create a key with the Amazon Kendra console, the policy is applied to the key for you. For more information, see [Using Key Policies in AWS KMS](#) in the *AWS Key Management Service Developer Guide*.

```
{
  "Effect": "Allow",
  "Principal": {
    "Service": "kendra.amazonaws.com"
  },
  "Action": [
    "kms:Encrypt",
```

```
    "kms:Decrypt",  
    "kms:ReEncrypt*",  
    "kms:GenerateDataKey*",  
    "kms:DescribeKey",  
    "kms:CreateGrant",  
    "kms:RetireGrant"  
  ],  
  "Resource": "*" }  
}
```

Identity and access management for Amazon Kendra

AWS Identity and Access Management (IAM) is an AWS service that helps an administrator securely control access to AWS resources. IAM administrators control who can be *authenticated* (signed in) and *authorized* (have permissions) to use Amazon Kendra resources. IAM is an AWS service that you can use with no additional charge.

Topics

- [Audience \(p. 109\)](#)
- [Authenticating with identities \(p. 109\)](#)
- [Managing access using policies \(p. 111\)](#)
- [How Amazon Kendra works with IAM \(p. 113\)](#)
- [Amazon Kendra Identity-based policy examples \(p. 116\)](#)

Audience

How you use AWS Identity and Access Management (IAM) differs, depending on the work you do in Amazon Kendra.

Service user – If you use the Amazon Kendra service to do your job, then your administrator provides you with the credentials and permissions that you need. As you use more Amazon Kendra features to do your work, you might need additional permissions. Understanding how access is managed can help you request the right permissions from your administrator. If you cannot access a feature in Amazon Kendra, see [Troubleshooting Amazon Kendra Identity and Access \(p. 119\)](#).

Service administrator – If you're in charge of Amazon Kendra resources at your company, you probably have full access to Amazon Kendra. It's your job to determine which Amazon Kendra features and resources your employees should access. You must then submit requests to your IAM administrator to change the permissions of your service users. Review the information on this page to understand the basic concepts of IAM. To learn more about how your company can use IAM with Amazon Kendra, see [How Amazon Kendra works with IAM \(p. 113\)](#).

IAM administrator – If you're an IAM administrator, you might want to learn details about how you can write policies to manage access to Amazon Kendra. To view example Amazon Kendra identity-based policies that you can use in IAM, see [Amazon Kendra Identity-based policy examples \(p. 116\)](#).

Authenticating with identities

Authentication is how you sign in to AWS using your identity credentials. For more information about signing in using the AWS Management Console, see [The IAM Console and Sign-in Page](#) in the *IAM User Guide*.

You must be *authenticated* (signed in to AWS) as the AWS account root user, an IAM user, or by assuming an IAM role. You can also use your company's single sign-on authentication, or even sign in using Google or Facebook. In these cases, your administrator previously set up identity federation using IAM roles. When you access AWS using credentials from another company, you are assuming a role indirectly.

To sign in directly to the [AWS Management Console](#), use your password with your root user email or your IAM user name. You can access AWS programmatically using your root user or IAM user access keys. AWS provides SDK and command line tools to cryptographically sign your request using your credentials. If you don't use AWS tools, you must sign the request yourself. Do this using *Signature Version 4*, a protocol for authenticating inbound API requests. For more information about authenticating requests, see [Signature Version 4 Signing Process](#) in the *AWS General Reference*.

Regardless of the authentication method that you use, you might also be required to provide additional security information. For example, AWS recommends that you use multi-factor authentication (MFA) to increase the security of your account. To learn more, see [Using Multi-Factor Authentication \(MFA\) in AWS](#) in the *IAM User Guide*.

AWS Account Root User

When you first create an AWS account, you begin with a single sign-in identity that has complete access to all AWS services and resources in the account. This identity is called the AWS account *root user* and is accessed by signing in with the email address and password that you used to create the account. We strongly recommend that you do not use the root user for your everyday tasks, even the administrative ones. Instead, adhere to the [best practice of using the root user only to create your first IAM user](#). Then securely lock away the root user credentials and use them to perform only a few account and service management tasks.

IAM Users and groups

An *IAM user* is an identity within your AWS account that has specific permissions for a single person or application. An IAM user can have long-term credentials such as a user name and password or a set of access keys. To learn how to generate access keys, see [Managing Access Keys for IAM Users](#) in the *IAM User Guide*. When you generate access keys for an IAM user, make sure you view and securely save the key pair. You cannot recover the secret access key in the future. Instead, you must generate a new access key pair.

An *IAM group* is an identity that specifies a collection of IAM users. You can't sign in as a group. You can use groups to specify permissions for multiple users at a time. Groups make permissions easier to manage for large sets of users. For example, you could have a group named *IAMAdmins* and give that group permissions to administer IAM resources.

Users are different from roles. A user is uniquely associated with one person or application, but a role is intended to be assumable by anyone who needs it. Users have permanent long-term credentials, but roles provide temporary credentials. To learn more, see [When to Create an IAM User \(Instead of a Role\)](#) in the *IAM User Guide*.

IAM roles

An *IAM role* is an identity within your AWS account that has specific permissions. It is similar to an IAM user, but is not associated with a specific person. You can temporarily assume an IAM role in the AWS Management Console by [switching roles](#). You can assume a role by calling an AWS CLI or AWS API operation or by using a custom URL. For more information about methods for using roles, see [Using IAM Roles](#) in the *IAM User Guide*.

IAM roles with temporary credentials are useful in the following situations:

- **Temporary IAM user permissions** – An IAM user can assume an IAM role to temporarily take on different permissions for a specific task.

- **Federated user access** – Instead of creating an IAM user, you can use existing identities from AWS Directory Service, your enterprise user directory, or a web identity provider. These are known as *federated users*. AWS assigns a role to a federated user when access is requested through an [identity provider](#). For more information about federated users, see [Federated Users and Roles](#) in the *IAM User Guide*.
- **Cross-account access** – You can use an IAM role to allow someone (a trusted principal) in a different account to access resources in your account. Roles are the primary way to grant cross-account access. However, with some AWS services, you can attach a policy directly to a resource (instead of using a role as a proxy). To learn the difference between roles and resource-based policies for cross-account access, see [How IAM Roles Differ from Resource-based Policies](#) in the *IAM User Guide*.
- **AWS service access** – A service role is an IAM role that a service assumes to perform actions in your account on your behalf. When you set up some AWS service environments, you must define a role for the service to assume. This service role must include all the permissions that are required for the service to access the AWS resources that it needs. Service roles vary from service to service, but many allow you to choose your permissions as long as you meet the documented requirements for that service. Service roles provide access only within your account and cannot be used to grant access to services in other accounts. You can create, modify, and delete a service role from within IAM. For example, you can create a role that allows Amazon Redshift to access an Amazon S3 bucket on your behalf and then load data from that bucket into an Amazon Redshift cluster. For more information, see [Creating a Role to Delegate Permissions to an AWS Service](#) in the *IAM User Guide*.
- **Applications running on Amazon EC2** – You can use an IAM role to manage temporary credentials for applications that are running on an EC2 instance and making AWS CLI or AWS API requests. This is preferable to storing access keys within the EC2 instance. To assign an AWS role to an EC2 instance and make it available to all of its applications, you create an instance profile that is attached to the instance. An instance profile contains the role and enables programs that are running on the EC2 instance to get temporary credentials. For more information, see [Using an IAM Role to Grant Permissions to Applications Running on Amazon EC2 Instances](#) in the *IAM User Guide*.

To learn whether to use IAM roles, see [When to Create an IAM Role \(Instead of a User\)](#) in the *IAM User Guide*.

Managing access using policies

You control access in AWS by creating policies and attaching them to IAM identities or AWS resources. A policy is an object in AWS that, when associated with an identity or resource, defines their permissions. AWS evaluates these policies when an entity (root user, IAM user, or IAM role) makes a request. Permissions in the policies determine whether the request is allowed or denied. Most policies are stored in AWS as JSON documents. For more information about the structure and contents of JSON policy documents, see [Overview of JSON Policies](#) in the *IAM User Guide*.

An IAM administrator can use policies to specify who has access to AWS resources, and what actions they can perform on those resources. Every IAM entity (user or role) starts with no permissions. In other words, by default, users can do nothing, not even change their own password. To give a user permission to do something, an administrator must attach a permissions policy to a user. Or the administrator can add the user to a group that has the intended permissions. When an administrator gives permissions to a group, all users in that group are granted those permissions.

IAM policies define permissions for an action regardless of the method that you use to perform the operation. For example, suppose that you have a policy that allows the `iam:GetRole` action. A user with that policy can get role information from the AWS Management Console, the AWS CLI, or the AWS API.

Identity-based policies

Identity-based policies are JSON permissions policy documents that you can attach to an identity, such as an IAM user, role, or group. These policies control what actions that identity can perform, on which

resources, and under what conditions. To learn how to create an identity-based policy, see [Creating IAM Policies](#) in the *IAM User Guide*.

Identity-based policies can be further categorized as *inline policies* or *managed policies*. Inline policies are embedded directly into a single user, group, or role. Managed policies are standalone policies that you can attach to multiple users, groups, and roles in your AWS account. Managed policies include AWS managed policies and customer managed policies. To learn how to choose between a managed policy or an inline policy, see [Choosing Between Managed Policies and Inline Policies](#) in the *IAM User Guide*.

Resource-based policies

Resource-based policies are JSON policy documents that you attach to a resource such as an Amazon S3 bucket. Service administrators can use these policies to define what actions a specified principal (account member, user, or role) can perform on that resource and under what conditions. Resource-based policies are inline policies. There are no managed resource-based policies.

Access control lists (ACLs)

Access control lists (ACLs) are a type of policy that controls which principals (account members, users, or roles) have permissions to access a resource. ACLs are similar to resource-based policies, although they do not use the JSON policy document format. Amazon S3, AWS WAF, and Amazon VPC are examples of services that support ACLs. To learn more about ACLs, see [Access Control List \(ACL\) Overview](#) in the *Amazon Simple Storage Service Developer Guide*.

Other policy types

AWS supports additional, less-common policy types. These policy types can set the maximum permissions granted to you by the more common policy types.

- **Permissions boundaries** – A permissions boundary is an advanced feature in which you set the maximum permissions that an identity-based policy can grant to an IAM entity (IAM user or role). You can set a permissions boundary for an entity. The resulting permissions are the intersection of entity's identity-based policies and its permissions boundaries. Resource-based policies that specify the user or role in the `Principal` field are not limited by the permissions boundary. An explicit deny in any of these policies overrides the allow. For more information about permissions boundaries, see [Permissions Boundaries for IAM Entities](#) in the *IAM User Guide*.
- **Service control policies (SCPs)** – SCPs are JSON policies that specify the maximum permissions for an organization or organizational unit (OU) in AWS Organizations. AWS Organizations is a service for grouping and centrally managing multiple AWS accounts that your business owns. If you enable all features in an organization, then you can apply service control policies (SCPs) to any or all of your accounts. The SCP limits permissions for entities in member accounts, including each AWS account root user. For more information about Organizations and SCPs, see [How SCPs Work](#) in the *AWS Organizations User Guide*.
- **Session policies** – Session policies are advanced policies that you pass as a parameter when you programmatically create a temporary session for a role or federated user. The resulting session's permissions are the intersection of the user or role's identity-based policies and the session policies. Permissions can also come from a resource-based policy. An explicit deny in any of these policies overrides the allow. For more information, see [Session Policies](#) in the *IAM User Guide*.

Multiple policy types

When multiple types of policies apply to a request, the resulting permissions are more complicated to understand. To learn how AWS determines whether to allow a request when multiple policy types are involved, see [Policy Evaluation Logic](#) in the *IAM User Guide*.

How Amazon Kendra works with IAM

Before you use IAM to manage access to Amazon Kendra, you should understand what IAM features are available to use with Amazon Kendra. To get a high-level view of how Amazon Kendra and other AWS services work with IAM, see [AWS Services That Work with IAM](#) in the *IAM User Guide*.

Topics

- [Amazon Kendra identity-based policies](#) (p. 113)
- [Amazon Kendra Resource-based policies](#) (p. 114)
- [Access control lists \(ACLs\)](#) (p. 114)
- [Authorization based on Amazon Kendra tags](#) (p. 115)
- [Amazon Kendra IAM Roles](#) (p. 115)

Amazon Kendra identity-based policies

With IAM identity-based policies, you can specify allowed or denied actions and resources as well as the conditions under which actions are allowed or denied. Amazon Kendra supports specific actions, resources, and condition keys. To learn about all of the elements that you use in a JSON policy, see [IAM JSON Policy Elements Reference](#) in the *IAM User Guide*.

Actions

The `Action` element of an IAM identity-based policy describes the specific action or actions that will be allowed or denied by the policy. Policy actions usually have the same name as the associated AWS API operation. The action is used in a policy to grant permissions to perform the associated operation.

Policy actions in Amazon Kendra use the following prefix before the action: `kendra:`. For example, to grant someone permission to list Amazon Kendra indexes with the [ListIndices](#) (p. 175) API operation, you include the `kendra:ListIndices` action in their policy. Policy statements must include either an `Action` or `NotAction` element. Amazon Kendra defines its own set of actions that describe tasks that you can perform with this service.

To specify multiple actions in a single statement, separate them with commas as follows:

```
"Action": [
    "kendra:action1",
    "kendra:action2"
```

You can specify multiple actions using wildcards (*). For example, to specify all actions that begin with the word `Describe`, include the following action:

```
"Action": "kendra:Describe*"
```

To see a list of Amazon Kendra actions, see [Actions Defined by Amazon Kendra](#) in the *IAM User Guide*.

Resources

The `Resource` element specifies the object or objects to which the action applies. Statements must include either a `Resource` or a `NotResource` element. You specify a resource using an ARN or using the wildcard (*) to indicate that the statement applies to all resources.

The Amazon Kendra index resource has the following ARN:

```
arn:${Partition}:kendra:${Region}:${Account}:index/${IndexId}
```

For more information about the format of ARNs, see [Amazon Resource Names \(ARNs\) and AWS Service Namespaces](#).

For example, to specify an index in your statement, use the GUID of the index in the following ARN:

```
"Resource": "arn:aws:kendra:${Region}:${Account}:index/${GUID}"
```

To specify all indexes that belong to a specific account, use the wildcard (*):

```
"Resource": "arn:aws:${Region}:${Account}:index/*"
```

Some Amazon Kendra actions, such as those for creating resources, cannot be performed on a specific resource. In those cases, you must use the wildcard (*).

```
"Resource": "*" 
```

To see a list of Amazon Kendra resource types and their ARNs, see [Resources Defined by Amazon Kendra](#) in the *IAM User Guide*. To learn with which actions you can specify the ARN of each resource, see [Actions Defined by Amazon Kendra](#).

Condition keys

The `Condition` element (or *Condition block*) lets you specify conditions in which a statement is in effect. The `Condition` element is optional. You can build conditional expressions that use [condition operators](#), such as equals or less than, to match the condition in the policy with values in the request.

If you specify multiple `Condition` elements in a statement, or multiple keys in a single `Condition` element, AWS evaluates them using a logical AND operation. If you specify multiple values for a single condition key, AWS evaluates the condition using a logical OR operation. All of the conditions must be met before the statement's permissions are granted.

You can also use placeholder variables when you specify conditions. For example, you can grant an IAM user permission to access a resource only if it is tagged with their IAM user name. For more information, see [IAM Policy Elements: Variables and Tags](#) in the *IAM User Guide*.

Amazon Kendra does not provide any service-specific condition keys, but it does support using some global condition keys. To see all AWS global condition keys, see [AWS Global Condition Context Keys](#) in the *IAM User Guide*.

Examples

To view examples of Amazon Kendra identity-based policies, see [Amazon Kendra Identity-based policy examples \(p. 116\)](#).

Amazon Kendra Resource-based policies

Amazon Kendra does not support resource-based policies.

Access control lists (ACLs)

Amazon Kendra does not support access control lists (ACLs) for access to AWS services and resources.

Authorization based on Amazon Kendra tags

You can associate tags with certain types of Amazon Kendra resources to authorize access to those resources. To control access based on tags, provide tag information in the condition element of a policy by using the `aws:RequestTag/key-name`, or `aws:TagKeys` condition keys.

The following table lists the actions, corresponding resource types, and condition keys for tag-based access control. Each action is authorized based on the tags associated with the corresponding resource type.

Action	Resource type	Condition keys
CreateDataSource (p. 132)		<code>aws:RequestTag</code> , <code>aws:TagKeys</code>
CreateFaq (p. 139)		<code>aws:RequestTag</code> , <code>aws:TagKeys</code>
CreateIndex (p. 142)		<code>aws:RequestTag</code> , <code>aws:TagKeys</code>
ListTagsForResource (p. 177)	data source, FAQ, index	
TagResource (p. 191)	data source, FAQ, index	<code>aws:RequestTag</code> , <code>aws:TagKeys</code>
UntagResource (p. 193)	data source, FAQ, index	<code>aws:TagKeys</code>

For information about tagging Amazon Kendra resources, see [Tags \(p. 7\)](#). For an example identity-based policy that limits access to a resource based on resource tags, see [Tag-based policy examples \(p. 118\)](#). For more information about using tags to limit access to resources, see [Controlling access using tags](#) in the *IAM User Guide*.

Amazon Kendra IAM Roles

An [IAM role](#) is an entity within your AWS account that has specific permissions.

Using temporary credentials with Amazon Kendra

You can use temporary credentials to sign in with federation, assume an IAM role, or to assume a cross-account role. You obtain temporary security credentials by calling AWS STS API operations such as [AssumeRole](#) or [GetFederationToken](#).

Amazon Kendra supports using temporary credentials.

Service roles

This feature allows a service to assume a [service role](#) on your behalf. This role allows the service to access resources in other services to complete an action on your behalf. Service roles appear in your IAM account and are owned by the account. This means that an IAM administrator can change the permissions for this role. However, doing so might break the functionality of the service.

Amazon Kendra supports service roles.

Choosing an IAM role in Amazon Kendra

When you create an index, call the `BatchPutDocument` operation, create a data source or create an FAQ, you must provide an access role Amazon Resource Name (ARN) that Amazon Kendra uses to access the

required resources on your behalf. If you have previously created a role, then the Amazon Kendra console provides you with a list of roles to choose from. It's important to choose a role that allows access to the resources that you require. For more information, see [IAM access roles for Amazon Kendra \(p. 27\)](#).

Amazon Kendra Identity-based policy examples

By default, IAM users and roles don't have permission to create or modify Amazon Kendra resources. They also can't perform tasks using the AWS Management Console, AWS CLI, or AWS API. An IAM administrator must create IAM policies that grant users and roles permission to perform specific API operations on the specified resources they need. The administrator must then attach those policies to the IAM users or groups that require those permissions.

To learn how to create an IAM identity-based policy using these example JSON policy documents, see [Creating Policies on the JSON Tab](#) in the *IAM User Guide*.

Topics

- [Policy best practices \(p. 116\)](#)
- [AWS Managed \(Predefined\) Policies for Amazon Kendra \(p. 116\)](#)
- [Allow users to view their own permissions \(p. 117\)](#)
- [Accessing one Amazon Kendra index \(p. 117\)](#)
- [Tag-based policy examples \(p. 118\)](#)
- [Troubleshooting Amazon Kendra Identity and Access \(p. 119\)](#)

Policy best practices

Identity-based policies are very powerful. They determine whether someone can create, access, or delete Amazon Kendra resources in your account. These actions can incur costs for your AWS account. When you create or edit identity-based policies, follow these guidelines and recommendations:

- **Get Started Using AWS Managed Policies** – To start using Amazon Kendra quickly, use AWS managed policies to give your employees the permissions they need. These policies are already available in your account and are maintained and updated by AWS. For more information, see [Get Started Using Permissions With AWS Managed Policies](#) in the *IAM User Guide*.
- **Grant Least Privilege** – When you create custom policies, grant only the permissions required to perform a task. Start with a minimum set of permissions and grant additional permissions as necessary. Doing so is more secure than starting with permissions that are too lenient and then trying to tighten them later. For more information, see [Grant Least Privilege](#) in the *IAM User Guide*.
- **Enable MFA for Sensitive Operations** – For extra security, require IAM users to use multi-factor authentication (MFA) to access sensitive resources or API operations. For more information, see [Using Multi-Factor Authentication \(MFA\) in AWS](#) in the *IAM User Guide*.
- **Use Policy Conditions for Extra Security** – To the extent that it's practical, define the conditions under which your identity-based policies allow access to a resource. For example, you can write conditions to specify a range of allowable IP addresses that a request must come from. You can also write conditions to allow requests only within a specified date or time range, or to require the use of SSL or MFA. For more information, see [IAM JSON Policy Elements: Condition](#) in the *IAM User Guide*.

AWS Managed (Predefined) Policies for Amazon Kendra

AWS addresses many common use cases by providing standalone IAM policies that are created and administered by AWS. These policies are called AWS managed policies. AWS managed policies make it easier for you to assign permissions to users, groups, and roles than if you had to write the policies yourself. For more information, see [Adding Permissions to a User](#) in the *IAM User Guide*.

The following AWS managed policies, which you can attach to groups and roles in your account, are specific to Amazon Kendra:

- **AmazonKendraReadOnly** — Grants read-only access to Amazon Kendra resources.
- **AmazonKendraFullAccess** — Grants full access to create, read, update, delete, tag, and run all Amazon Kendra resources.

For the console, your role must also have `iam:CreateRole`, `iam:CreatePolicy`, `iam:AttachRolePolicy`, and `s3:ListBucket` permissions.

Note

You can review these permissions by signing in to the IAM console and searching for specific policies.

You can also create your own custom policies to allow permissions for Amazon Kendra API actions. You can attach these custom policies to the IAM roles or groups that require those permissions. For examples of IAM policies for Amazon Kendra, see [Amazon Kendra Identity-based policy examples \(p. 116\)](#).

Allow users to view their own permissions

This example shows how you might create a policy that allows IAM users to view the inline and managed policies that are attached to their user identity. This policy includes permissions to complete this action on the console or programmatically using the AWS CLI or AWS API.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Sid": "ViewOwnUserInfo",
      "Effect": "Allow",
      "Action": [
        "iam:GetUserPolicy",
        "iam:ListGroupsWithUser",
        "iam:ListAttachedUserPolicies",
        "iam:ListUserPolicies",
        "iam:GetUser"
      ],
      "Resource": ["arn:aws:iam::*:user/${aws:username}"]
    },
    {
      "Sid": "NavigateInConsole",
      "Effect": "Allow",
      "Action": [
        "iam:GetGroupPolicy",
        "iam:GetPolicyVersion",
        "iam:GetPolicy",
        "iam:ListAttachedGroupPolicies",
        "iam:ListGroupPolicies",
        "iam:ListPolicyVersions",
        "iam:ListPolicies",
        "iam:ListUsers"
      ],
      "Resource": "*"
    }
  ]
}
```

Accessing one Amazon Kendra index

In this example, you want to grant an IAM user in your AWS account access to query an index.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Sid": "QueryIndex",
      "Effect": "Allow",
      "Action": [
        "kendra:Query"
      ],
      "Resource": "arn:aws:kendra:${Region}:${Account}:index/${Index ID}"
    }
  ]
}
```

Tag-based policy examples

Tag-based policies are JSON policy documents that specify the actions that a principal can perform on tagged resources.

Example: Use a tag to access a resource

This example policy grants an IAM user or role in your AWS account permission to use the `Query` operation with any resource tagged with the key **department** and the value **finance**.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": [
        "kendra:Query"
      ],
      "Resource": "*",
      "Condition": {
        "StringEquals": {
          "aws:ResourceTag/department": "finance"
        }
      }
    }
  ]
}
```

Example: Use a tag to enable Amazon Kendra operations

This example policy grants an IAM user or role in your AWS account permission to use any Amazon Kendra operation except `TagResource` operation with any resource tagged with the key **department** and the value **finance**.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": "kendra:*",
      "Resource": "*"
    },
    {
      "Effect": "Deny",
      "Action": [
        "kendra:TagResource"
      ],
      "Resource": "*"
    }
  ]
}
```

```

        "Resource": "*",
        "Condition": {
            "StringEquals": {
                "aws:ResourceTag/department": "finance"
            }
        }
    }
]
}

```

Example: Use a tag to restrict access to an operation

This example policy restricts access for an IAM user or role in your AWS account to use the `CreateIndex` operation unless the user provides the **department** tag and it has the allowed values **finance** and **IT**.

```

{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": "kendra:CreateIndex",
      "Resource": "*"
    },
    {
      "Effect": "Deny",
      "Action": "kendra:CreateIndex",
      "Resource": "*",
      "Condition": {
        "Null": {
          "aws:RequestTag/department": "true"
        }
      }
    },
    {
      "Effect": "Deny",
      "Action": "kendra:CreateIndex",
      "Resource": "*",
      "Condition": {
        "ForAnyValue:StringNotEquals": {
          "aws:RequestTag/department": [
            "finance",
            "IT"
          ]
        }
      }
    }
  ]
}

```

Troubleshooting Amazon Kendra Identity and Access

Use the following information to help you diagnose and fix common issues that you might encounter when working with Amazon Kendra and IAM.

Topics

- [I am not authorized to perform an action in Amazon Kendra \(p. 120\)](#)
- [I am not authorized to perform iam:PassRole \(p. 120\)](#)
- [I want to view my access keys \(p. 120\)](#)
- [I'm an administrator and I want to allow others to access Amazon Kendra \(p. 120\)](#)
- [I want to allow people outside of my AWS account to access my Amazon Kendra resources \(p. 121\)](#)

I am not authorized to perform an action in Amazon Kendra

If the AWS Management Console tells you that you're not authorized to perform an action, then you must contact your administrator for assistance. Your administrator is the person that provided you with your user name and password.

The following example error occurs when the `mateojackson` IAM user tries to use the console to view details about an index but does not have `kendra:DescribeIndex` permissions.

```
User: arn:aws:iam::123456789012:user/mateojackson is not authorized to perform:
kendra:DescribeIndex on resource: index ARN
```

In this case, Mateo asks his administrator to update his policies to allow him to access the index resource using the `kendra:DescribeIndex` action.

I am not authorized to perform iam:PassRole

If you receive an error that you're not authorized to perform the `iam:PassRole` action, then you must contact your administrator for assistance. Your administrator is the person that provided you with your user name and password. Ask that person to update your policies to allow you to pass a role to Amazon Kendra.

Some AWS services allow you to pass an existing role to that service, instead of creating a new service role or service-linked role. To do this, you must have permissions to pass the role to the service.

The following example error occurs when an IAM user named `marymajor` tries to use the console to perform an action in Amazon Kendra. However, the action requires the service to have permissions granted by a service role. Mary does not have permissions to pass the role to the service.

```
User: arn:aws:iam::123456789012:user/marymajor is not authorized to perform: iam:PassRole
```

In this case, Mary asks her administrator to update her policies to allow her to perform the `iam:PassRole` action.

I want to view my access keys

After you create your IAM user access keys, you can view your access key ID at any time. However, you can't view your secret access key again. If you lose your secret key, you must create a new access key pair.

Access keys consist of two parts: an access key ID (for example, `AKIAIOSFODNN7EXAMPLE`) and a secret access key (for example, `wJalrXUtnFEMI/K7MDENG/bPxrFiCYEXAMPLEKEY`). Like a user name and password, you must use both the access key ID and secret access key together to authenticate your requests. Manage your access keys as securely as you do your user name and password.

Important

Do not provide your access keys to a third party, even to help [find your canonical user ID](#). By doing this, you might give someone permanent access to your account.

When you create an access key pair, you are prompted to save the access key ID and secret access key in a secure location. The secret access key is available only at the time you create it. If you lose your secret access key, you must add new access keys to your IAM user. You can have a maximum of two access keys. If you already have two, you must delete one key pair before creating a new one. To view instructions, see [Managing Access Keys](#) in the *IAM User Guide*.

I'm an administrator and I want to allow others to access Amazon Kendra

To allow others to access Amazon Kendra, you must create an IAM entity (user or role) for the person or application that needs access. They will use the credentials for that entity to access AWS. You must then attach a policy to the entity that grants them the correct permissions in Amazon Kendra.

To get started right away, see [Creating Your First IAM Delegated User and Group](#) in the *IAM User Guide*.

I want to allow people outside of my AWS account to access my Amazon Kendra resources

You can create a role that users in other accounts or people outside of your organization can use to access your resources. You can specify who is trusted to assume the role. For services that support resource-based policies or access control lists (ACLs), you can use those policies to grant people access to your resources.

To learn more, consult the following:

- To learn whether Amazon Kendra supports these features, see [How Amazon Kendra works with IAM](#) (p. 113).
- To learn how to provide access to your resources across AWS accounts that you own, see [Providing Access to an IAM User in Another AWS Account That You Own](#) in the *IAM User Guide*.
- To learn how to provide access to your resources to third-party AWS accounts, see [Providing Access to AWS Accounts Owned by Third Parties](#) in the *IAM User Guide*.
- To learn how to provide access through identity federation, see [Providing Access to Externally Authenticated Users \(Identity Federation\)](#) in the *IAM User Guide*.
- To learn the difference between using roles and resource-based policies for cross-account access, see [How IAM Roles Differ from Resource-based Policies](#) in the *IAM User Guide*.

Logging and monitoring in Amazon Kendra

Monitoring is an important part of maintaining the reliability, availability, and performance of your Amazon Kendra applications. To monitor Amazon Kendra API calls, you can use AWS CloudTrail. To monitor the status of your jobs, use Amazon CloudWatch Logs.

- **Amazon CloudWatch Alarms** — Using CloudWatch alarms, you watch a single metric over a time period that you specify. If the metric exceeds a given threshold, a notification is sent to an Amazon SNS topic or AWS Auto Scaling policy. CloudWatch alarms do not invoke actions when a metric is in a particular state. Rather the state must have changed and been maintained for a specified number of periods. For more information, see [Monitoring Amazon Kendra with Amazon CloudWatch](#) (p. 100).
- **AWS CloudTrail Logs** — CloudTrail provides a record of actions taken by a user, role, or an AWS service in Amazon Kendra. Using the information collected by CloudTrail, you can determine the request that was made to Amazon Kendra, the IP address from which the request was made, who made the request, when it was made, and additional details. For more information, see [Logging Amazon Kendra API calls with AWS CloudTrail logs](#) (p. 98).

Compliance validation for Amazon Kendra

Third-party auditors assess the security and compliance of Amazon Kendra as part of multiple AWS compliance programs. Amazon Kendra is not in scope of any AWS compliance programs.

For a list of AWS services in scope of specific compliance programs, see [AWS Services in Scope by Compliance Program](#). For general information, see [AWS Compliance Programs](#).

You can download third-party audit reports using AWS Artifact. For more information, see [Downloading Reports in AWS Artifact](#).

Your compliance responsibility when using Amazon Kendra is determined by the sensitivity of your data, your company's compliance objectives, and applicable laws and regulations. AWS provides the following resources to help with compliance:

- [Security and Compliance Quick Start Guides](#) – These deployment guides discuss architectural considerations and provide steps for deploying security- and compliance-focused baseline environments on AWS.
- [AWS Compliance Resources](#) – This collection of workbooks and guides might apply to your industry and location.
- [Evaluating Resources with Rules](#) in the *AWS Config Developer Guide* – The AWS Config service assesses how well your resource configurations comply with internal practices, industry guidelines, and regulations.
- [AWS Security Hub](#) – This AWS service provides a comprehensive view of your security state within AWS that helps you check your compliance with security industry standards and best practices.

Resilience in Amazon Kendra

The AWS global infrastructure is built around AWS Regions and Availability Zones. AWS Regions provide multiple physically separated and isolated Availability Zones, which are connected with low-latency, high-throughput, and highly redundant networking. With Availability Zones, you can design and operate applications and databases that automatically fail over between zones without interruption. Availability Zones are more highly available, fault tolerant, and scalable than traditional single or multiple data center infrastructures.

For more information about AWS Regions and Availability Zones, see [AWS Global Infrastructure](#).

In addition to the AWS global infrastructure, Amazon Kendra offers several features to help support your data resiliency and backup needs.

Infrastructure security in Amazon Kendra

As a managed service, Amazon Kendra is protected by the AWS global network security procedures that are described in the [Amazon Web Services: Overview of Security Processes](#) whitepaper.

You use AWS published API calls to access Amazon Kendra through the network. Clients must support Transport Layer Security (TLS) 1.0 or later. We recommend TLS 1.2 or later. Clients must also support cipher suites with perfect forward secrecy (PFS) such as Ephemeral Diffie-Hellman (DHE) or Elliptic Curve Ephemeral Diffie-Hellman (ECDHE). Most modern systems such as Java 7 and later support these modes.

Additionally, requests must be signed by using an access key ID and a secret access key that is associated with an IAM principal. Or you can use the [AWS Security Token Service](#) (AWS STS) to generate temporary security credentials to sign requests.

Quotas for Amazon Kendra

Supported regions

For a list of AWS Regions where Amazon Kendra is available, see [AWS Regions and Endpoints](#) in the *Amazon Web Services General Reference*.

Quotas

Service quotas, also referred to as limits, are the maximum number of service resources for your AWS account. For more information, see [AWS Service Quotas](#) in the *AWS General Reference*.

Description	Default	Adjustable
Maximum number of indexes per account	5	Yes
Maximum number of data sources per enterprise edition index	50	Yes
Maximum number of storage capacity units per enterprise edition index	20	Yes
Maximum number of query capacity units per enterprise edition index	20	Yes
Maximum number of data sources per developer edition index	5	No
Maximum size of a single document	50 MB	Yes
Maximum amount of text extracted from a document	5 MB	Yes
Maximum user group list size per query attribute	10	Yes
Maximum string list size per query attribute	10	Yes
Maximum number of FAQs per index	30	Yes

For more information about Amazon Kendra service quotas and to request a quota increase, see [Service Quotas](#).

Document history for Amazon Kendra

- **Latest documentation update:** May 11, 2020

The following table describes important changes in each release of Amazon Kendra. For notification about updates to this documentation, you can subscribe to an RSS feed.

update-history-change	update-history-description	update-history-date
New guide	This is the first release of the <i>Amazon Kendra Developer Guide</i> .	May 11, 2020

API Reference

This section contains the API Reference documentation.

Actions

The following actions are supported:

- [BatchDeleteDocument](#) (p. 126)
- [BatchPutDocument](#) (p. 129)
- [CreateDataSource](#) (p. 132)
- [CreateFaq](#) (p. 139)
- [CreateIndex](#) (p. 142)
- [DeleteDataSource](#) (p. 145)
- [DeleteFaq](#) (p. 147)
- [DeleteIndex](#) (p. 149)
- [DescribeDataSource](#) (p. 151)
- [DescribeFaq](#) (p. 158)
- [DescribeIndex](#) (p. 162)
- [ListDataSources](#) (p. 166)
- [ListDataSourceSyncJobs](#) (p. 169)
- [ListFaqs](#) (p. 172)
- [ListIndices](#) (p. 175)
- [ListTagsForResource](#) (p. 177)
- [Query](#) (p. 179)
- [StartDataSourceSyncJob](#) (p. 185)
- [StopDataSourceSyncJob](#) (p. 187)
- [SubmitFeedback](#) (p. 189)
- [TagResource](#) (p. 191)
- [UntagResource](#) (p. 193)
- [UpdateDataSource](#) (p. 195)
- [UpdateIndex](#) (p. 201)

BatchDeleteDocument

Removes one or more documents from an index. The documents must have been added with the [BatchPutDocument](#) (p. 129) operation.

The documents are deleted asynchronously. You can see the progress of the deletion by using AWS CloudWatch. Any error messages related to the processing of the batch are sent to you CloudWatch log.

Request Syntax

```
{
  "DataSourceSyncJobMetricTarget": {
    "DataSourceId": "string",
    "DataSourceSyncJobId": "string"
  },
  "DocumentIdList": [ "string" ],
  "IndexId": "string"
}
```

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 288).

The request accepts the following data in JSON format.

[DataSourceSyncJobMetricTarget](#) (p. 126)

Maps a particular data source sync job to a particular data source.

Type: [DataSourceSyncJobMetricTarget](#) (p. 229) object

Required: No

[DocumentIdList](#) (p. 126)

One or more identifiers for documents to delete from the index.

Type: Array of strings

Array Members: Minimum number of 1 item. Maximum number of 10 items.

Length Constraints: Minimum length of 1. Maximum length of 2048.

Required: Yes

[IndexId](#) (p. 126)

The identifier of the index that contains the documents to delete.

Type: String

Length Constraints: Fixed length of 36.

Pattern: [a-zA-Z0-9][a-zA-Z0-9-]*

Required: Yes

Response Syntax

```
{
```

```
"FailedDocuments": [  
  {  
    "ErrorCode": "string",  
    "ErrorMessage": "string",  
    "Id": "string"  
  }  
]
```

Response Elements

If the action is successful, the service sends back an HTTP 200 response.

The following data is returned in JSON format by the service.

FailedDocuments (p. 126)

A list of documents that could not be removed from the index. Each entry contains an error message that indicates why the document couldn't be removed from the index.

Type: Array of [BatchDeleteDocumentResponseFailedDocument \(p. 212\)](#) objects

Errors

For information about the errors that are common to all actions, see [Common Errors \(p. 286\)](#).

AccessDeniedException

HTTP Status Code: 400

ConflictException

HTTP Status Code: 400

InternalServerErrorException

HTTP Status Code: 500

ResourceNotFoundException

HTTP Status Code: 400

ThrottlingException

HTTP Status Code: 400

ValidationException

HTTP Status Code: 400

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for JavaScript](#)

- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

BatchPutDocument

Adds one or more documents to an index.

The `BatchPutDocument` operation enables you to ingest inline documents or a set of documents stored in an Amazon S3 bucket. Use this operation to ingest your text and unstructured text into an index, add custom attributes to the documents, and to attach an access control list to the documents added to the index.

The documents are indexed asynchronously. You can see the progress of the batch using AWS CloudWatch. Any error messages related to processing the batch are sent to your AWS CloudWatch log.

Request Syntax

```
{
  "Documents": [
    {
      "AccessControlList": [
        {
          "Access": "string",
          "Name": "string",
          "Type": "string"
        }
      ],
      "Attributes": [
        {
          "Key": "string",
          "Value": {
            "DateValue": number,
            "LongValue": number,
            "StringListValue": [ "string" ],
            "StringValue": "string"
          }
        }
      ],
      "Blob": blob,
      "ContentType": "string",
      "Id": "string",
      "S3Path": {
        "Bucket": "string",
        "Key": "string"
      },
      "Title": "string"
    }
  ],
  "IndexId": "string",
  "RoleArn": "string"
}
```

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 288).

The request accepts the following data in JSON format.

Documents (p. 129)

One or more documents to add to the index.

Documents have the following file size limits.

- 5 MB total size for inline documents
- 50 MB total size for files from an S3 bucket
- 5 MB extracted text for any file

For more information about file size and transaction per second quotas, see [Quotas](#).

Type: Array of [Document \(p. 232\)](#) objects

Array Members: Minimum number of 1 item. Maximum number of 10 items.

Required: Yes

IndexId (p. 129)

The identifier of the index to add the documents to. You need to create the index first using the [CreateIndex \(p. 142\)](#) operation.

Type: String

Length Constraints: Fixed length of 36.

Pattern: `[a-zA-Z0-9][a-zA-Z0-9-]*`

Required: Yes

RoleArn (p. 129)

The Amazon Resource Name (ARN) of a role that is allowed to run the `BatchPutDocument` operation. For more information, see [IAM Roles for Amazon Kendra](#).

Type: String

Length Constraints: Minimum length of 1. Maximum length of 1284.

Pattern: `arn:[a-z0-9-\.]{1,63}:[a-z0-9-\.]{0,63}:[a-z0-9-\.]{0,63}:[a-z0-9-\.]{0,63}:[^/].{0,1023}`

Required: No

Response Syntax

```
{
  "FailedDocuments": [
    {
      "ErrorCode": "string",
      "ErrorMessage": "string",
      "Id": "string"
    }
  ]
}
```

Response Elements

If the action is successful, the service sends back an HTTP 200 response.

The following data is returned in JSON format by the service.

FailedDocuments (p. 130)

A list of documents that were not added to the index because the document failed a validation check. Each document contains an error message that indicates why the document couldn't be added to the index.

If there was an error adding a document to an index the error is reported in your AWS CloudWatch log. For more information, see [Monitoring Amazon Kendra with Amazon CloudWatch Logs](#)

Type: Array of [BatchPutDocumentResponseFailedDocument](#) (p. 213) objects

Errors

For information about the errors that are common to all actions, see [Common Errors](#) (p. 286).

AccessDeniedException

HTTP Status Code: 400

ConflictException

HTTP Status Code: 400

InternalServerErrorException

HTTP Status Code: 500

ResourceNotFoundException

HTTP Status Code: 400

ServiceQuotaExceededException

HTTP Status Code: 400

ThrottlingException

HTTP Status Code: 400

ValidationException

HTTP Status Code: 400

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

CreateDataSource

Creates a data source that you use to with an Amazon Kendra index.

You specify a name, connector type and description for your data source. You can choose between an S3 connector, a SharePoint Online connector, and a database connector.

You also specify configuration information such as document metadata (author, source URI, and so on) and user context information.

CreateDataSource is a synchronous operation. The operation returns 200 if the data source was successfully created. Otherwise, an exception is raised.

Request Syntax

```
{
  "Configuration": {
    "DatabaseConfiguration": {
      "AclConfiguration": {
        "AllowedGroupsColumnName": "string"
      },
      "ColumnConfiguration": {
        "ChangeDetectingColumns": [ "string" ],
        "DocumentDataColumnName": "string",
        "DocumentIdColumnName": "string",
        "DocumentTitleColumnName": "string",
        "FieldMappings": [
          {
            "DataSourceFieldName": "string",
            "DateFieldFormat": "string",
            "IndexFieldName": "string"
          }
        ]
      },
    },
    "ConnectionConfiguration": {
      "DatabaseHost": "string",
      "DatabaseName": "string",
      "DatabasePort": number,
      "SecretArn": "string",
      "TableName": "string"
    },
    "DatabaseEngineType": "string",
    "VpcConfiguration": {
      "SecurityGroupIds": [ "string" ],
      "SubnetIds": [ "string" ]
    }
  },
  "OneDriveConfiguration": {
    "ExclusionPatterns": [ "string" ],
    "FieldMappings": [
      {
        "DataSourceFieldName": "string",
        "DateFieldFormat": "string",
        "IndexFieldName": "string"
      }
    ],
    "InclusionPatterns": [ "string" ],
    "OneDriveUsers": {
      "OneDriveUserList": [ "string" ],
      "OneDriveUserS3Path": {
        "Bucket": "string",
        "Key": "string"
      }
    }
  }
}
```

```

    },
    "SecretArn": "string",
    "TenantDomain": "string"
  },
  "S3Configuration": {
    "AccessControlListConfiguration": {
      "KeyPath": "string"
    },
    "BucketName": "string",
    "DocumentsMetadataConfiguration": {
      "S3Prefix": "string"
    },
    "ExclusionPatterns": [ "string" ],
    "InclusionPrefixes": [ "string" ]
  },
  "SalesforceConfiguration": {
    "ChatterFeedConfiguration": {
      "DocumentDataFieldName": "string",
      "DocumentTitleFieldName": "string",
      "FieldMappings": [
        {
          "DataSourceFieldName": "string",
          "DateFieldFormat": "string",
          "IndexFieldName": "string"
        }
      ],
      "IncludeFilterTypes": [ "string" ]
    },
    "CrawlAttachments": boolean,
    "ExcludeAttachmentFilePatterns": [ "string" ],
    "IncludeAttachmentFilePatterns": [ "string" ],
    "KnowledgeArticleConfiguration": {
      "CustomKnowledgeArticleTypeConfigurations": [
        {
          "DocumentDataFieldName": "string",
          "DocumentTitleFieldName": "string",
          "FieldMappings": [
            {
              "DataSourceFieldName": "string",
              "DateFieldFormat": "string",
              "IndexFieldName": "string"
            }
          ],
          "Name": "string"
        }
      ],
      "IncludedStates": [ "string" ],
      "StandardKnowledgeArticleTypeConfiguration": {
        "DocumentDataFieldName": "string",
        "DocumentTitleFieldName": "string",
        "FieldMappings": [
          {
            "DataSourceFieldName": "string",
            "DateFieldFormat": "string",
            "IndexFieldName": "string"
          }
        ]
      }
    },
    "SecretArn": "string",
    "ServerUrl": "string",
    "StandardObjectAttachmentConfiguration": {
      "DocumentTitleFieldName": "string",
      "FieldMappings": [
        {

```

```

        "DataSourceFieldName": "string",
        "DateFieldFormat": "string",
        "IndexFieldName": "string"
    }
]
},
"StandardObjectConfigurations": [
    {
        "DocumentDataFieldName": "string",
        "DocumentTitleFieldName": "string",
        "FieldMappings": [
            {
                "DataSourceFieldName": "string",
                "DateFieldFormat": "string",
                "IndexFieldName": "string"
            }
        ],
        "Name": "string"
    }
]
},
"ServiceNowConfiguration": {
    "HostUrl": "string",
    "KnowledgeArticleConfiguration": {
        "CrawlAttachments": boolean,
        "DocumentDataFieldName": "string",
        "DocumentTitleFieldName": "string",
        "ExcludeAttachmentFilePatterns": [ "string" ],
        "FieldMappings": [
            {
                "DataSourceFieldName": "string",
                "DateFieldFormat": "string",
                "IndexFieldName": "string"
            }
        ],
        "IncludeAttachmentFilePatterns": [ "string" ]
    },
    "SecretArn": "string",
    "ServiceCatalogConfiguration": {
        "CrawlAttachments": boolean,
        "DocumentDataFieldName": "string",
        "DocumentTitleFieldName": "string",
        "ExcludeAttachmentFilePatterns": [ "string" ],
        "FieldMappings": [
            {
                "DataSourceFieldName": "string",
                "DateFieldFormat": "string",
                "IndexFieldName": "string"
            }
        ],
        "IncludeAttachmentFilePatterns": [ "string" ]
    },
    "ServiceNowBuildVersion": "string"
},
"SharePointConfiguration": {
    "CrawlAttachments": boolean,
    "DocumentTitleFieldName": "string",
    "ExclusionPatterns": [ "string" ],
    "FieldMappings": [
        {
            "DataSourceFieldName": "string",
            "DateFieldFormat": "string",
            "IndexFieldName": "string"
        }
    ],
    "InclusionPatterns": [ "string" ],

```

```

        "SecretArn": "string",
        "SharePointVersion": "string",
        "Urls": [ "string" ],
        "UseChangeLog": boolean,
        "VpcConfiguration": {
            "SecurityGroupIds": [ "string" ],
            "SubnetIds": [ "string" ]
        }
    },
    "Description": "string",
    "IndexId": "string",
    "Name": "string",
    "RoleArn": "string",
    "Schedule": "string",
    "Tags": [
        {
            "Key": "string",
            "Value": "string"
        }
    ],
    "Type": "string"
}

```

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters \(p. 288\)](#).

The request accepts the following data in JSON format.

Configuration (p. 132)

The connector configuration information that is required to access the repository.

Type: [DataSourceConfiguration \(p. 221\)](#) object

Required: Yes

Description (p. 132)

A description for the data source.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 1000.

Pattern: `^\P{C}*$`

Required: No

IndexId (p. 132)

The identifier of the index that should be associated with this data source.

Type: String

Length Constraints: Fixed length of 36.

Pattern: `[a-zA-Z0-9][a-zA-Z0-9-]*`

Required: Yes

Name (p. 132)

A unique name for the data source. A data source name can't be changed without deleting and recreating the data source.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 1000.

Pattern: `[a-zA-Z0-9][a-zA-Z0-9_-]*`

Required: Yes

RoleArn (p. 132)

The Amazon Resource Name (ARN) of a role with permission to access the data source. For more information, see [IAM Roles for Amazon Kendra](#).

Type: String

Length Constraints: Minimum length of 1. Maximum length of 1284.

Pattern: `arn:[a-z0-9-\.]{1,63}:[a-z0-9-\.]{0,63}:[a-z0-9-\.]{0,63}:[a-z0-9-\.]{0,63}:[^/]{0,1023}`

Required: Yes

Schedule (p. 132)

Sets the frequency that Amazon Kendra will check the documents in your repository and update the index. If you don't set a schedule Amazon Kendra will not periodically update the index. You can call the `StartDataSourceSyncJob` operation to update the index.

Type: String

Required: No

Tags (p. 132)

A list of key-value pairs that identify the data source. You can use the tags to identify and organize your resources and to control access to resources.

Type: Array of [Tag \(p. 283\)](#) objects

Array Members: Minimum number of 0 items. Maximum number of 200 items.

Required: No

Type (p. 132)

The type of repository that contains the data source.

Type: String

Valid Values: `S3` | `SHAREPOINT` | `DATABASE` | `SALESFORCE` | `ONEDRIVE` | `SERVICENOW`

Required: Yes

Response Syntax

```
{  
  "Id": "string"
```



```
}
```

Response Elements

If the action is successful, the service sends back an HTTP 200 response.

The following data is returned in JSON format by the service.

Id (p. 136)

A unique identifier for the data source.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 100.

Pattern: `[a-zA-Z0-9][a-zA-Z0-9_-]*`

Errors

For information about the errors that are common to all actions, see [Common Errors \(p. 286\)](#).

AccessDeniedException

HTTP Status Code: 400

ConflictException

HTTP Status Code: 400

InternalServerErrorException

HTTP Status Code: 500

ResourceAlreadyExistsException

HTTP Status Code: 400

ResourceNotFoundException

HTTP Status Code: 400

ServiceQuotaExceededException

HTTP Status Code: 400

ThrottlingException

HTTP Status Code: 400

ValidationException

HTTP Status Code: 400

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)

- [AWS SDK for Java](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

CreateFaq

Creates a new set of frequently asked question (FAQ) questions and answers.

Request Syntax

```
{
  "Description": "string",
  "IndexId": "string",
  "Name": "string",
  "RoleArn": "string",
  "S3Path": {
    "Bucket": "string",
    "Key": "string"
  },
  "Tags": [
    {
      "Key": "string",
      "Value": "string"
    }
  ]
}
```

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters \(p. 288\)](#).

The request accepts the following data in JSON format.

Description (p. 139)

A description of the FAQ.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 1000.

Pattern: `^\P{C}*$`

Required: No

IndexId (p. 139)

The identifier of the index that contains the FAQ.

Type: String

Length Constraints: Fixed length of 36.

Pattern: `[a-zA-Z0-9][a-zA-Z0-9-]*`

Required: Yes

Name (p. 139)

The name that should be associated with the FAQ.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 100.

Pattern: `[a-zA-Z0-9][a-zA-Z0-9_-]*`

Required: Yes

RoleArn (p. 139)

The Amazon Resource Name (ARN) of a role with permission to access the S3 bucket that contains the FAQs. For more information, see [IAM Roles for Amazon Kendra](#).

Type: String

Length Constraints: Minimum length of 1. Maximum length of 1284.

Pattern: `arn:[a-z0-9-\.\]{1,63}:[a-z0-9-\.\]{0,63}:[a-z0-9-\.\]{0,63}:[a-z0-9-\.\]{0,63}:[^/].{0,1023}`

Required: Yes

S3Path (p. 139)

The S3 location of the FAQ input data.

Type: [S3Path \(p. 259\)](#) object

Required: Yes

Tags (p. 139)

A list of key-value pairs that identify the FAQ. You can use the tags to identify and organize your resources and to control access to resources.

Type: Array of [Tag \(p. 283\)](#) objects

Array Members: Minimum number of 0 items. Maximum number of 200 items.

Required: No

Response Syntax

```
{  
  "Id": "string"  
}
```

Response Elements

If the action is successful, the service sends back an HTTP 200 response.

The following data is returned in JSON format by the service.

Id (p. 140)

The unique identifier of the FAQ.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 100.

Pattern: `[a-zA-Z0-9][a-zA-Z0-9_-]*`

Errors

For information about the errors that are common to all actions, see [Common Errors \(p. 286\)](#).

AccessDeniedException

HTTP Status Code: 400

ConflictException

HTTP Status Code: 400

InternalServerErrorException

HTTP Status Code: 500

ResourceNotFoundException

HTTP Status Code: 400

ServiceQuotaExceededException

HTTP Status Code: 400

ThrottlingException

HTTP Status Code: 400

ValidationException

HTTP Status Code: 400

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

CreateIndex

Creates a new Amazon Kendra index. Index creation is an asynchronous operation. To determine if index creation has completed, check the `Status` field returned from a call to [DescribeIndex \(p. 162\)](#). The `Status` field is set to `ACTIVE` when the index is ready to use.

Once the index is active you can index your documents using the [BatchPutDocument \(p. 129\)](#) operation or using one of the supported data sources.

Request Syntax

```
{
  "ClientToken": "string",
  "Description": "string",
  "Edition": "string",
  "Name": "string",
  "RoleArn": "string",
  "ServerSideEncryptionConfiguration": {
    "KmsKeyId": "string"
  },
  "Tags": [
    {
      "Key": "string",
      "Value": "string"
    }
  ]
}
```

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters \(p. 288\)](#).

The request accepts the following data in JSON format.

ClientToken (p. 142)

A token that you provide to identify the request to create an index. Multiple calls to the `CreateIndex` operation with the same client token will create only one index."

Type: String

Length Constraints: Minimum length of 1. Maximum length of 100.

Required: No

Description (p. 142)

A description for the index.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 1000.

Pattern: `^\P{C}*$`

Required: No

Edition (p. 142)

The Amazon Kendra edition to use for the index. Choose `DEVELOPER_EDITION` for indexes intended for development, testing, or proof of concept. Use `ENTERPRISE_EDITION` for your production databases. Once you set the edition for an index, it can't be changed.

Type: String

Valid Values: `DEVELOPER_EDITION` | `ENTERPRISE_EDITION`

Required: No

Name (p. 142)

The name for the new index.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 1000.

Pattern: `[a-zA-Z0-9][a-zA-Z0-9_-]*`

Required: Yes

RoleArn (p. 142)

An IAM role that gives Amazon Kendra permissions to access your Amazon CloudWatch logs and metrics. This is also the role used when you use the `BatchPutDocument` operation to index documents from an Amazon S3 bucket.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 1284.

Pattern: `arn:[a-z0-9-\.\.]{1,63}:[a-z0-9-\.\.]{0,63}:[a-z0-9-\.\.]{0,63}:[a-z0-9-\.\.]{0,63}:[^/].{0,1023}`

Required: Yes

ServerSideEncryptionConfiguration (p. 142)

The identifier of the AWS KMS customer managed key (CMK) to use to encrypt data indexed by Amazon Kendra. Amazon Kendra doesn't support asymmetric CMKs.

Type: [ServerSideEncryptionConfiguration \(p. 273\)](#) object

Required: No

Tags (p. 142)

A list of key-value pairs that identify the index. You can use the tags to identify and organize your resources and to control access to resources.

Type: Array of [Tag \(p. 283\)](#) objects

Array Members: Minimum number of 0 items. Maximum number of 200 items.

Required: No

Response Syntax

```
{
  "Id": "string"
}
```

Response Elements

If the action is successful, the service sends back an HTTP 200 response.

The following data is returned in JSON format by the service.

Id (p. 143)

The unique identifier of the index. Use this identifier when you query an index, set up a data source, or index a document.

Type: String

Length Constraints: Fixed length of 36.

Pattern: [a-zA-Z0-9][a-zA-Z0-9-]*

Errors

For information about the errors that are common to all actions, see [Common Errors \(p. 286\)](#).

AccessDeniedException

HTTP Status Code: 400

ConflictException

HTTP Status Code: 400

InternalServerErrorException

HTTP Status Code: 500

ResourceAlreadyExistsException

HTTP Status Code: 400

ServiceQuotaExceededException

HTTP Status Code: 400

ThrottlingException

HTTP Status Code: 400

ValidationException

HTTP Status Code: 400

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

DeleteDataSource

Deletes an Amazon Kendra data source. An exception is not thrown if the data source is already being deleted. While the data source is being deleted, the `Status` field returned by a call to the [DescribeDataSource](#) (p. 151) operation is set to `DELETING`. For more information, see [Deleting Data Sources](#).

Request Syntax

```
{  
  "Id": "string",  
  "IndexId": "string"  
}
```

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 288).

The request accepts the following data in JSON format.

[Id](#) (p. 145)

The unique identifier of the data source to delete.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 100.

Pattern: `[a-zA-Z0-9][a-zA-Z0-9_-]*`

Required: Yes

[IndexId](#) (p. 145)

The unique identifier of the index associated with the data source.

Type: String

Length Constraints: Fixed length of 36.

Pattern: `[a-zA-Z0-9][a-zA-Z0-9-]*`

Required: Yes

Response Elements

If the action is successful, the service sends back an HTTP 200 response with an empty HTTP body.

Errors

For information about the errors that are common to all actions, see [Common Errors](#) (p. 286).

AccessDeniedException

HTTP Status Code: 400

ConflictException

HTTP Status Code: 400

InternalServerErrorException

HTTP Status Code: 500

ResourceNotFoundException

HTTP Status Code: 400

ThrottlingException

HTTP Status Code: 400

ValidationException

HTTP Status Code: 400

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

DeleteFaq

Removes an FAQ from an index.

Request Syntax

```
{  
  "Id": "string",  
  "IndexId": "string"  
}
```

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters \(p. 288\)](#).

The request accepts the following data in JSON format.

[Id \(p. 147\)](#)

The identifier of the FAQ to remove.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 100.

Pattern: `[a-zA-Z0-9][a-zA-Z0-9_-]*`

Required: Yes

[IndexId \(p. 147\)](#)

The index to remove the FAQ from.

Type: String

Length Constraints: Fixed length of 36.

Pattern: `[a-zA-Z0-9][a-zA-Z0-9-]*`

Required: Yes

Response Elements

If the action is successful, the service sends back an HTTP 200 response with an empty HTTP body.

Errors

For information about the errors that are common to all actions, see [Common Errors \(p. 286\)](#).

AccessDeniedException

HTTP Status Code: 400

ConflictException

HTTP Status Code: 400

InternalServerErrorException

HTTP Status Code: 500

ResourceNotFoundException

HTTP Status Code: 400

ThrottlingException

HTTP Status Code: 400

ValidationException

HTTP Status Code: 400

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

DeleteIndex

Deletes an existing Amazon Kendra index. An exception is not thrown if the index is already being deleted. While the index is being deleted, the `Status` field returned by a call to the [DescribeIndex \(p. 162\)](#) operation is set to `DELETING`.

Request Syntax

```
{  
  "Id": "string"  
}
```

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters \(p. 288\)](#).

The request accepts the following data in JSON format.

Id (p. 149)

The identifier of the index to delete.

Type: String

Length Constraints: Fixed length of 36.

Pattern: `[a-zA-Z0-9][a-zA-Z0-9-]*`

Required: Yes

Response Elements

If the action is successful, the service sends back an HTTP 200 response with an empty HTTP body.

Errors

For information about the errors that are common to all actions, see [Common Errors \(p. 286\)](#).

AccessDeniedException

HTTP Status Code: 400

ConflictException

HTTP Status Code: 400

InternalServerErrorException

HTTP Status Code: 500

ResourceNotFoundException

HTTP Status Code: 400

ThrottlingException

HTTP Status Code: 400

ValidationException

HTTP Status Code: 400

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

DescribeDataSource

Gets information about a Amazon Kendra data source.

Request Syntax

```
{  
  "Id": "string",  
  "IndexId": "string"  
}
```

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 288).

The request accepts the following data in JSON format.

Id (p. 151)

The unique identifier of the data source to describe.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 100.

Pattern: [a-zA-Z0-9][a-zA-Z0-9_-]*

Required: Yes

IndexId (p. 151)

The identifier of the index that contains the data source.

Type: String

Length Constraints: Fixed length of 36.

Pattern: [a-zA-Z0-9][a-zA-Z0-9-]*

Required: Yes

Response Syntax

```
{  
  "Configuration": {  
    "DatabaseConfiguration": {  
      "AclConfiguration": {  
        "AllowedGroupsColumnName": "string"  
      },  
      "ColumnConfiguration": {  
        "ChangeDetectingColumns": [ "string" ],  
        "DocumentDataColumnName": "string",  
        "DocumentIdColumnName": "string",  
        "DocumentTitleColumnName": "string",  
        "FieldMappings": [  
          {  
            "DataSourceFieldName": "string",  
            "DateFieldFormat": "string",  
            "IndexFieldName": "string"  
          }  
        ]  
      }  
    }  
  }  
}
```

```

    }
  ]
},
"ConnectionConfiguration": {
  "DatabaseHost": "string",
  "DatabaseName": "string",
  "DatabasePort": number,
  "SecretArn": "string",
  "TableName": "string"
},
"DatabaseEngineType": "string",
"VpcConfiguration": {
  "SecurityGroupIds": [ "string" ],
  "SubnetIds": [ "string" ]
}
},
"OneDriveConfiguration": {
  "ExclusionPatterns": [ "string" ],
  "FieldMappings": [
    {
      "DataSourceFieldName": "string",
      "DateFieldFormat": "string",
      "IndexFieldName": "string"
    }
  ],
  "InclusionPatterns": [ "string" ],
  "OneDriveUsers": {
    "OneDriveUserList": [ "string" ],
    "OneDriveUserS3Path": {
      "Bucket": "string",
      "Key": "string"
    }
  },
  "SecretArn": "string",
  "TenantDomain": "string"
},
"S3Configuration": {
  "AccessControlListConfiguration": {
    "KeyPath": "string"
  },
  "BucketName": "string",
  "DocumentsMetadataConfiguration": {
    "S3Prefix": "string"
  },
  "ExclusionPatterns": [ "string" ],
  "InclusionPrefixes": [ "string" ]
},
"SalesforceConfiguration": {
  "ChatterFeedConfiguration": {
    "DocumentDataFieldName": "string",
    "DocumentTitleFieldName": "string",
    "FieldMappings": [
      {
        "DataSourceFieldName": "string",
        "DateFieldFormat": "string",
        "IndexFieldName": "string"
      }
    ]
  },
  "IncludeFilterTypes": [ "string" ]
},
"CrawlAttachments": boolean,
"ExcludeAttachmentFilePatterns": [ "string" ],
"IncludeAttachmentFilePatterns": [ "string" ],
"KnowledgeArticleConfiguration": {
  "CustomKnowledgeArticleTypeConfigurations": [
    {

```



```

        "DocumentDataFieldName": "string",
        "DocumentTitleFieldName": "string",
        "FieldMappings": [
            {
                "DataSourceFieldName": "string",
                "DateFieldFormat": "string",
                "IndexFieldName": "string"
            }
        ],
        "Name": "string"
    }
],
"IncludedStates": [ "string" ],
"StandardKnowledgeArticleTypeConfiguration": {
    "DocumentDataFieldName": "string",
    "DocumentTitleFieldName": "string",
    "FieldMappings": [
        {
            "DataSourceFieldName": "string",
            "DateFieldFormat": "string",
            "IndexFieldName": "string"
        }
    ]
},
"SecretArn": "string",
"ServerUrl": "string",
"StandardObjectAttachmentConfiguration": {
    "DocumentTitleFieldName": "string",
    "FieldMappings": [
        {
            "DataSourceFieldName": "string",
            "DateFieldFormat": "string",
            "IndexFieldName": "string"
        }
    ]
},
"StandardObjectConfigurations": [
    {
        "DocumentDataFieldName": "string",
        "DocumentTitleFieldName": "string",
        "FieldMappings": [
            {
                "DataSourceFieldName": "string",
                "DateFieldFormat": "string",
                "IndexFieldName": "string"
            }
        ],
        "Name": "string"
    }
],
"ServiceNowConfiguration": {
    "HostUrl": "string",
    "KnowledgeArticleConfiguration": {
        "CrawlAttachments": boolean,
        "DocumentDataFieldName": "string",
        "DocumentTitleFieldName": "string",
        "ExcludeAttachmentFilePatterns": [ "string" ],
        "FieldMappings": [
            {
                "DataSourceFieldName": "string",
                "DateFieldFormat": "string",
                "IndexFieldName": "string"
            }
        ]
    }
},

```

```

        "IncludeAttachmentFilePatterns": [ "string" ]
    },
    "SecretArn": "string",
    "ServiceCatalogConfiguration": {
        "CrawlAttachments": boolean,
        "DocumentDataFieldName": "string",
        "DocumentTitleFieldName": "string",
        "ExcludeAttachmentFilePatterns": [ "string" ],
        "FieldMappings": [
            {
                "DataSourceFieldName": "string",
                "DateFieldFormat": "string",
                "IndexFieldName": "string"
            }
        ],
        "IncludeAttachmentFilePatterns": [ "string" ]
    },
    "ServiceNowBuildVersion": "string"
},
"SharePointConfiguration": {
    "CrawlAttachments": boolean,
    "DocumentTitleFieldName": "string",
    "ExclusionPatterns": [ "string" ],
    "FieldMappings": [
        {
            "DataSourceFieldName": "string",
            "DateFieldFormat": "string",
            "IndexFieldName": "string"
        }
    ],
    "InclusionPatterns": [ "string" ],
    "SecretArn": "string",
    "SharePointVersion": "string",
    "Urls": [ "string" ],
    "UseChangeLog": boolean,
    "VpcConfiguration": {
        "SecurityGroupIds": [ "string" ],
        "SubnetIds": [ "string" ]
    }
}
},
"CreatedAt": number,
"Description": "string",
"ErrorMessage": "string",
"Id": "string",
"IndexId": "string",
"Name": "string",
"RoleArn": "string",
"Schedule": "string",
"Status": "string",
"Type": "string",
"UpdatedAt": number
}

```

Response Elements

If the action is successful, the service sends back an HTTP 200 response.

The following data is returned in JSON format by the service.

Configuration (p. 151)

Information that describes where the data source is located and how the data source is configured. The specific information in the description depends on the data source provider.

Type: [DataSourceConfiguration \(p. 221\)](#) object

CreatedAt (p. 151)

The Unix timestamp of when the data source was created.

Type: Timestamp

Description (p. 151)

The description of the data source.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 1000.

Pattern: `^\P{C}*$`

ErrorMessage (p. 151)

When the `Status` field value is `FAILED`, the `ErrorMessage` field contains a description of the error that caused the data source to fail.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 2048.

Pattern: `^\P{C}*$`

Id (p. 151)

The identifier of the data source.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 100.

Pattern: `[a-zA-Z0-9][a-zA-Z0-9_-]*`

IndexId (p. 151)

The identifier of the index that contains the data source.

Type: String

Length Constraints: Fixed length of 36.

Pattern: `[a-zA-Z0-9][a-zA-Z0-9_-]*`

Name (p. 151)

The name that you gave the data source when it was created.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 1000.

Pattern: `[a-zA-Z0-9][a-zA-Z0-9_-]*`

RoleArn (p. 151)

The Amazon Resource Name (ARN) of the role that enables the data source to access its resources.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 1284.

Pattern: `arn:[a-z0-9-\.]{1,63}:[a-z0-9-\.]{0,63}:[a-z0-9-\.]{0,63}:[a-z0-9-\.]{0,63}:[^/]{0,1023}`

Schedule (p. 151)

The schedule that Amazon Kendra will update the data source.

Type: String

Status (p. 151)

The current status of the data source. When the status is `ACTIVE` the data source is ready to use. When the status is `FAILED`, the `ErrorMessage` field contains the reason that the data source failed.

Type: String

Valid Values: `CREATING` | `DELETING` | `FAILED` | `UPDATING` | `ACTIVE`

Type (p. 151)

The type of the data source.

Type: String

Valid Values: `S3` | `SHAREPOINT` | `DATABASE` | `SALESFORCE` | `ONEDRIVE` | `SERVICENOW`

UpdatedAt (p. 151)

The Unix timestamp of when the data source was last updated.

Type: Timestamp

Errors

For information about the errors that are common to all actions, see [Common Errors \(p. 286\)](#).

AccessDeniedException

HTTP Status Code: 400

InternalServerErrorException

HTTP Status Code: 500

ResourceNotFoundException

HTTP Status Code: 400

ThrottlingException

HTTP Status Code: 400

ValidationException

HTTP Status Code: 400

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)

- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

DescribeFaq

Gets information about an FAQ list.

Request Syntax

```
{  
  "Id": "string",  
  "IndexId": "string"  
}
```

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 288).

The request accepts the following data in JSON format.

Id (p. 158)

The unique identifier of the FAQ.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 100.

Pattern: [a-zA-Z0-9][a-zA-Z0-9_-]*

Required: Yes

IndexId (p. 158)

The identifier of the index that contains the FAQ.

Type: String

Length Constraints: Fixed length of 36.

Pattern: [a-zA-Z0-9][a-zA-Z0-9-]*

Required: Yes

Response Syntax

```
{  
  "CreatedAt": number,  
  "Description": "string",  
  "ErrorMessage": "string",  
  "Id": "string",  
  "IndexId": "string",  
  "Name": "string",  
  "RoleArn": "string",  
  "S3Path": {  
    "Bucket": "string",  
    "Key": "string"  
  },  
  "Status": "string",  
  "UpdatedAt": number  
}
```

```
}
```

Response Elements

If the action is successful, the service sends back an HTTP 200 response.

The following data is returned in JSON format by the service.

CreatedAt (p. 158)

The date and time that the FAQ was created.

Type: Timestamp

Description (p. 158)

The description of the FAQ that you provided when it was created.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 1000.

Pattern: `^\P{C}*$`

ErrorMessage (p. 158)

If the `Status` field is `FAILED`, the `ErrorMessage` field contains the reason why the FAQ failed.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 2048.

Pattern: `^\P{C}*$`

Id (p. 158)

The identifier of the FAQ.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 100.

Pattern: `[a-zA-Z0-9][a-zA-Z0-9_-]*`

IndexId (p. 158)

The identifier of the index that contains the FAQ.

Type: String

Length Constraints: Fixed length of 36.

Pattern: `[a-zA-Z0-9][a-zA-Z0-9-]*`

Name (p. 158)

The name that you gave the FAQ when it was created.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 100.

Pattern: `[a-zA-Z0-9][a-zA-Z0-9_-]*`

RoleArn (p. 158)

The Amazon Resource Name (ARN) of the role that provides access to the S3 bucket containing the input files for the FAQ.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 1284.

Pattern: `arn:[a-z0-9-\.]{1,63}:[a-z0-9-\.]{0,63}:[a-z0-9-\.]{0,63}:[a-z0-9-\.]{0,63}:[^/]{0,1023}`

S3Path (p. 158)

Information required to find a specific file in an Amazon S3 bucket.

Type: [S3Path \(p. 259\)](#) object

Status (p. 158)

The status of the FAQ. It is ready to use when the status is `ACTIVE`.

Type: String

Valid Values: `CREATING` | `UPDATING` | `ACTIVE` | `DELETING` | `FAILED`

UpdatedAt (p. 158)

The date and time that the FAQ was last updated.

Type: Timestamp

Errors

For information about the errors that are common to all actions, see [Common Errors \(p. 286\)](#).

AccessDeniedException

HTTP Status Code: 400

InternalServerErrorException

HTTP Status Code: 500

ResourceNotFoundException

HTTP Status Code: 400

ThrottlingException

HTTP Status Code: 400

ValidationException

HTTP Status Code: 400

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)

- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

DescribeIndex

Describes an existing Amazon Kendra index

Request Syntax

```
{  
  "Id": "string"  
}
```

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 288).

The request accepts the following data in JSON format.

Id (p. 162)

The name of the index to describe.

Type: String

Length Constraints: Fixed length of 36.

Pattern: [a-zA-Z0-9][a-zA-Z0-9-]*

Required: Yes

Response Syntax

```
{  
  "CapacityUnits": {  
    "QueryCapacityUnits": number,  
    "StorageCapacityUnits": number  
  },  
  "CreatedAt": number,  
  "Description": "string",  
  "DocumentMetadataConfigurations": [  
    {  
      "Name": "string",  
      "Relevance": {  
        "Duration": "string",  
        "Freshness": boolean,  
        "Importance": number,  
        "RankOrder": "string",  
        "ValueImportanceMap": {  
          "string" : number  
        }  
      },  
    },  
    {  
      "Search": {  
        "Displayable": boolean,  
        "Facetable": boolean,  
        "Searchable": boolean  
      },  
      "Type": "string"  
    }  
  ],  
  "Edition": "string",  
}
```

```
"ErrorMessage": "string",
"Id": "string",
"IndexStatistics": {
  "FaqStatistics": {
    "IndexedQuestionAnswersCount": number
  },
  "TextDocumentStatistics": {
    "IndexedTextBytes": number,
    "IndexedTextDocumentsCount": number
  }
},
"Name": "string",
"RoleArn": "string",
"ServerSideEncryptionConfiguration": {
  "KmsKeyId": "string"
},
"Status": "string",
"UpdatedAt": number
}
```

Response Elements

If the action is successful, the service sends back an HTTP 200 response.

The following data is returned in JSON format by the service.

CapacityUnits (p. 162)

For enterprise edition indexes, you can choose to use additional capacity to meet the needs of your application. This contains the capacity units used for the index. A 0 for the query capacity or the storage capacity indicates that the index is using the default capacity for the index.

Type: [CapacityUnitsConfiguration \(p. 214\)](#) object

CreatedAt (p. 162)

The Unix datetime that the index was created.

Type: Timestamp

Description (p. 162)

The description of the index.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 1000.

Pattern: `^\P{C}*$`

DocumentMetadataConfigurations (p. 162)

Configuration settings for any metadata applied to the documents in the index.

Type: Array of [DocumentMetadataConfiguration \(p. 237\)](#) objects

Array Members: Minimum number of 0 items. Maximum number of 500 items.

Edition (p. 162)

The Amazon Kendra edition used for the index. You decide the edition when you create the index.

Type: String

Valid Values: `DEVELOPER_EDITION` | `ENTERPRISE_EDITION`

ErrorMessage (p. 162)

When the `status` field value is `FAILED`, the `ErrorMessage` field contains a message that explains why.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 2048.

Pattern: `^\P{C}*$`

Id (p. 162)

the name of the index.

Type: String

Length Constraints: Fixed length of 36.

Pattern: `[a-zA-Z0-9][a-zA-Z0-9-]*`

IndexStatistics (p. 162)

Provides information about the number of FAQ questions and answers and the number of text documents indexed.

Type: [IndexStatistics \(p. 247\)](#) object

Name (p. 162)

The name of the index.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 1000.

Pattern: `[a-zA-Z0-9][a-zA-Z0-9_-]*`

RoleArn (p. 162)

The Amazon Resource Name (ARN) of the IAM role that gives Amazon Kendra permission to write to your Amazon Cloudwatch logs.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 1284.

Pattern: `arn:[a-z0-9-\.\.]{1,63}:[a-z0-9-\.\.]{0,63}:[a-z0-9-\.\.]{0,63}:[a-z0-9-\.\.]{0,63}:[^/].{0,1023}`

ServerSideEncryptionConfiguration (p. 162)

The identifier of the AWS KMS customer master key (CMK) used to encrypt your data. Amazon Kendra doesn't support asymmetric CMKs.

Type: [ServerSideEncryptionConfiguration \(p. 273\)](#) object

Status (p. 162)

The current status of the index. When the value is `ACTIVE`, the index is ready for use. If the `Status` field value is `FAILED`, the `ErrorMessage` field contains a message that explains why.

Type: String

Valid Values: `CREATING` | `ACTIVE` | `DELETING` | `FAILED` | `UPDATING` | `SYSTEM_UPDATING`

UpdatedAt (p. 162)

The Unix datetime that the index was last updated.

Type: Timestamp

Errors

For information about the errors that are common to all actions, see [Common Errors \(p. 286\)](#).

AccessDeniedException

HTTP Status Code: 400

InternalServerErrorException

HTTP Status Code: 500

ResourceNotFoundException

HTTP Status Code: 400

ThrottlingException

HTTP Status Code: 400

ValidationException

HTTP Status Code: 400

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

ListDataSources

Lists the data sources that you have created.

Request Syntax

```
{  
  "IndexId": "string",  
  "MaxResults": number,  
  "NextToken": "string"  
}
```

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 288).

The request accepts the following data in JSON format.

IndexId (p. 166)

The identifier of the index that contains the data source.

Type: String

Length Constraints: Fixed length of 36.

Pattern: [a-zA-Z0-9][a-zA-Z0-9-]*

Required: Yes

MaxResults (p. 166)

The maximum number of data sources to return.

Type: Integer

Valid Range: Minimum value of 1. Maximum value of 100.

Required: No

NextToken (p. 166)

If the previous response was incomplete (because there is more data to retrieve), Amazon Kendra returns a pagination token in the response. You can use this pagination token to retrieve the next set of data sources (`DataSourceSummaryItems`).

Type: String

Length Constraints: Minimum length of 1. Maximum length of 800.

Required: No

Response Syntax

```
{  
  "NextToken": "string",  
  "SummaryItems": [  
    {  
      ...  
    }  
  ]  
}
```

```
    "CreatedAt": number,  
    "Id": "string",  
    "Name": "string",  
    "Status": "string",  
    "Type": "string",  
    "UpdatedAt": number  
  }  
]  
}
```

Response Elements

If the action is successful, the service sends back an HTTP 200 response.

The following data is returned in JSON format by the service.

NextToken (p. 166)

If the response is truncated, Amazon Kendra returns this token that you can use in the subsequent request to retrieve the next set of data sources.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 800.

SummaryItems (p. 166)

An array of summary information for one or more data sources.

Type: Array of [DataSourceSummary \(p. 223\)](#) objects

Errors

For information about the errors that are common to all actions, see [Common Errors \(p. 286\)](#).

AccessDeniedException

HTTP Status Code: 400

InternalServerErrorException

HTTP Status Code: 500

ResourceNotFoundException

HTTP Status Code: 400

ThrottlingException

HTTP Status Code: 400

ValidationException

HTTP Status Code: 400

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

ListDataSourceSyncJobs

Gets statistics about synchronizing Amazon Kendra with a data source.

Request Syntax

```
{  
  "Id": "string",  
  "IndexId": "string",  
  "MaxResults": number,  
  "NextToken": "string",  
  "StartTimeFilter": {  
    "EndTime": number,  
    "StartTime": number  
  },  
  "StatusFilter": "string"  
}
```

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 288).

The request accepts the following data in JSON format.

Id (p. 169)

The identifier of the data source.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 100.

Pattern: [a-zA-Z0-9][a-zA-Z0-9_-]*

Required: Yes

IndexId (p. 169)

The identifier of the index that contains the data source.

Type: String

Length Constraints: Fixed length of 36.

Pattern: [a-zA-Z0-9][a-zA-Z0-9-]*

Required: Yes

MaxResults (p. 169)

The maximum number of synchronization jobs to return in the response. If there are fewer results in the list, this response contains only the actual results.

Type: Integer

Valid Range: Minimum value of 1. Maximum value of 10.

Required: No

NextToken (p. 169)

If the result of the previous request to `GetDataSourceSyncJobHistory` was truncated, include the `NextToken` to fetch the next set of jobs.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 800.

Required: No

StartTimeFilter (p. 169)

When specified, the synchronization jobs returned in the list are limited to jobs between the specified dates.

Type: [TimeRange \(p. 286\)](#) object

Required: No

StatusFilter (p. 169)

When specified, only returns synchronization jobs with the `Status` field equal to the specified status.

Type: String

Valid Values: `FAILED` | `SUCCEEDED` | `SYNCING` | `INCOMPLETE` | `STOPPING` | `ABORTED` | `SYNCING_INDEXING`

Required: No

Response Syntax

```
{
  "History": [
    {
      "DataSourceErrorCode": "string",
      "EndTime": number,
      "ErrorCode": "string",
      "ErrorMessage": "string",
      "ExecutionId": "string",
      "Metrics": {
        "DocumentsAdded": "string",
        "DocumentsDeleted": "string",
        "DocumentsFailed": "string",
        "DocumentsModified": "string",
        "DocumentsScanned": "string"
      },
      "StartTime": number,
      "Status": "string"
    }
  ],
  "NextToken": "string"
}
```

Response Elements

If the action is successful, the service sends back an HTTP 200 response.

The following data is returned in JSON format by the service.

History (p. 170)

A history of synchronization jobs for the data source.

Type: Array of [DataSourceSyncJob \(p. 225\)](#) objects

NextToken (p. 170)

The `GetDataSourceSyncJobHistory` operation returns a page of vocabularies at a time. The maximum size of the page is set by the `MaxResults` parameter. If there are more jobs in the list than the page size, Amazon Kendra returns the `NextPage` token. Include the token in the next request to the `GetDataSourceSyncJobHistory` operation to return in the next page of jobs.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 800.

Errors

For information about the errors that are common to all actions, see [Common Errors \(p. 286\)](#).

AccessDeniedException

HTTP Status Code: 400

ConflictException

HTTP Status Code: 400

InternalServerErrorException

HTTP Status Code: 500

ResourceNotFoundException

HTTP Status Code: 400

ThrottlingException

HTTP Status Code: 400

ValidationException

HTTP Status Code: 400

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

ListFaq

Gets a list of FAQ lists associated with an index.

Request Syntax

```
{  
  "IndexId": "string",  
  "MaxResults": number,  
  "NextToken": "string"  
}
```

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 288).

The request accepts the following data in JSON format.

IndexId (p. 172)

The index that contains the FAQ lists.

Type: String

Length Constraints: Fixed length of 36.

Pattern: [a-zA-Z0-9][a-zA-Z0-9-]*

Required: Yes

MaxResults (p. 172)

The maximum number of FAQs to return in the response. If there are fewer results in the list, this response contains only the actual results.

Type: Integer

Valid Range: Minimum value of 1. Maximum value of 100.

Required: No

NextToken (p. 172)

If the result of the previous request to `ListFaq` was truncated, include the `NextToken` to fetch the next set of FAQs.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 800.

Required: No

Response Syntax

```
{  
  "FaqSummaryItems": [  
    {  
      "CreatedAt": number,  
      ...  
    }  
  ]  
}
```

```
        "Id": "string",  
        "Name": "string",  
        "Status": "string",  
        "UpdatedAt": number  
    }  
  ],  
  "NextToken": "string"  
}
```

Response Elements

If the action is successful, the service sends back an HTTP 200 response.

The following data is returned in JSON format by the service.

FaqSummaryItems (p. 172)

information about the FAQs associated with the specified index.

Type: Array of [FaqSummary](#) (p. 242) objects

NextToken (p. 172)

The `ListFaq` operation returns a page of FAQs at a time. The maximum size of the page is set by the `MaxResults` parameter. If there are more jobs in the list than the page size, Amazon Kendra returns the `NextPage` token. Include the token in the next request to the `ListFaq` operation to return the next page of FAQs.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 800.

Errors

For information about the errors that are common to all actions, see [Common Errors](#) (p. 286).

AccessDeniedException

HTTP Status Code: 400

InternalServerErrorException

HTTP Status Code: 500

ResourceNotFoundException

HTTP Status Code: 400

ThrottlingException

HTTP Status Code: 400

ValidationException

HTTP Status Code: 400

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)

- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

ListIndices

Lists the Amazon Kendra indexes that you have created.

Request Syntax

```
{  
  "MaxResults": number,  
  "NextToken": "string"  
}
```

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 288).

The request accepts the following data in JSON format.

MaxResults (p. 175)

The maximum number of data sources to return.

Type: Integer

Valid Range: Minimum value of 1. Maximum value of 100.

Required: No

NextToken (p. 175)

If the previous response was incomplete (because there is more data to retrieve), Amazon Kendra returns a pagination token in the response. You can use this pagination token to retrieve the next set of indexes (DataSourceSummaryItems).

Type: String

Length Constraints: Minimum length of 1. Maximum length of 800.

Required: No

Response Syntax

```
{  
  "IndexConfigurationSummaryItems": [  
    {  
      "CreatedAt": number,  
      "Edition": "string",  
      "Id": "string",  
      "Name": "string",  
      "Status": "string",  
      "UpdatedAt": number  
    }  
  ],  
  "NextToken": "string"  
}
```

Response Elements

If the action is successful, the service sends back an HTTP 200 response.

The following data is returned in JSON format by the service.

IndexConfigurationSummaryItems (p. 175)

An array of summary information for one or more indexes.

Type: Array of [IndexConfigurationSummary \(p. 245\)](#) objects

NextToken (p. 175)

If the response is truncated, Amazon Kendra returns this token that you can use in the subsequent request to retrieve the next set of indexes.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 800.

Errors

For information about the errors that are common to all actions, see [Common Errors \(p. 286\)](#).

AccessDeniedException

HTTP Status Code: 400

InternalServerErrorException

HTTP Status Code: 500

ThrottlingException

HTTP Status Code: 400

ValidationException

HTTP Status Code: 400

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

ListTagsForResource

Gets a list of tags associated with a specified resource. Indexes, FAQs, and data sources can have tags associated with them.

Request Syntax

```
{  
  "ResourceARN": "string"  
}
```

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 288).

The request accepts the following data in JSON format.

ResourceARN (p. 177)

The Amazon Resource Name (ARN) of the index, FAQ, or data source to get a list of tags for.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 1011.

Required: Yes

Response Syntax

```
{  
  "Tags": [  
    {  
      "Key": "string",  
      "Value": "string"  
    }  
  ]  
}
```

Response Elements

If the action is successful, the service sends back an HTTP 200 response.

The following data is returned in JSON format by the service.

Tags (p. 177)

A list of tags associated with the index, FAQ, or data source.

Type: Array of [Tag](#) (p. 283) objects

Array Members: Minimum number of 0 items. Maximum number of 200 items.

Errors

For information about the errors that are common to all actions, see [Common Errors](#) (p. 286).

AccessDeniedException

HTTP Status Code: 400

InternalServerErrorException

HTTP Status Code: 500

ResourceUnavailableException

HTTP Status Code: 400

ThrottlingException

HTTP Status Code: 400

ValidationException

HTTP Status Code: 400

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

Query

Searches an active index. Use this API to search your documents using query. The `Query` operation enables to do faceted search and to filter results based on document attributes.

It also enables you to provide user context that Amazon Kendra uses to enforce document access control in the search results.

Amazon Kendra searches your index for text content and question and answer (FAQ) content. By default the response contains three types of results.

- Relevant passages
- Matching FAQs
- Relevant documents

You can specify that the query return only one type of result using the `QueryResultTypeConfig` parameter.

Request Syntax

```
{
  "AttributeFilter": {
    "AndAllFilters": [
      "AttributeFilter"
    ],
    "ContainsAll": {
      "Key": "string",
      "Value": {
        "DateValue": number,
        "LongValue": number,
        "StringListValue": [ "string" ],
        "StringValue": "string"
      }
    },
    "ContainsAny": {
      "Key": "string",
      "Value": {
        "DateValue": number,
        "LongValue": number,
        "StringListValue": [ "string" ],
        "StringValue": "string"
      }
    },
    "EqualsTo": {
      "Key": "string",
      "Value": {
        "DateValue": number,
        "LongValue": number,
        "StringListValue": [ "string" ],
        "StringValue": "string"
      }
    },
    "GreaterThan": {
      "Key": "string",
      "Value": {
        "DateValue": number,
        "LongValue": number,
        "StringListValue": [ "string" ],
        "StringValue": "string"
      }
    }
  }
}
```

```

    },
    "GreaterThanOrEquals": {
      "Key": "string",
      "Value": {
        "DateValue": number,
        "LongValue": number,
        "StringListValue": [ "string" ],
        "StringValue": "string"
      }
    },
    "LessThan": {
      "Key": "string",
      "Value": {
        "DateValue": number,
        "LongValue": number,
        "StringListValue": [ "string" ],
        "StringValue": "string"
      }
    },
    "LessThanOrEquals": {
      "Key": "string",
      "Value": {
        "DateValue": number,
        "LongValue": number,
        "StringListValue": [ "string" ],
        "StringValue": "string"
      }
    },
    "NotFilter": "AttributeFilter",
    "OrAllFilters": [
      "AttributeFilter"
    ]
  },
  "Facets": [
    {
      "DocumentAttributeKey": "string"
    }
  ],
  "IndexId": "string",
  "PageNumber": number,
  "PageSize": number,
  "QueryResultTypeFilter": "string",
  "QueryText": "string",
  "RequestedDocumentAttributes": [ "string" ]
}

```

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 288).

The request accepts the following data in JSON format.

AttributeFilter (p. 179)

Enables filtered searches based on document attributes. You can only provide one attribute filter; however, the `AndAllFilters`, `NotFilter`, and `OrAllFilters` parameters contain a list of other filters.

The `AttributeFilter` parameter enables you to create a set of filtering rules that a document must satisfy to be included in the query results.

Type: [AttributeFilter](#) (p. 210) object

Required: No

Facets (p. 179)

An array of documents attributes. Amazon Kendra returns a count for each attribute key specified. You can use this information to help narrow the search for your user.

Type: Array of [Facet \(p. 239\)](#) objects

Required: No

IndexId (p. 179)

The unique identifier of the index to search. The identifier is returned in the response from the [CreateIndex \(p. 142\)](#) operation.

Type: String

Length Constraints: Fixed length of 36.

Pattern: `[a-zA-Z0-9][a-zA-Z0-9-]*`

Required: Yes

PageNumber (p. 179)

Query results are returned in pages the size of the `PageSize` parameter. By default, Amazon Kendra returns the first page of results. Use this parameter to get result pages after the first one.

Type: Integer

Required: No

PageSize (p. 179)

Sets the number of results that are returned in each page of results. The default page size is 10. The maximum number of results returned is 100. If you ask for more than 100 results, only 100 are returned.

Type: Integer

Required: No

QueryResultTypeFilter (p. 179)

Sets the type of query. Only results for the specified query type are returned.

Type: String

Valid Values: `DOCUMENT` | `QUESTION_ANSWER` | `ANSWER`

Required: No

QueryText (p. 179)

The text to search for.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 1000.

Pattern: `^\P{C}*$`

Required: Yes

RequestedDocumentAttributes (p. 179)

An array of document attributes to include in the response. No other document attributes are included in the response. By default all document attributes are included in the response.

Type: Array of strings

Array Members: Minimum number of 1 item. Maximum number of 100 items.

Length Constraints: Minimum length of 1. Maximum length of 200.

Pattern: [a-zA-Z0-9_][a-zA-Z0-9_-]*

Required: No

Response Syntax

```
{
  "FacetResults": [
    {
      "DocumentAttributeKey": "string",
      "DocumentAttributeValueCountPairs": [
        {
          "Count": number,
          "DocumentAttributeValue": {
            "DateValue": number,
            "LongValue": number,
            "StringListValue": [ "string" ],
            "StringValue": "string"
          }
        }
      ]
    }
  ],
  "QueryId": "string",
  "ResultItems": [
    {
      "AdditionalAttributes": [
        {
          "Key": "string",
          "Value": {
            "TextWithHighlightsValue": {
              "Highlights": [
                {
                  "BeginOffset": number,
                  "EndOffset": number,
                  "TopAnswer": boolean
                }
              ],
              "Text": "string"
            }
          },
          "ValueType": "string"
        }
      ],
      "DocumentAttributes": [
        {
          "Key": "string",
          "Value": {
            "DateValue": number,
            "LongValue": number,
            "StringListValue": [ "string" ],
            "StringValue": "string"
          }
        }
      ]
    }
  ]
}
```

```

    }
  },
  "DocumentExcerpt": {
    "Highlights": [
      {
        "BeginOffset": number,
        "EndOffset": number,
        "TopAnswer": boolean
      }
    ],
    "Text": "string"
  },
  "DocumentId": "string",
  "DocumentTitle": {
    "Highlights": [
      {
        "BeginOffset": number,
        "EndOffset": number,
        "TopAnswer": boolean
      }
    ],
    "Text": "string"
  },
  "DocumentURI": "string",
  "Id": "string",
  "Type": "string"
}
],
"TotalNumberOfResults": number
}

```

Response Elements

If the action is successful, the service sends back an HTTP 200 response.

The following data is returned in JSON format by the service.

FacetResults (p. 182)

Contains the facet results. A `FacetResult` contains the counts for each attribute key that was specified in the `Facets` input parameter.

Type: Array of [FacetResult \(p. 240\)](#) objects

QueryId (p. 182)

The unique identifier for the search. You use `QueryId` to identify the search when using the feedback API.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 36.

ResultItems (p. 182)

The results of the search.

Type: Array of [QueryResultItem \(p. 252\)](#) objects

TotalNumberOfResults (p. 182)

The number of items returned by the search. Use this to determine when you have requested the last set of results.

Type: Integer

Errors

For information about the errors that are common to all actions, see [Common Errors \(p. 286\)](#).

AccessDeniedException

HTTP Status Code: 400

ConflictException

HTTP Status Code: 400

InternalServerErrorException

HTTP Status Code: 500

ResourceNotFoundException

HTTP Status Code: 400

ServiceQuotaExceededException

HTTP Status Code: 400

ThrottlingException

HTTP Status Code: 400

ValidationException

HTTP Status Code: 400

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

StartDataSourceSyncJob

Starts a synchronization job for a data source. If a synchronization job is already in progress, Amazon Kendra returns a `ResourceInUseException` exception.

Request Syntax

```
{  
  "Id": "string",  
  "IndexId": "string"  
}
```

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters \(p. 288\)](#).

The request accepts the following data in JSON format.

Id (p. 185)

The identifier of the data source to synchronize.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 100.

Pattern: `[a-zA-Z0-9][a-zA-Z0-9_-]*`

Required: Yes

IndexId (p. 185)

The identifier of the index that contains the data source.

Type: String

Length Constraints: Fixed length of 36.

Pattern: `[a-zA-Z0-9][a-zA-Z0-9-]*`

Required: Yes

Response Syntax

```
{  
  "ExecutionId": "string"  
}
```

Response Elements

If the action is successful, the service sends back an HTTP 200 response.

The following data is returned in JSON format by the service.

ExecutionId (p. 185)

Identifies a particular synchronization job.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 2048.

Errors

For information about the errors that are common to all actions, see [Common Errors \(p. 286\)](#).

AccessDeniedException

HTTP Status Code: 400

ConflictException

HTTP Status Code: 400

InternalServerErrorException

HTTP Status Code: 500

ResourceInUseException

HTTP Status Code: 400

ResourceNotFoundException

HTTP Status Code: 400

ThrottlingException

HTTP Status Code: 400

ValidationException

HTTP Status Code: 400

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

StopDataSourceSyncJob

Stops a running synchronization job. You can't stop a scheduled synchronization job.

Request Syntax

```
{  
  "Id": "string",  
  "IndexId": "string"  
}
```

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters \(p. 288\)](#).

The request accepts the following data in JSON format.

[Id \(p. 187\)](#)

The identifier of the data source for which to stop the synchronization jobs.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 100.

Pattern: `[a-zA-Z0-9][a-zA-Z0-9_-]*`

Required: Yes

[IndexId \(p. 187\)](#)

The identifier of the index that contains the data source.

Type: String

Length Constraints: Fixed length of 36.

Pattern: `[a-zA-Z0-9][a-zA-Z0-9-]*`

Required: Yes

Response Elements

If the action is successful, the service sends back an HTTP 200 response with an empty HTTP body.

Errors

For information about the errors that are common to all actions, see [Common Errors \(p. 286\)](#).

AccessDeniedException

HTTP Status Code: 400

InternalServerErrorException

HTTP Status Code: 500

ResourceNotFoundException

HTTP Status Code: 400

ThrottlingException

HTTP Status Code: 400

ValidationException

HTTP Status Code: 400

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

SubmitFeedback

Enables you to provide feedback to Amazon Kendra to improve the performance of the service.

Request Syntax

```
{
  "ClickFeedbackItems": [
    {
      "ClickTime": number,
      "ResultId": "string"
    }
  ],
  "IndexId": "string",
  "QueryId": "string",
  "RelevanceFeedbackItems": [
    {
      "RelevanceValue": "string",
      "ResultId": "string"
    }
  ]
}
```

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 288).

The request accepts the following data in JSON format.

[ClickFeedbackItems](#) (p. 189)

Tells Amazon Kendra that a particular search result link was chosen by the user.

Type: Array of [ClickFeedback](#) (p. 215) objects

Required: No

[IndexId](#) (p. 189)

The identifier of the index that was queried.

Type: String

Length Constraints: Fixed length of 36.

Pattern: [a-zA-Z0-9][a-zA-Z0-9-]*

Required: Yes

[QueryId](#) (p. 189)

The identifier of the specific query for which you are submitting feedback. The query ID is returned in the response to the [Query](#) (p. 179) operation.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 36.

Required: Yes

RelevanceFeedbackItems (p. 189)

Provides Amazon Kendra with relevant or not relevant feedback for whether a particular item was relevant to the search.

Type: Array of [RelevanceFeedback \(p. 256\)](#) objects

Required: No

Response Elements

If the action is successful, the service sends back an HTTP 200 response with an empty HTTP body.

Errors

For information about the errors that are common to all actions, see [Common Errors \(p. 286\)](#).

AccessDeniedException

HTTP Status Code: 400

InternalServerErrorException

HTTP Status Code: 500

ResourceNotFoundException

HTTP Status Code: 400

ResourceUnavailableException

HTTP Status Code: 400

ThrottlingException

HTTP Status Code: 400

ValidationException

HTTP Status Code: 400

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

TagResource

Adds the specified tag to the specified index, FAQ, or data source resource. If the tag already exists, the existing value is replaced with the new value.

Request Syntax

```
{
  "ResourceARN": "string",
  "Tags": [
    {
      "Key": "string",
      "Value": "string"
    }
  ]
}
```

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters \(p. 288\)](#).

The request accepts the following data in JSON format.

ResourceARN (p. 191)

The Amazon Resource Name (ARN) of the index, FAQ, or data source to tag.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 1011.

Required: Yes

Tags (p. 191)

A list of tag keys to add to the index, FAQ, or data source. If a tag already exists, the existing value is replaced with the new value.

Type: Array of [Tag \(p. 283\)](#) objects

Array Members: Minimum number of 0 items. Maximum number of 200 items.

Required: Yes

Response Elements

If the action is successful, the service sends back an HTTP 200 response with an empty HTTP body.

Errors

For information about the errors that are common to all actions, see [Common Errors \(p. 286\)](#).

AccessDeniedException

HTTP Status Code: 400

InternalServerErrorException

HTTP Status Code: 500

ResourceUnavailableException

HTTP Status Code: 400

ThrottlingException

HTTP Status Code: 400

ValidationException

HTTP Status Code: 400

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

UntagResource

Removes a tag from an index, FAQ, or a data source.

Request Syntax

```
{  
  "ResourceARN": "string",  
  "TagKeys": [ "string" ]  
}
```

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters \(p. 288\)](#).

The request accepts the following data in JSON format.

ResourceARN (p. 193)

The Amazon Resource Name (ARN) of the index, FAQ, or data source to remove the tag from.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 1011.

Required: Yes

TagKeys (p. 193)

A list of tag keys to remove from the index, FAQ, or data source. If a tag key does not exist on the resource, it is ignored.

Type: Array of strings

Array Members: Minimum number of 0 items. Maximum number of 200 items.

Length Constraints: Minimum length of 1. Maximum length of 128.

Required: Yes

Response Elements

If the action is successful, the service sends back an HTTP 200 response with an empty HTTP body.

Errors

For information about the errors that are common to all actions, see [Common Errors \(p. 286\)](#).

AccessDeniedException

HTTP Status Code: 400

InternalServerError

HTTP Status Code: 500

ResourceUnavailableException

HTTP Status Code: 400

ThrottlingException

HTTP Status Code: 400

ValidationException

HTTP Status Code: 400

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

UpdateDataSource

Updates an existing Amazon Kendra data source.

Request Syntax

```
{
  "Configuration": {
    "DatabaseConfiguration": {
      "AclConfiguration": {
        "AllowedGroupsColumnName": "string"
      },
      "ColumnConfiguration": {
        "ChangeDetectingColumns": [ "string" ],
        "DocumentDataColumnName": "string",
        "DocumentIdColumnName": "string",
        "DocumentTitleColumnName": "string",
        "FieldMappings": [
          {
            "DataSourceFieldName": "string",
            "DateFieldFormat": "string",
            "IndexFieldName": "string"
          }
        ]
      },
      "ConnectionConfiguration": {
        "DatabaseHost": "string",
        "DatabaseName": "string",
        "DatabasePort": number,
        "SecretArn": "string",
        "TableName": "string"
      },
      "DatabaseEngineType": "string",
      "VpcConfiguration": {
        "SecurityGroupIds": [ "string" ],
        "SubnetIds": [ "string" ]
      }
    },
    "OneDriveConfiguration": {
      "ExclusionPatterns": [ "string" ],
      "FieldMappings": [
        {
          "DataSourceFieldName": "string",
          "DateFieldFormat": "string",
          "IndexFieldName": "string"
        }
      ],
      "InclusionPatterns": [ "string" ],
      "OneDriveUsers": {
        "OneDriveUserList": [ "string" ],
        "OneDriveUserS3Path": {
          "Bucket": "string",
          "Key": "string"
        }
      },
      "SecretArn": "string",
      "TenantDomain": "string"
    },
    "S3Configuration": {
      "AccessControlListConfiguration": {
        "KeyPath": "string"
      },
      "BucketName": "string",
```

```

    "DocumentsMetadataConfiguration": {
      "S3Prefix": "string"
    },
    "ExclusionPatterns": [ "string" ],
    "InclusionPrefixes": [ "string" ]
  },
  "SalesforceConfiguration": {
    "ChatterFeedConfiguration": {
      "DocumentDataFieldName": "string",
      "DocumentTitleFieldName": "string",
      "FieldMappings": [
        {
          "DataSourceFieldName": "string",
          "DateFieldFormat": "string",
          "IndexFieldName": "string"
        }
      ],
      "IncludeFilterTypes": [ "string" ]
    },
    "CrawlAttachments": boolean,
    "ExcludeAttachmentFilePatterns": [ "string" ],
    "IncludeAttachmentFilePatterns": [ "string" ],
    "KnowledgeArticleConfiguration": {
      "CustomKnowledgeArticleTypeConfigurations": [
        {
          "DocumentDataFieldName": "string",
          "DocumentTitleFieldName": "string",
          "FieldMappings": [
            {
              "DataSourceFieldName": "string",
              "DateFieldFormat": "string",
              "IndexFieldName": "string"
            }
          ],
          "Name": "string"
        }
      ],
      "IncludedStates": [ "string" ],
      "StandardKnowledgeArticleTypeConfiguration": {
        "DocumentDataFieldName": "string",
        "DocumentTitleFieldName": "string",
        "FieldMappings": [
          {
            "DataSourceFieldName": "string",
            "DateFieldFormat": "string",
            "IndexFieldName": "string"
          }
        ]
      }
    },
    "SecretArn": "string",
    "ServerUrl": "string",
    "StandardObjectAttachmentConfiguration": {
      "DocumentTitleFieldName": "string",
      "FieldMappings": [
        {
          "DataSourceFieldName": "string",
          "DateFieldFormat": "string",
          "IndexFieldName": "string"
        }
      ]
    },
    "StandardObjectConfigurations": [
      {
        "DocumentDataFieldName": "string",
        "DocumentTitleFieldName": "string",

```

```

        "FieldMappings": [
            {
                "DataSourceFieldName": "string",
                "DateFieldFormat": "string",
                "IndexFieldName": "string"
            }
        ],
        "Name": "string"
    }
]
},
"ServiceNowConfiguration": {
    "HostUrl": "string",
    "KnowledgeArticleConfiguration": {
        "CrawlAttachments": boolean,
        "DocumentDataFieldName": "string",
        "DocumentTitleFieldName": "string",
        "ExcludeAttachmentFilePatterns": [ "string" ],
        "FieldMappings": [
            {
                "DataSourceFieldName": "string",
                "DateFieldFormat": "string",
                "IndexFieldName": "string"
            }
        ],
        "IncludeAttachmentFilePatterns": [ "string" ]
    },
    "SecretArn": "string",
    "ServiceCatalogConfiguration": {
        "CrawlAttachments": boolean,
        "DocumentDataFieldName": "string",
        "DocumentTitleFieldName": "string",
        "ExcludeAttachmentFilePatterns": [ "string" ],
        "FieldMappings": [
            {
                "DataSourceFieldName": "string",
                "DateFieldFormat": "string",
                "IndexFieldName": "string"
            }
        ],
        "IncludeAttachmentFilePatterns": [ "string" ]
    },
    "ServiceNowBuildVersion": "string"
},
"SharePointConfiguration": {
    "CrawlAttachments": boolean,
    "DocumentTitleFieldName": "string",
    "ExclusionPatterns": [ "string" ],
    "FieldMappings": [
        {
            "DataSourceFieldName": "string",
            "DateFieldFormat": "string",
            "IndexFieldName": "string"
        }
    ],
    "InclusionPatterns": [ "string" ],
    "SecretArn": "string",
    "SharePointVersion": "string",
    "Urls": [ "string" ],
    "UseChangeLog": boolean,
    "VpcConfiguration": {
        "SecurityGroupIds": [ "string" ],
        "SubnetIds": [ "string" ]
    }
}
},
},

```

```
"Description": "string",  
"Id": "string",  
"IndexId": "string",  
"Name": "string",  
"RoleArn": "string",  
"Schedule": "string"  
}
```

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters \(p. 288\)](#).

The request accepts the following data in JSON format.

Configuration (p. 195)

Configuration information for a Amazon Kendra data source.

Type: [DataSourceConfiguration \(p. 221\)](#) object

Required: No

Description (p. 195)

The new description for the data source.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 1000.

Pattern: `^\P{C}*$`

Required: No

Id (p. 195)

The unique identifier of the data source to update.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 100.

Pattern: `[a-zA-Z0-9][a-zA-Z0-9_-]*`

Required: Yes

IndexId (p. 195)

The identifier of the index that contains the data source to update.

Type: String

Length Constraints: Fixed length of 36.

Pattern: `[a-zA-Z0-9][a-zA-Z0-9-]*`

Required: Yes

Name (p. 195)

The name of the data source to update. The name of the data source can't be updated. To rename a data source you must delete the data source and re-create it.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 1000.

Pattern: `[a-zA-Z0-9][a-zA-Z0-9_-]*`

Required: No

RoleArn (p. 195)

The Amazon Resource Name (ARN) of the new role to use when the data source is accessing resources on your behalf.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 1284.

Pattern: `arn:[a-z0-9-\.\.]{1,63}:[a-z0-9-\.\.]{0,63}:[a-z0-9-\.\.]{0,63}:[a-z0-9-\.\.]{0,63}:[^/].{0,1023}`

Required: No

Schedule (p. 195)

The new update schedule for the data source.

Type: String

Required: No

Response Elements

If the action is successful, the service sends back an HTTP 200 response with an empty HTTP body.

Errors

For information about the errors that are common to all actions, see [Common Errors \(p. 286\)](#).

AccessDeniedException

HTTP Status Code: 400

ConflictException

HTTP Status Code: 400

InternalServerErrorException

HTTP Status Code: 500

ResourceNotFoundException

HTTP Status Code: 400

ThrottlingException

HTTP Status Code: 400

ValidationException

HTTP Status Code: 400

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

UpdateIndex

Updates an existing Amazon Kendra index.

Request Syntax

```
{
  "CapacityUnits": {
    "QueryCapacityUnits": number,
    "StorageCapacityUnits": number
  },
  "Description": "string",
  "DocumentMetadataConfigurationUpdates": [
    {
      "Name": "string",
      "Relevance": {
        "Duration": "string",
        "Freshness": boolean,
        "Importance": number,
        "RankOrder": "string",
        "ValueImportanceMap": {
          "string": number
        }
      },
      "Search": {
        "Displayable": boolean,
        "Facetable": boolean,
        "Searchable": boolean
      },
      "Type": "string"
    }
  ],
  "Id": "string",
  "Name": "string",
  "RoleArn": "string"
}
```

Request Parameters

For information about the parameters that are common to all actions, see [Common Parameters](#) (p. 288).

The request accepts the following data in JSON format.

CapacityUnits (p. 201)

Sets the number of additional storage and query capacity units that should be used by the index. You can change the capacity of the index up to 5 times per day.

If you are using extra storage units, you can't reduce the storage capacity below that required to meet the storage needs for your index.

Type: [CapacityUnitsConfiguration](#) (p. 214) object

Required: No

Description (p. 201)

A new description for the index.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 1000.

Pattern: `^\P{C}*$`

Required: No

DocumentMetadataConfigurationUpdates (p. 201)

The document metadata to update.

Type: Array of [DocumentMetadataConfiguration \(p. 237\)](#) objects

Array Members: Minimum number of 0 items. Maximum number of 500 items.

Required: No

Id (p. 201)

The identifier of the index to update.

Type: String

Length Constraints: Fixed length of 36.

Pattern: `[a-zA-Z0-9][a-zA-Z0-9-]*`

Required: Yes

Name (p. 201)

The name of the index to update.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 1000.

Pattern: `[a-zA-Z0-9][a-zA-Z0-9_-]*`

Required: No

RoleArn (p. 201)

A new IAM role that gives Amazon Kendra permission to access your Amazon CloudWatch logs.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 1284.

Pattern: `arn:[a-z0-9-\.]{1,63}:[a-z0-9-\.]{0,63}:[a-z0-9-\.]{0,63}:[a-z0-9-\.]{0,63}:[^/].{0,1023}`

Required: No

Response Elements

If the action is successful, the service sends back an HTTP 200 response with an empty HTTP body.

Errors

For information about the errors that are common to all actions, see [Common Errors \(p. 286\)](#).

AccessDeniedException

HTTP Status Code: 400

ConflictException

HTTP Status Code: 400

InternalServerErrorException

HTTP Status Code: 500

ResourceNotFoundException

HTTP Status Code: 400

ServiceQuotaExceededException

HTTP Status Code: 400

ThrottlingException

HTTP Status Code: 400

ValidationException

HTTP Status Code: 400

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS Command Line Interface](#)
- [AWS SDK for .NET](#)
- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for JavaScript](#)
- [AWS SDK for PHP V3](#)
- [AWS SDK for Python](#)
- [AWS SDK for Ruby V3](#)

Data Types

The following data types are supported:

- [AccessControlListConfiguration](#) (p. 206)
- [AclConfiguration](#) (p. 207)
- [AdditionalResultAttribute](#) (p. 208)
- [AdditionalResultAttributeValue](#) (p. 209)
- [AttributeFilter](#) (p. 210)
- [BatchDeleteDocumentResponseFailedDocument](#) (p. 212)
- [BatchPutDocumentResponseFailedDocument](#) (p. 213)
- [CapacityUnitsConfiguration](#) (p. 214)
- [ClickFeedback](#) (p. 215)
- [ColumnConfiguration](#) (p. 216)
- [ConnectionConfiguration](#) (p. 218)
- [DatabaseConfiguration](#) (p. 220)

- [DataSourceConfiguration](#) (p. 221)
- [DataSourceSummary](#) (p. 223)
- [DataSourceSyncJob](#) (p. 225)
- [DataSourceSyncJobMetrics](#) (p. 227)
- [DataSourceSyncJobMetricTarget](#) (p. 229)
- [DataSourceToIndexFieldMapping](#) (p. 230)
- [DataSourceVpcConfiguration](#) (p. 231)
- [Document](#) (p. 232)
- [DocumentAttribute](#) (p. 234)
- [DocumentAttributeValue](#) (p. 235)
- [DocumentAttributeValueCountPair](#) (p. 236)
- [DocumentMetadataConfiguration](#) (p. 237)
- [DocumentsMetadataConfiguration](#) (p. 238)
- [Facet](#) (p. 239)
- [FacetResult](#) (p. 240)
- [FAQStatistics](#) (p. 241)
- [FAQSummary](#) (p. 242)
- [Highlight](#) (p. 244)
- [IndexConfigurationSummary](#) (p. 245)
- [IndexStatistics](#) (p. 247)
- [OneDriveConfiguration](#) (p. 248)
- [OneDriveUsers](#) (p. 250)
- [Principal](#) (p. 251)
- [QueryResultItem](#) (p. 252)
- [Relevance](#) (p. 254)
- [RelevanceFeedback](#) (p. 256)
- [S3DataSourceConfiguration](#) (p. 257)
- [S3Path](#) (p. 259)
- [SalesforceChatterFeedConfiguration](#) (p. 260)
- [SalesforceConfiguration](#) (p. 262)
- [SalesforceCustomKnowledgeArticleTypeConfiguration](#) (p. 265)
- [SalesforceKnowledgeArticleConfiguration](#) (p. 267)
- [SalesforceStandardKnowledgeArticleTypeConfiguration](#) (p. 268)
- [SalesforceStandardObjectAttachmentConfiguration](#) (p. 269)
- [SalesforceStandardObjectConfiguration](#) (p. 270)
- [Search](#) (p. 272)
- [ServerSideEncryptionConfiguration](#) (p. 273)
- [ServiceNowConfiguration](#) (p. 274)
- [ServiceNowKnowledgeArticleConfiguration](#) (p. 276)
- [ServiceNowServiceCatalogConfiguration](#) (p. 278)
- [SharePointConfiguration](#) (p. 280)
- [Tag](#) (p. 283)
- [TextDocumentStatistics](#) (p. 284)
- [TextWithHighlights](#) (p. 285)
- [TimeRange](#) (p. 286)

AccessControlListConfiguration

Access Control List files for the documents in a data source.

Contents

KeyPath

Path to the AWS S3 bucket that contains the ACL files.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 1024.

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

AclConfiguration

Provides information about the column that should be used for filtering the query response by groups.

Contents

AllowedGroupsColumnName

A list of groups, separated by semi-colons, that filters a query response based on user context. The document is only returned to users that are in one of the groups specified in the `UserContext` field of the [Query \(p. 179\)](#) operation.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 100.

Pattern: `^[a-zA-Z][a-zA-Z0-9_]*$`

Required: Yes

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

AdditionalResultAttribute

An attribute returned from an index query.

Contents

Key

The key that identifies the attribute.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 2048.

Required: Yes

Value

An object that contains the attribute value.

Type: [AdditionalResultAttributeValue](#) (p. 209) object

Required: Yes

ValueType

The data type of the `Value` property.

Type: String

Valid Values: `TEXT_WITH_HIGHLIGHTS_VALUE`

Required: Yes

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

AdditionalResultAttributeValue

An attribute returned with a document from a search.

Contents

TextWithHighlightsValue

The text associated with the attribute and information about the highlight to apply to the text.

Type: [TextWithHighlights](#) (p. 285) object

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

AttributeFilter

Provides filtering the query results based on document attributes.

When you use the `AndAllFilters` or `OrAllFilters`, filters you can use 2 layers under the first attribute filter. For example, you can use:

```
<AndAllFilters>
```

1. `<OrAllFilters>`
2. `<EqualTo>`

If you use more than 2 layers, you receive a `ValidationException` exception with the message "AttributeFilter cannot have a depth of more than 2."

Contents

AndAllFilters

Performs a logical AND operation on all supplied filters.

Type: Array of [AttributeFilter \(p. 210\)](#) objects

Required: No

ContainsAll

Returns true when a document contains all of the specified document attributes. This filter is only applicable to `StringListValue` metadata.

Type: [DocumentAttribute \(p. 234\)](#) object

Required: No

ContainsAny

Returns true when a document contains any of the specified document attributes. This filter is only applicable to `StringListValue` metadata.

Type: [DocumentAttribute \(p. 234\)](#) object

Required: No

EqualsTo

Performs an equals operation on two document attributes.

Type: [DocumentAttribute \(p. 234\)](#) object

Required: No

GreaterThan

Performs a greater than operation on two document attributes. Use with a document attribute of type `Integer` or `Long`.

Type: [DocumentAttribute \(p. 234\)](#) object

Required: No

GreaterThanOrEquals

Performs a greater or equals than operation on two document attributes. Use with a document attribute of type `Integer` or `Long`.

Type: [DocumentAttribute \(p. 234\)](#) object

Required: No

LessThan

Performs a less than operation on two document attributes. Use with a document attribute of type `Integer` or `Long`.

Type: [DocumentAttribute \(p. 234\)](#) object

Required: No

LessThanOrEquals

Performs a less than or equals operation on two document attributes. Use with a document attribute of type `Integer` or `Long`.

Type: [DocumentAttribute \(p. 234\)](#) object

Required: No

NotFilter

Performs a logical NOT operation on all supplied filters.

Type: [AttributeFilter \(p. 210\)](#) object

Required: No

OrAllFilters

Performs a logical OR operation on all supplied filters.

Type: Array of [AttributeFilter \(p. 210\)](#) objects

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

BatchDeleteDocumentResponseFailedDocument

Provides information about documents that could not be removed from an index by the [BatchDeleteDocument](#) (p. 126) operation.

Contents

ErrorCode

The error code for why the document couldn't be removed from the index.

Type: String

Valid Values: `InternalServerError` | `InvalidRequest`

Required: No

ErrorMessage

An explanation for why the document couldn't be removed from the index.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 2048.

Pattern: `^\P{C}*$`

Required: No

Id

The identifier of the document that couldn't be removed from the index.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 2048.

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

BatchPutDocumentResponseFailedDocument

Provides information about a document that could not be indexed.

Contents

ErrorCode

The type of error that caused the document to fail to be indexed.

Type: String

Valid Values: `InternalServerError` | `InvalidRequest`

Required: No

ErrorMessage

A description of the reason why the document could not be indexed.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 2048.

Pattern: `^\P{C}*$`

Required: No

Id

The unique identifier of the document.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 2048.

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

CapacityUnitsConfiguration

Specifies capacity units configured for your index. You can add and remove capacity units to tune an index to your requirements.

Contents

QueryCapacityUnits

The amount of extra query capacity for an index. Each capacity unit provides 0.5 queries per second and 40,000 queries per day.

Type: Integer

Valid Range: Minimum value of 0.

Required: Yes

StorageCapacityUnits

The amount of extra storage capacity for an index. Each capacity unit provides 150 Gb of storage space or 500,000 documents, whichever is reached first.

Type: Integer

Valid Range: Minimum value of 0.

Required: Yes

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

ClickFeedback

Gathers information about when a particular result was clicked by a user. Your application uses the [SubmitFeedback \(p. 189\)](#) operation to provide click information.

Contents

ClickTime

The Unix timestamp of the date and time that the result was clicked.

Type: Timestamp

Required: Yes

ResultId

The unique identifier of the search result that was clicked.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 73.

Required: Yes

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

ColumnConfiguration

Provides information about how Amazon Kendra should use the columns of a database in an index.

Contents

ChangeDetectingColumns

One to five columns that indicate when a document in the database has changed.

Type: Array of strings

Array Members: Minimum number of 1 item. Maximum number of 5 items.

Length Constraints: Minimum length of 1. Maximum length of 100.

Pattern: `^[a-zA-Z][a-zA-Z0-9_]*$`

Required: Yes

DocumentDataColumnName

The column that contains the contents of the document.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 100.

Pattern: `^[a-zA-Z][a-zA-Z0-9_]*$`

Required: Yes

DocumentIdColumnName

The column that provides the document's unique identifier.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 100.

Pattern: `^[a-zA-Z][a-zA-Z0-9_]*$`

Required: Yes

DocumentTitleColumnName

The column that contains the title of the document.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 100.

Pattern: `^[a-zA-Z][a-zA-Z0-9_]*$`

Required: No

FieldMappings

An array of objects that map database column names to the corresponding fields in an index. You must first create the fields in the index using the [UpdateIndex \(p. 201\)](#) operation.

Type: Array of [DataSourceToIndexFieldMapping \(p. 230\)](#) objects

Array Members: Minimum number of 1 item. Maximum number of 100 items.

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

ConnectionConfiguration

Provides the information necessary to connect to a database.

Contents

DatabaseHost

The name of the host for the database. Can be either a string (host.subdomain.domain.tld) or an IPv4 or IPv6 address.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 253.

Required: Yes

DatabaseName

The name of the database containing the document data.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 100.

Pattern: `^[a-zA-Z][a-zA-Z0-9_]*$`

Required: Yes

DatabasePort

The port that the database uses for connections.

Type: Integer

Valid Range: Minimum value of 1. Maximum value of 65535.

Required: Yes

SecretArn

The Amazon Resource Name (ARN) of credentials stored in AWS Secrets Manager. The credentials should be a user/password pair. For more information, see [Using a Database Data Source](#). For more information about AWS Secrets Manager, see [What Is AWS Secrets Manager](#) in the *AWS Secrets Manager* user guide.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 1284.

Pattern: `arn:[a-z0-9-\.]{1,63}:[a-z0-9-\.]{0,63}:[a-z0-9-\.]{0,63}:[a-z0-9-\.]{0,63}:[^/].{0,1023}`

Required: Yes

TableName

The name of the table that contains the document data.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 100.

Pattern: `^[a-zA-Z][a-zA-Z0-9_]*$`

Required: Yes

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

DatabaseConfiguration

Provides the information necessary to connect a database to an index.

Contents

AclConfiguration

Information about the database column that provides information for user context filtering.

Type: [AclConfiguration \(p. 207\)](#) object

Required: No

ColumnConfiguration

Information about where the index should get the document information from the database.

Type: [ColumnConfiguration \(p. 216\)](#) object

Required: Yes

ConnectionConfiguration

The information necessary to connect to a database.

Type: [ConnectionConfiguration \(p. 218\)](#) object

Required: Yes

DatabaseEngineType

The type of database engine that runs the database.

Type: String

Valid Values: `RDS_AURORA_MYSQL` | `RDS_AURORA_POSTGRESQL` | `RDS_MYSQL` | `RDS_POSTGRESQL`

Required: Yes

VpcConfiguration

Provides information for connecting to an Amazon VPC.

Type: [DataSourceVpcConfiguration \(p. 231\)](#) object

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

DataSourceConfiguration

Configuration information for a Amazon Kendra data source.

Contents

DatabaseConfiguration

Provides information necessary to create a connector for a database.

Type: [DatabaseConfiguration](#) (p. 220) object

Required: No

OneDriveConfiguration

Provided configuration for data sources that connect to Microsoft OneDrive.

Type: [OneDriveConfiguration](#) (p. 248) object

Required: No

S3Configuration

Provides information to create a connector for a document repository in an Amazon S3 bucket.

Type: [S3DataSourceConfiguration](#) (p. 257) object

Required: No

SalesforceConfiguration

Provides configuration information for data sources that connect to a Salesforce site.

Type: [SalesforceConfiguration](#) (p. 262) object

Required: No

ServiceNowConfiguration

Provides configuration for data sources that connect to ServiceNow instances.

Type: [ServiceNowConfiguration](#) (p. 274) object

Required: No

SharePointConfiguration

Provides information necessary to create a connector for a Microsoft SharePoint site.

Type: [SharePointConfiguration](#) (p. 280) object

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)

- [AWS SDK for Ruby V3](#)

DataSourceSummary

Summary information for a Amazon Kendra data source. Returned in a call to [DescribeDataSource](#) (p. 151).

Contents

CreatedAt

The UNIX datetime that the data source was created.

Type: Timestamp

Required: No

Id

The unique identifier for the data source.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 100.

Pattern: [a-zA-Z0-9][a-zA-Z0-9_-]*

Required: No

Name

The name of the data source.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 1000.

Pattern: [a-zA-Z0-9][a-zA-Z0-9_-]*

Required: No

Status

The status of the data source. When the status is `ACTIVE` the data source is ready to use.

Type: String

Valid Values: `CREATING` | `DELETING` | `FAILED` | `UPDATING` | `ACTIVE`

Required: No

Type

The type of the data source.

Type: String

Valid Values: `S3` | `SHAREPOINT` | `DATABASE` | `SALESFORCE` | `ONEDRIVE` | `SERVICENOW`

Required: No

UpdatedAt

The UNIX datetime that the data source was last updated.

Type: Timestamp

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

DataSourceSyncJob

Provides information about a synchronization job.

Contents

DataSourceErrorCode

If the reason that the synchronization failed is due to an error with the underlying data source, this field contains a code that identifies the error.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 2048.

Required: No

EndTime

The UNIX datetime that the synchronization job was completed.

Type: Timestamp

Required: No

ErrorCode

If the `Status` field is set to `FAILED`, the `ErrorCode` field contains a the reason that the synchronization failed.

Type: String

Valid Values: `InternalError` | `InvalidRequest`

Required: No

ErrorMessage

If the `Status` field is set to `ERROR`, the `ErrorMessage` field contains a description of the error that caused the synchronization to fail.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 2048.

Pattern: `^\P{C}*$`

Required: No

ExecutionId

A unique identifier for the synchronization job.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 2048.

Required: No

Metrics

Maps a batch delete document request to a specific data source sync job. This is optional and should only be supplied when documents are deleted by a connector.

Type: [DataSourceSyncJobMetrics](#) (p. 227) object

Required: No

StartTime

The UNIX datetime that the synchronization job was started.

Type: Timestamp

Required: No

Status

The execution status of the synchronization job. When the `Status` field is set to `SUCCEEDED`, the synchronization job is done. If the status code is set to `FAILED`, the `ErrorCode` and `ErrorMessage` fields give you the reason for the failure.

Type: String

Valid Values: `FAILED` | `SUCCEEDED` | `SYNCING` | `INCOMPLETE` | `STOPPING` | `ABORTED` | `SYNCING_INDEXING`

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

DataSourceSyncJobMetrics

Maps a batch delete document request to a specific data source sync job. This is optional and should only be supplied when documents are deleted by a connector.

Contents

DocumentsAdded

The number of documents added from the data source up to now in the data source sync.

Type: String

Pattern: `(([1-9][0-9]*)|0)`

Required: No

DocumentsDeleted

The number of documents deleted from the data source up to now in the data source sync run.

Type: String

Pattern: `(([1-9][0-9]*)|0)`

Required: No

DocumentsFailed

The number of documents that failed to sync from the data source up to now in the data source sync run.

Type: String

Pattern: `(([1-9][0-9]*)|0)`

Required: No

DocumentsModified

The number of documents modified in the data source up to now in the data source sync run.

Type: String

Pattern: `(([1-9][0-9]*)|0)`

Required: No

DocumentsScanned

The current number of documents crawled by the current sync job in the data source.

Type: String

Pattern: `(([1-9][0-9]*)|0)`

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

DataSourceSyncJobMetricTarget

Maps a particular data source sync job to a particular data source.

Contents

DataSourceId

The ID of the data source that is running the sync job.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 100.

Pattern: `[a-zA-Z0-9][a-zA-Z0-9_-]*`

Required: Yes

DataSourceSyncJobId

The ID of the sync job that is running on the data source.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 100.

Pattern: `[a-zA-Z0-9][a-zA-Z0-9_-]*`

Required: Yes

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

DataSourceToIndexFieldMapping

Maps a column or attribute in the data source to an index field. You must first create the fields in the index using the [UpdateIndex](#) (p. 201) operation.

Contents

DataSourceFieldName

The name of the column or attribute in the data source.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 100.

Pattern: `^[a-zA-Z][a-zA-Z0-9_]*$`

Required: Yes

DateFieldFormat

The type of data stored in the column or attribute.

Type: String

Length Constraints: Minimum length of 4. Maximum length of 40.

Pattern: `^(?!\\s)\\.*(?!\\s)$`

Required: No

IndexFieldName

The name of the field in the index.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 30.

Pattern: `^\\P{C}*$`

Required: Yes

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

DataSourceVpcConfiguration

Provides information for connecting to an Amazon VPC.

Contents

SecurityGroupIds

A list of identifiers of security groups within your Amazon VPC. The security groups should enable Amazon Kendra to connect to the data source.

Type: Array of strings

Array Members: Minimum number of 1 item. Maximum number of 10 items.

Length Constraints: Minimum length of 1. Maximum length of 200.

Pattern: [-0-9a-zA-Z] +

Required: Yes

SubnetIds

A list of identifiers for subnets within your Amazon VPC. The subnets should be able to connect to each other in the VPC, and they should have outgoing access to the Internet through a NAT device.

Type: Array of strings

Array Members: Minimum number of 1 item. Maximum number of 6 items.

Length Constraints: Minimum length of 1. Maximum length of 200.

Pattern: [\-0-9a-zA-Z] +

Required: Yes

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

Document

A document in an index.

Contents

AccessControlList

Information to use for user context filtering.

Type: Array of [Principal \(p. 251\)](#) objects

Required: No

Attributes

Custom attributes to apply to the document. Use the custom attributes to provide additional information for searching, to provide facets for refining searches, and to provide additional information in the query response.

Type: Array of [DocumentAttribute \(p. 234\)](#) objects

Required: No

Blob

The contents of the document.

Documents passed to the `Blob` parameter must be base64 encoded. Your code might not need to encode the document file bytes if you're using an AWS SDK to call Amazon Kendra operations. If you are calling the Amazon Kendra endpoint directly using REST, you must base64 encode the contents before sending.

Type: Base64-encoded binary data object

Required: No

ContentType

The file type of the document in the `Blob` field.

Type: String

Valid Values: `PDF` | `HTML` | `MS_WORD` | `PLAIN_TEXT` | `PPT`

Required: No

Id

A unique identifier of the document in the index.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 2048.

Required: Yes

S3Path

Information required to find a specific file in an Amazon S3 bucket.

Type: [S3Path \(p. 259\)](#) object

Required: No

Title

The title of the document.

Type: String

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

DocumentAttribute

A custom attribute value assigned to a document.

Contents

Key

The identifier for the attribute.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 200.

Pattern: `[a-zA-Z0-9_][a-zA-Z0-9_-]*`

Required: Yes

Value

The value of the attribute.

Type: [DocumentAttributeValue \(p. 235\)](#) object

Required: Yes

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

DocumentAttributeValue

The value of a custom document attribute. You can only provide one value for a custom attribute.

Contents

DateValue

A date expressed as an ISO 8601 string.

Type: Timestamp

Required: No

LongValue

A long integer value.

Type: Long

Required: No

StringListValue

A list of strings.

Type: Array of strings

Length Constraints: Minimum length of 1. Maximum length of 2048.

Required: No

StringValue

A string, such as "department".

Type: String

Length Constraints: Minimum length of 1. Maximum length of 2048.

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

DocumentAttributeValueCountPair

Provides the count of documents that match a particular attribute when doing a faceted search.

Contents

Count

The number of documents in the response that have the attribute value for the key.

Type: Integer

Required: No

DocumentAttributeValue

The value of the attribute. For example, "HR."

Type: [DocumentAttributeValue](#) (p. 235) object

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

DocumentMetadataConfiguration

Specifies the properties of a custom index field.

Contents

Name

The name of the index field.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 30.

Required: Yes

Relevance

Provides manual tuning parameters to determine how the field affects the search results.

Type: [Relevance \(p. 254\)](#) object

Required: No

Search

Provides information about how the field is used during a search.

Type: [Search \(p. 272\)](#) object

Required: No

Type

The data type of the index field.

Type: String

Valid Values: `STRING_VALUE` | `STRING_LIST_VALUE` | `LONG_VALUE` | `DATE_VALUE`

Required: Yes

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

DocumentsMetadataConfiguration

Document metadata files that contain information such as the document access control information, source URI, document author, and custom attributes. Each metadata file contains metadata about a single document.

Contents

S3Prefix

A prefix used to filter metadata configuration files in the AWS S3 bucket. The S3 bucket might contain multiple metadata files. Use `S3Prefix` to include only the desired metadata files.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 1024.

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

Facet

Information about a document attribute

Contents

DocumentAttributeKey

The unique key for the document attribute.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 200.

Pattern: [a-zA-Z0-9_][a-zA-Z0-9_-]*

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

FacetResult

The facet values for the documents in the response.

Contents

DocumentAttributeKey

The key for the facet values. This is the same as the `DocumentAttributeKey` provided in the query.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 200.

Pattern: `[a-zA-Z0-9_][a-zA-Z0-9_-]*`

Required: No

DocumentAttributeValueCountPairs

An array of key/value pairs, where the key is the value of the attribute and the count is the number of documents that share the key value.

Type: Array of [DocumentAttributeValueCountPair](#) (p. 236) objects

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

FaqStatistics

Provides statistical information about the FAQ questions and answers contained in an index.

Contents

IndexedQuestionAnswersCount

The total number of FAQ questions and answers contained in the index.

Type: Integer

Valid Range: Minimum value of 0.

Required: Yes

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

FaqSummary

Provides information about a frequently asked questions and answer contained in an index.

Contents

CreatedAt

The UNIX datetime that the FAQ was added to the index.

Type: Timestamp

Required: No

Id

The unique identifier of the FAQ.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 100.

Pattern: [a-zA-Z0-9][a-zA-Z0-9_-]*

Required: No

Name

The name that you assigned the FAQ when you created or updated the FAQ.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 100.

Pattern: [a-zA-Z0-9][a-zA-Z0-9_-]*

Required: No

Status

The current status of the FAQ. When the status is `ACTIVE` the FAQ is ready for use.

Type: String

Valid Values: `CREATING` | `UPDATING` | `ACTIVE` | `DELETING` | `FAILED`

Required: No

UpdatedAt

The UNIX datetime that the FAQ was last updated.

Type: Timestamp

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)

- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

Highlight

Provides information that you can use to highlight a search result so that your users can quickly identify terms in the response.

Contents

BeginOffset

The zero-based location in the response string where the highlight starts.

Type: Integer

Required: Yes

EndOffset

The zero-based location in the response string where the highlight ends.

Type: Integer

Required: Yes

TopAnswer

Indicates whether the response is the best response. True if this is the best response; otherwise, false.

Type: Boolean

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

IndexConfigurationSummary

A summary of information about an index.

Contents

CreatedAt

The Unix timestamp when the index was created.

Type: Timestamp

Required: Yes

Edition

Indicates whether the index is a enterprise edition index or a developer edition index.

Type: String

Valid Values: `DEVELOPER_EDITION` | `ENTERPRISE_EDITION`

Required: No

Id

A unique identifier for the index. Use this to identify the index when you are using operations such as `Query`, `DescribeIndex`, `UpdateIndex`, and `DeleteIndex`.

Type: String

Length Constraints: Fixed length of 36.

Pattern: `[a-zA-Z0-9][a-zA-Z0-9-]*`

Required: No

Name

The name of the index.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 1000.

Pattern: `[a-zA-Z0-9][a-zA-Z0-9_-]*`

Required: No

Status

The current status of the index. When the status is `ACTIVE`, the index is ready to search.

Type: String

Valid Values: `CREATING` | `ACTIVE` | `DELETING` | `FAILED` | `UPDATING` | `SYSTEM_UPDATING`

Required: Yes

UpdatedAt

The Unix timestamp when the index was last updated by the `UpdateIndex` operation.

Type: Timestamp

Required: Yes

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

IndexStatistics

Provides information about the number of documents and the number of questions and answers in an index.

Contents

FaqStatistics

The number of question and answer topics in the index.

Type: [FaqStatistics \(p. 241\)](#) object

Required: Yes

TextDocumentStatistics

The number of text documents indexed.

Type: [TextDocumentStatistics \(p. 284\)](#) object

Required: Yes

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

OneDriveConfiguration

Provides configuration information for data sources that connect to OneDrive.

Contents

ExclusionPatterns

List of regular expressions applied to documents. Items that match the exclusion pattern are not indexed. If you provide both an inclusion pattern and an exclusion pattern, any item that matches the exclusion pattern isn't indexed.

The exclusion pattern is applied to the file name.

Type: Array of strings

Array Members: Minimum number of 0 items. Maximum number of 100 items.

Length Constraints: Minimum length of 1. Maximum length of 50.

Required: No

FieldMappings

A list of `DataSourceToIndexFieldMapping` objects that map Microsoft OneDrive fields to custom fields in the Amazon Kendra index. You must first create the index fields before you map OneDrive fields.

Type: Array of [DataSourceToIndexFieldMapping \(p. 230\)](#) objects

Array Members: Minimum number of 1 item. Maximum number of 100 items.

Required: No

InclusionPatterns

A list of regular expression patterns. Documents that match the pattern are included in the index. Documents that don't match the pattern are excluded from the index. If a document matches both an inclusion pattern and an exclusion pattern, the document is not included in the index.

The exclusion pattern is applied to the file name.

Type: Array of strings

Array Members: Minimum number of 0 items. Maximum number of 100 items.

Length Constraints: Minimum length of 1. Maximum length of 50.

Required: No

OneDriveUsers

A list of user accounts whose documents should be indexed.

Type: [OneDriveUsers \(p. 250\)](#) object

Required: Yes

SecretArn

The Amazon Resource Name (ARN) of an AWS Secrets Manager secret that contains the user name and password to connect to OneDrive. The user name should be the application ID for the OneDrive application, and the password is the application key for the OneDrive application.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 1284.

Pattern: `arn:[a-z0-9-\.]{1,63}:[a-z0-9-\.]{0,63}:[a-z0-9-\.]{0,63}:[a-z0-9-\.]{0,63}:[^/].{0,1023}`

Required: Yes

TenantDomain

The Azure Active Directory domain of the organization.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 256.

Pattern: `^([a-zA-Z0-9]+(-[a-zA-Z0-9]+)*\.)+[a-z]{2,}$`

Required: Yes

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

OneDriveUsers

User accounts whose documents should be indexed.

Contents

OneDriveUserList

A list of users whose documents should be indexed. Specify the user names in email format, for example, `username@tenantdomain`. If you need to index the documents of more than 100 users, use the `OneDriveUserS3Path` field to specify the location of a file containing a list of users.

Type: Array of strings

Array Members: Minimum number of 1 item. Maximum number of 100 items.

Length Constraints: Minimum length of 1. Maximum length of 256.

Pattern: `^(?!\\s)\\.+@([a-zA-Z0-9_\\-\\.]+)\\.([a-zA-Z]{2,5})$`

Required: No

OneDriveUserS3Path

The S3 bucket location of a file containing a list of users whose documents should be indexed.

Type: [S3Path](#) (p. 259) object

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

Principal

Provides user and group information for document access filtering.

Contents

Access

Whether to allow or deny access to the principal.

Type: String

Valid Values: `ALLOW` | `DENY`

Required: Yes

Name

The name of the user or group.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 200.

Pattern: `^\P{C}*$`

Required: Yes

Type

The type of principal.

Type: String

Valid Values: `USER` | `GROUP`

Required: Yes

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

QueryResultItem

A single query result.

A query result contains information about a document returned by the query. This includes the original location of the document, a list of attributes assigned to the document, and relevant text from the document that satisfies the query.

Contents

AdditionalAttributes

One or more additional attributes associated with the query result.

Type: Array of [AdditionalResultAttribute \(p. 208\)](#) objects

Required: No

DocumentAttributes

An array of document attributes for the document that the query result maps to. For example, the document author (Author) or the source URI (SourceUri) of the document.

Type: Array of [DocumentAttribute \(p. 234\)](#) objects

Required: No

DocumentExcerpt

An extract of the text in the document. Contains information about highlighting the relevant terms in the excerpt.

Type: [TextWithHighlights \(p. 285\)](#) object

Required: No

DocumentId

The unique identifier for the document.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 2048.

Required: No

DocumentTitle

The title of the document. Contains the text of the title and information for highlighting the relevant terms in the title.

Type: [TextWithHighlights \(p. 285\)](#) object

Required: No

DocumentURI

The URI of the original location of the document.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 2048.

Pattern: `^(https?|ftp|file):\\/(\\[\\^\\s\\]*\\)`

Required: No

Id

The unique identifier for the query result.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 73.

Required: No

Type

The type of document.

Type: String

Valid Values: DOCUMENT | QUESTION_ANSWER | ANSWER

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

Relevance

Provides information for manually tuning the relevance of a field in a search. When a query includes terms that match the field, the results are given a boost in the response based on these tuning parameters.

Contents

Duration

Specifies the time period that the boost applies to. For example, to make the boost apply to documents with the field value within the last month, you would use "2628000s". Once the field value is beyond the specified range, the effect of the boost drops off. The higher the importance, the faster the effect drops off. If you don't specify a value, the default is 3 months. The value of the field is a numeric string followed by the character "s", for example "86400s" for one day, or "604800s" for one week.

Only applies to `DATE` fields.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 10.

Pattern: `[0-9]+[s]`

Required: No

Freshness

Indicates that this field determines how "fresh" a document is. For example, if document 1 was created on November 5, and document 2 was created on October 31, document 1 is "fresher" than document 2. You can only set the `Freshness` field on one `DATE` type field. Only applies to `DATE` fields.

Type: Boolean

Required: No

Importance

The relative importance of the field in the search. Larger numbers provide more of a boost than smaller numbers.

Type: Integer

Valid Range: Minimum value of 1. Maximum value of 10.

Required: No

RankOrder

Determines how values should be interpreted.

When the `RankOrder` field is `ASCENDING`, higher numbers are better. For example, a document with a rating score of 10 is higher ranking than a document with a rating score of 1.

When the `RankOrder` field is `DESCENDING`, lower numbers are better. For example, in a task tracking application, a priority 1 task is more important than a priority 5 task.

Only applies to `LONG` and `DOUBLE` fields.

Type: String

Valid Values: ASCENDING | DESCENDING

Required: No

ValueImportanceMap

A list of values that should be given a different boost when they appear in the result list. For example, if you are boosting a field called "department," query terms that match the department field are boosted in the result. However, you can add entries from the department field to boost documents with those values higher.

For example, you can add entries to the map with names of departments. If you add "HR",5 and "Legal",3 those departments are given special attention when they appear in the metadata of a document. When those terms appear they are given the specified importance instead of the regular importance for the boost.

Type: String to integer map

Key Length Constraints: Minimum length of 1. Maximum length of 50.

Valid Range: Minimum value of 1. Maximum value of 10.

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

RelevanceFeedback

Provides feedback on how relevant a document is to a search. Your application uses the [SubmitFeedback \(p. 189\)](#) operation to provide relevance information.

Contents

RelevanceValue

Whether to document was relevant or not relevant to the search.

Type: String

Valid Values: `RELEVANT` | `NOT_RELEVANT`

Required: Yes

ResultId

The unique identifier of the search result that the user provided relevance feedback for.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 73.

Required: Yes

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

S3DataSourceConfiguration

Provides configuration information for a data source to index documents in an Amazon S3 bucket.

Contents

AccessControlListConfiguration

Provides the path to the S3 bucket that contains the user context filtering files for the data source.

Type: [AccessControlListConfiguration](#) (p. 206) object

Required: No

BucketName

The name of the bucket that contains the documents.

Type: String

Length Constraints: Minimum length of 3. Maximum length of 63.

Pattern: `[a-z0-9][\.\-a-z0-9]{1,61}[a-z0-9]`

Required: Yes

DocumentsMetadataConfiguration

Document metadata files that contain information such as the document access control information, source URI, document author, and custom attributes. Each metadata file contains metadata about a single document.

Type: [DocumentsMetadataConfiguration](#) (p. 238) object

Required: No

ExclusionPatterns

A list of glob patterns for documents that should not be indexed. If a document that matches an inclusion prefix also matches an exclusion pattern, the document is not indexed.

For more information about glob patterns, see [glob \(programming\)](#) in *Wikipedia*.

Type: Array of strings

Array Members: Minimum number of 0 items. Maximum number of 100 items.

Length Constraints: Minimum length of 1. Maximum length of 50.

Required: No

InclusionPrefixes

A list of S3 prefixes for the documents that should be included in the index.

Type: Array of strings

Array Members: Minimum number of 0 items. Maximum number of 100 items.

Length Constraints: Minimum length of 1. Maximum length of 50.

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

S3Path

Information required to find a specific file in an Amazon S3 bucket.

Contents

Bucket

The name of the S3 bucket that contains the file.

Type: String

Length Constraints: Minimum length of 3. Maximum length of 63.

Pattern: `[a-z0-9][\.\-a-z0-9]{1,61}[a-z0-9]`

Required: Yes

Key

The name of the file.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 1024.

Required: Yes

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

SalesforceChatterFeedConfiguration

Defines configuration for syncing a Salesforce chatter feed. The contents of the object comes from the Salesforce FeedItem table.

Contents

DocumentDataFieldName

The name of the column in the Salesforce FeedItem table that contains the content to index. Typically this is the `Body` column.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 100.

Pattern: `^[a-zA-Z][a-zA-Z0-9_\.]*$`

Required: Yes

DocumentTitleFieldName

The name of the column in the Salesforce FeedItem table that contains the title of the document. This is typically the `Title` column.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 100.

Pattern: `^[a-zA-Z][a-zA-Z0-9_\.]*$`

Required: No

FieldMappings

Maps fields from a Salesforce chatter feed into Amazon Kendra index fields.

Type: Array of [DataSourceToIndexFieldMapping \(p. 230\)](#) objects

Array Members: Minimum number of 1 item. Maximum number of 100 items.

Required: No

IncludeFilterTypes

Filters the documents in the feed based on status of the user. When you specify `ACTIVE_USERS` only documents from users who have an active account are indexed. When you specify `STANDARD_USER` only documents for Salesforce standard users are documented. You can specify both.

Type: Array of strings

Array Members: Minimum number of 1 item. Maximum number of 2 items.

Valid Values: `ACTIVE_USER` | `STANDARD_USER`

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

SalesforceConfiguration

Provides configuration information for connecting to a Salesforce data source.

Contents

ChatterFeedConfiguration

Specifies configuration information for Salesforce chatter feeds.

Type: [SalesforceChatterFeedConfiguration](#) (p. 260) object

Required: No

CrawlAttachments

Indicates whether Amazon Kendra should index attachments to Salesforce objects.

Type: Boolean

Required: No

ExcludeAttachmentFilePatterns

A list of regular expression patterns. Documents that match the patterns are excluded from the index. Documents that don't match the patterns are included in the index. If a document matches both an exclusion pattern and an inclusion pattern, the document is not included in the index.

The regex is applied to the name of the attached file.

Type: Array of strings

Array Members: Minimum number of 0 items. Maximum number of 100 items.

Length Constraints: Minimum length of 1. Maximum length of 50.

Required: No

IncludeAttachmentFilePatterns

A list of regular expression patterns. Documents that match the patterns are included in the index. Documents that don't match the patterns are excluded from the index. If a document matches both an inclusion pattern and an exclusion pattern, the document is not included in the index.

The regex is applied to the name of the attached file.

Type: Array of strings

Array Members: Minimum number of 0 items. Maximum number of 100 items.

Length Constraints: Minimum length of 1. Maximum length of 50.

Required: No

KnowledgeArticleConfiguration

Specifies configuration information for the knowledge article types that Amazon Kendra indexes. Amazon Kendra indexes standard knowledge articles and the standard fields of knowledge articles, or the custom fields of custom knowledge articles, but not both.

Type: [SalesforceKnowledgeArticleConfiguration](#) (p. 267) object

Required: No

SecretArn

The Amazon Resource Name (ARN) of an AWS Secrets Manager secret that contains the key/value pairs required to connect to your Salesforce instance. The secret must contain a JSON structure with the following keys:

- authenticationUrl - The OAUTH endpoint that Amazon Kendra connects to get an OAUTH token.
- consumerKey - The application public key generated when you created your Salesforce application.
- consumerSecret - The application private key generated when you created your Salesforce application.
- password - The password associated with the user logging in to the Salesforce instance.
- securityToken - The token associated with the user account logging in to the Salesforce instance.
- username - The user name of the user logging in to the Salesforce instance.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 1284.

Pattern: `arn:[a-z0-9-\.]{1,63}:[a-z0-9-\.]{0,63}:[a-z0-9-\.]{0,63}:[a-z0-9-\.]{0,63}:[^/]{0,1023}`

Required: Yes

ServerUrl

The instance URL for the Salesforce site that you want to index.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 2048.

Pattern: `^(https?|ftp|file):\/\/([^\s]*)`

Required: Yes

StandardObjectAttachmentConfiguration

Provides configuration information for processing attachments to Salesforce standard objects.

Type: [SalesforceStandardObjectAttachmentConfiguration \(p. 269\)](#) object

Required: No

StandardObjectConfigurations

Specifies the Salesforce standard objects that Amazon Kendra indexes.

Type: Array of [SalesforceStandardObjectConfiguration \(p. 270\)](#) objects

Array Members: Minimum number of 1 item. Maximum number of 17 items.

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)

- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

SalesforceCustomKnowledgeArticleTypeConfiguration

Provides configuration information for indexing Salesforce custom articles.

Contents

DocumentDataFieldName

The name of the field in the custom knowledge article that contains the document data to index.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 100.

Pattern: `^[a-zA-Z][a-zA-Z0-9_]*$`

Required: Yes

DocumentTitleFieldName

The name of the field in the custom knowledge article that contains the document title.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 100.

Pattern: `^[a-zA-Z][a-zA-Z0-9_]*$`

Required: No

FieldMappings

One or more objects that map fields in the custom knowledge article to fields in the Amazon Kendra index.

Type: Array of [DataSourceToIndexFieldMapping \(p. 230\)](#) objects

Array Members: Minimum number of 1 item. Maximum number of 100 items.

Required: No

Name

The name of the configuration.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 100.

Pattern: `^[a-zA-Z][a-zA-Z0-9_]*$`

Required: Yes

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)

- [AWS SDK for Ruby V3](#)

SalesforceKnowledgeArticleConfiguration

Specifies configuration information for the knowledge article types that Amazon Kendra indexes. Amazon Kendra indexes standard knowledge articles and the standard fields of knowledge articles, or the custom fields of custom knowledge articles, but not both

Contents

CustomKnowledgeArticleTypeConfigurations

Provides configuration information for custom Salesforce knowledge articles.

Type: Array of [SalesforceCustomKnowledgeArticleTypeConfiguration \(p. 265\)](#) objects

Array Members: Minimum number of 1 item. Maximum number of 10 items.

Required: No

IncludedStates

Specifies the document states that should be included when Amazon Kendra indexes knowledge articles. You must specify at least one state.

Type: Array of strings

Array Members: Minimum number of 1 item. Maximum number of 3 items.

Valid Values: DRAFT | PUBLISHED | ARCHIVED

Required: Yes

StandardKnowledgeArticleTypeConfiguration

Provides configuration information for standard Salesforce knowledge articles.

Type: [SalesforceStandardKnowledgeArticleTypeConfiguration \(p. 268\)](#) object

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

SalesforceStandardKnowledgeArticleTypeConfiguration

Provides configuration information for standard Salesforce knowledge articles.

Contents

DocumentDataFieldName

The name of the field that contains the document data to index.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 100.

Pattern: `^[a-zA-Z][a-zA-Z0-9_\.]*$`

Required: Yes

DocumentTitleFieldName

The name of the field that contains the document title.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 100.

Pattern: `^[a-zA-Z][a-zA-Z0-9_\.]*$`

Required: No

FieldMappings

One or more objects that map fields in the knowledge article to Amazon Kendra index fields. The index field must exist before you can map a Salesforce field to it.

Type: Array of [DataSourceToIndexFieldMapping \(p. 230\)](#) objects

Array Members: Minimum number of 1 item. Maximum number of 100 items.

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

SalesforceStandardObjectAttachmentConfiguration

Provides configuration information for processing attachments to Salesforce standard objects.

Contents

DocumentTitleFieldName

The name of the field used for the document title.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 100.

Pattern: `^[a-zA-Z][a-zA-Z0-9_.]*$`

Required: No

FieldMappings

One or more objects that map fields in attachments to Amazon Kendra index fields.

Type: Array of [DataSourceToIndexFieldMapping \(p. 230\)](#) objects

Array Members: Minimum number of 1 item. Maximum number of 100 items.

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

SalesforceStandardObjectConfiguration

Specifies configuration information for indexing a single standard object.

Contents

DocumentDataFieldName

The name of the field in the standard object table that contains the document contents.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 100.

Pattern: `^[a-zA-Z][a-zA-Z0-9_]*$`

Required: Yes

DocumentTitleFieldName

The name of the field in the standard object table that contains the document titleB.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 100.

Pattern: `^[a-zA-Z][a-zA-Z0-9_]*$`

Required: No

FieldMappings

One or more objects that map fields in the standard object to Amazon Kendra index fields. The index field must exist before you can map a Salesforce field to it.

Type: Array of [DataSourceToIndexFieldMapping \(p. 230\)](#) objects

Array Members: Minimum number of 1 item. Maximum number of 100 items.

Required: No

Name

The name of the standard object.

Type: String

Valid Values: ACCOUNT | CAMPAIGN | CASE | CONTACT | CONTRACT | DOCUMENT | GROUP
| IDEA | LEAD | OPPORTUNITY | PARTNER | PRICEBOOK | PRODUCT | PROFILE |
SOLUTION | TASK | USER

Required: Yes

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)

- [AWS SDK for Ruby V3](#)

Search

Provides information about how a custom index field is used during a search.

Contents

Displayable

Determines whether the field is returned in the query response. The default is `true`.

Type: Boolean

Required: No

Facetable

Indicates that the field can be used to create search facets, a count of results for each value in the field. The default is `false`.

Type: Boolean

Required: No

Searchable

Determines whether the field is used in the search. If the `Searchable` field is `true`, you can use relevance tuning to manually tune how Amazon Kendra weights the field in the search. The default is `true` for string fields and `false` for number and date fields.

Type: Boolean

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

ServerSideEncryptionConfiguration

Provides the identifier of the AWS KMS customer master key (CMK) used to encrypt data indexed by Amazon Kendra. Amazon Kendra doesn't support asymmetric CMKs.

Contents

KmsKeyId

The identifier of the AWS KMS customer master key (CMK). Amazon Kendra doesn't support asymmetric CMKs.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 2048.

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

ServiceNowConfiguration

Provides configuration information required to connect to a ServiceNow data source.

Contents

HostUrl

The ServiceNow instance that the data source connects to. The host endpoint should look like the following: `{instance}.service-now.com`.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 2048.

Pattern: `^(?!^(https?|ftp|file):\\\/\\\/)[a-z0-9-]+(\\.service-now\\.com)$`

Required: Yes

KnowledgeArticleConfiguration

Provides configuration information for crawling knowledge articles in the ServiceNow site.

Type: [ServiceNowKnowledgeArticleConfiguration \(p. 276\)](#) object

Required: No

SecretArn

The Amazon Resource Name (ARN) of the AWS Secret Manager secret that contains the user name and password required to connect to the ServiceNow instance.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 1284.

Pattern: `arn:[a-z0-9-\\.]{1,63}:[a-z0-9-\\.]{0,63}:[a-z0-9-\\.]{0,63}:[a-z0-9-\\.]{0,63}:[^/]{0,1023}`

Required: Yes

ServiceCatalogConfiguration

Provides configuration information for crawling service catalogs in the ServiceNow site.

Type: [ServiceNowServiceCatalogConfiguration \(p. 278\)](#) object

Required: No

ServiceNowBuildVersion

The identifier of the release that the ServiceNow host is running. If the host is not running the LONDON release, use OTHERS.

Type: String

Valid Values: LONDON | OTHERS

Required: Yes

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

ServiceNowKnowledgeArticleConfiguration

Provides configuration information for crawling knowledge articles in the ServiceNow site.

Contents

CrawlAttachments

Indicates whether Amazon Kendra should index attachments to knowledge articles.

Type: Boolean

Required: No

DocumentDataFieldName

The name of the ServiceNow field that is mapped to the index document contents field in the Amazon Kendra index.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 100.

Pattern: `^[a-zA-Z][a-zA-Z0-9_\.]*$`

Required: Yes

DocumentTitleFieldName

The name of the ServiceNow field that is mapped to the index document title field.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 100.

Pattern: `^[a-zA-Z][a-zA-Z0-9_\.]*$`

Required: No

ExcludeAttachmentFilePatterns

List of regular expressions applied to knowledge articles. Items that don't match the inclusion pattern are not indexed. The regex is applied to the field specified in the `PatternTargetField`

Type: Array of strings

Array Members: Minimum number of 0 items. Maximum number of 100 items.

Length Constraints: Minimum length of 1. Maximum length of 50.

Required: No

FieldMappings

Mapping between ServiceNow fields and Amazon Kendra index fields. You must create the index field before you map the field.

Type: Array of [DataSourceToIndexFieldMapping \(p. 230\)](#) objects

Array Members: Minimum number of 1 item. Maximum number of 100 items.

Required: No

IncludeAttachmentFilePatterns

List of regular expressions applied to knowledge articles. Items that don't match the inclusion pattern are not indexed. The regex is applied to the field specified in the `PatternTargetField`.

Type: Array of strings

Array Members: Minimum number of 0 items. Maximum number of 100 items.

Length Constraints: Minimum length of 1. Maximum length of 50.

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

ServiceNowServiceCatalogConfiguration

Provides configuration information for crawling service catalog items in the ServiceNow site

Contents

CrawlAttachments

Indicates whether Amazon Kendra should crawl attachments to the service catalog items.

Type: Boolean

Required: No

DocumentDataFieldName

The name of the ServiceNow field that is mapped to the index document contents field in the Amazon Kendra index.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 100.

Pattern: `^[a-zA-Z][a-zA-Z0-9_\.]*$`

Required: Yes

DocumentTitleFieldName

The name of the ServiceNow field that is mapped to the index document title field.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 100.

Pattern: `^[a-zA-Z][a-zA-Z0-9_\.]*$`

Required: No

ExcludeAttachmentFilePatterns

Determines the types of file attachments that are excluded from the index.

Type: Array of strings

Array Members: Minimum number of 0 items. Maximum number of 100 items.

Length Constraints: Minimum length of 1. Maximum length of 50.

Required: No

FieldMappings

Mapping between ServiceNow fields and Amazon Kendra index fields. You must create the index field before you map the field.

Type: Array of [DataSourceToIndexFieldMapping \(p. 230\)](#) objects

Array Members: Minimum number of 1 item. Maximum number of 100 items.

Required: No

IncludeAttachmentFilePatterns

Determines the types of file attachments that are included in the index.

Type: Array of strings

Array Members: Minimum number of 0 items. Maximum number of 100 items.

Length Constraints: Minimum length of 1. Maximum length of 50.

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

SharePointConfiguration

Provides configuration information for connecting to a Microsoft SharePoint data source.

Contents

CrawlAttachments

`TRUE` to include attachments to documents stored in your Microsoft SharePoint site in the index; otherwise, `FALSE`.

Type: Boolean

Required: No

DocumentTitleFieldName

The Microsoft SharePoint attribute field that contains the title of the document.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 100.

Pattern: `^[a-zA-Z][a-zA-Z0-9_.*]*$`

Required: No

ExclusionPatterns

A list of regular expression patterns. Documents that match the patterns are excluded from the index. Documents that don't match the patterns are included in the index. If a document matches both an exclusion pattern and an inclusion pattern, the document is not included in the index.

The regex is applied to the display URL of the SharePoint document.

Type: Array of strings

Array Members: Minimum number of 0 items. Maximum number of 100 items.

Length Constraints: Minimum length of 1. Maximum length of 50.

Required: No

FieldMappings

A list of `DataSourceToIndexFieldMapping` objects that map Microsoft SharePoint attributes to custom fields in the Amazon Kendra index. You must first create the index fields using the [UpdateIndex \(p. 201\)](#) operation before you map SharePoint attributes. For more information, see [Mapping Data Source Fields](#).

Type: Array of [DataSourceToIndexFieldMapping \(p. 230\)](#) objects

Array Members: Minimum number of 1 item. Maximum number of 100 items.

Required: No

InclusionPatterns

A list of regular expression patterns. Documents that match the patterns are included in the index. Documents that don't match the patterns are excluded from the index. If a document matches both an inclusion pattern and an exclusion pattern, the document is not included in the index.

The regex is applied to the display URL of the SharePoint document.

Type: Array of strings

Array Members: Minimum number of 0 items. Maximum number of 100 items.

Length Constraints: Minimum length of 1. Maximum length of 50.

Required: No

SecretArn

The Amazon Resource Name (ARN) of credentials stored in AWS Secrets Manager. The credentials should be a user/password pair. For more information, see [Using a Microsoft SharePoint Data Source](#). For more information about AWS Secrets Manager, see [What Is AWS Secrets Manager](#) in the *AWS Secrets Manager* user guide.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 1284.

Pattern: `arn:[a-z0-9-\.]{1,63}:[a-z0-9-\.]{0,63}:[a-z0-9-\.]{0,63}:[a-z0-9-\.]{0,63}:[^/].{0,1023}`

Required: Yes

SharePointVersion

The version of Microsoft SharePoint that you are using as a data source.

Type: String

Valid Values: `SHAREPOINT_ONLINE`

Required: Yes

Urls

The URLs of the Microsoft SharePoint site that contains the documents that should be indexed.

Type: Array of strings

Array Members: Minimum number of 1 item. Maximum number of 100 items.

Length Constraints: Minimum length of 1. Maximum length of 2048.

Pattern: `^(https?|ftp|file):\\\/\\\/([^\s]*)`

Required: Yes

UseChangeLog

Set to `TRUE` to use the Microsoft SharePoint change log to determine the documents that need to be updated in the index. Depending on the size of the SharePoint change log, it may take longer for Amazon Kendra to use the change log than it takes it to determine the changed documents using the Amazon Kendra document crawler.

Type: Boolean

Required: No

VpcConfiguration

Provides information for connecting to an Amazon VPC.

Type: [DataSourceVpcConfiguration \(p. 231\)](#) object

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

Tag

A list of key/value pairs that identify an index, FAQ, or data source. Tag keys and values can consist of Unicode letters, digits, white space, and any of the following symbols: _ . : / = + - @.

Contents

Key

The key for the tag. Keys are not case sensitive and must be unique for the index, FAQ, or data source.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 128.

Required: Yes

Value

The value associated with the tag. The value may be an empty string but it can't be null.

Type: String

Length Constraints: Minimum length of 0. Maximum length of 256.

Required: Yes

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

TextDocumentStatistics

Provides information about text documents indexed in an index.

Contents

IndexedTextBytes

The total size, in bytes, of the indexed documents.

Type: Long

Valid Range: Minimum value of 0.

Required: Yes

IndexedTextDocumentsCount

The number of text documents indexed.

Type: Integer

Valid Range: Minimum value of 0.

Required: Yes

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

TextWithHighlights

Provides text and information about where to highlight the text.

Contents

Highlights

The beginning and end of the text that should be highlighted.

Type: Array of [Highlight \(p. 244\)](#) objects

Required: No

Text

The text to display to the user.

Type: String

Length Constraints: Minimum length of 1. Maximum length of 2048.

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

TimeRange

Provides a range of time.

Contents

EndTime

The UNIX datetime of the end of the time range.

Type: Timestamp

Required: No

StartTime

The UNIX datetime of the beginning of the time range.

Type: Timestamp

Required: No

See Also

For more information about using this API in one of the language-specific AWS SDKs, see the following:

- [AWS SDK for C++](#)
- [AWS SDK for Go](#)
- [AWS SDK for Java](#)
- [AWS SDK for Ruby V3](#)

Common Errors

This section lists the errors common to the API actions of all AWS services. For errors specific to an API action for this service, see the topic for that API action.

AccessDeniedException

You do not have sufficient access to perform this action.

HTTP Status Code: 400

IncompleteSignature

The request signature does not conform to AWS standards.

HTTP Status Code: 400

InternalFailure

The request processing has failed because of an unknown error, exception or failure.

HTTP Status Code: 500

InvalidAction

The action or operation requested is invalid. Verify that the action is typed correctly.

HTTP Status Code: 400

InvalidClientTokenId

The X.509 certificate or AWS access key ID provided does not exist in our records.

HTTP Status Code: 403

InvalidParameterCombination

Parameters that must not be used together were used together.

HTTP Status Code: 400

InvalidParameterValue

An invalid or out-of-range value was supplied for the input parameter.

HTTP Status Code: 400

InvalidQueryParameter

The AWS query string is malformed or does not adhere to AWS standards.

HTTP Status Code: 400

MalformedQueryString

The query string contains a syntax error.

HTTP Status Code: 404

MissingAction

The request is missing an action or a required parameter.

HTTP Status Code: 400

MissingAuthenticationToken

The request must contain either a valid (registered) AWS access key ID or X.509 certificate.

HTTP Status Code: 403

MissingParameter

A required parameter for the specified action is not supplied.

HTTP Status Code: 400

OptInRequired

The AWS access key ID needs a subscription for the service.

HTTP Status Code: 403

RequestExpired

The request reached the service more than 15 minutes after the date stamp on the request or more than 15 minutes after the request expiration date (such as for pre-signed URLs), or the date stamp on the request is more than 15 minutes in the future.

HTTP Status Code: 400

ServiceUnavailable

The request has failed due to a temporary failure of the server.

HTTP Status Code: 503

ThrottlingException

The request was denied due to request throttling.

HTTP Status Code: 400

ValidationError

The input fails to satisfy the constraints specified by an AWS service.

HTTP Status Code: 400

Common Parameters

The following list contains the parameters that all actions use for signing Signature Version 4 requests with a query string. Any action-specific parameters are listed in the topic for that action. For more information about Signature Version 4, see [Signature Version 4 Signing Process](#) in the *Amazon Web Services General Reference*.

Action

The action to be performed.

Type: string

Required: Yes

Version

The API version that the request is written for, expressed in the format YYYY-MM-DD.

Type: string

Required: Yes

X-Amz-Algorithm

The hash algorithm that you used to create the request signature.

Condition: Specify this parameter when you include authentication information in a query string instead of in the HTTP authorization header.

Type: string

Valid Values: AWS4-HMAC-SHA256

Required: Conditional

X-Amz-Credential

The credential scope value, which is a string that includes your access key, the date, the region you are targeting, the service you are requesting, and a termination string ("aws4_request"). The value is expressed in the following format: *access_key/YYYYMMDD/region/service/aws4_request*.

For more information, see [Task 2: Create a String to Sign for Signature Version 4](#) in the *Amazon Web Services General Reference*.

Condition: Specify this parameter when you include authentication information in a query string instead of in the HTTP authorization header.

Type: string

Required: Conditional

X-Amz-Date

The date that is used to create the signature. The format must be ISO 8601 basic format (YYYYMMDD'T'HHMMSS'Z'). For example, the following date time is a valid X-Amz-Date value: 20120325T120000Z.

Condition: X-Amz-Date is optional for all requests; it can be used to override the date used for signing requests. If the Date header is specified in the ISO 8601 basic format, X-Amz-Date is not required. When X-Amz-Date is used, it always overrides the value of the Date header. For more information, see [Handling Dates in Signature Version 4](#) in the *Amazon Web Services General Reference*.

Type: string

Required: Conditional

X-Amz-Security-Token

The temporary security token that was obtained through a call to AWS Security Token Service (AWS STS). For a list of services that support temporary security credentials from AWS Security Token Service, go to [AWS Services That Work with IAM](#) in the *IAM User Guide*.

Condition: If you're using temporary security credentials from the AWS Security Token Service, you must include the security token.

Type: string

Required: Conditional

X-Amz-Signature

Specifies the hex-encoded signature that was calculated from the string to sign and the derived signing key.

Condition: Specify this parameter when you include authentication information in a query string instead of in the HTTP authorization header.

Type: string

Required: Conditional

X-Amz-SignedHeaders

Specifies all the HTTP headers that were included as part of the canonical request. For more information about specifying signed headers, see [Task 1: Create a Canonical Request For Signature Version 4](#) in the *Amazon Web Services General Reference*.

Condition: Specify this parameter when you include authentication information in a query string instead of in the HTTP authorization header.

Type: string

Required: Conditional

AWS glossary

For the latest AWS terminology, see the [AWS glossary](#) in the *AWS General Reference*.